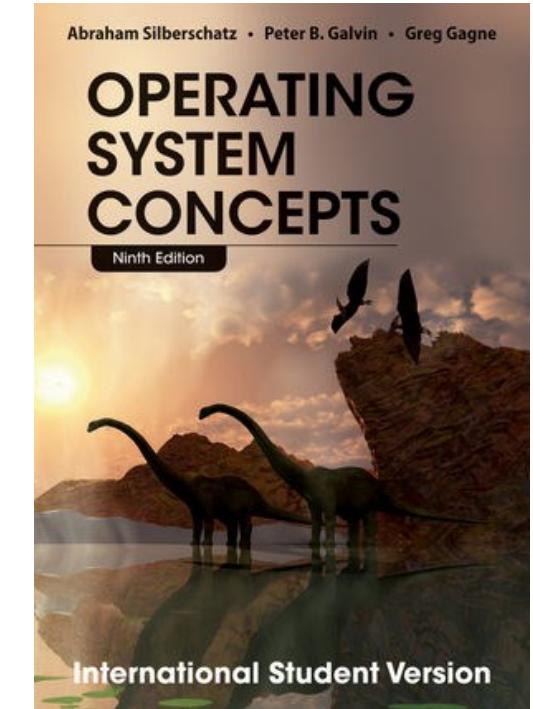


Chapter 9: Memory-Management Strategies

Dept. of Software, Gachon Univ.
Joon Yoo



Chapter 9: Memory Management Strategies

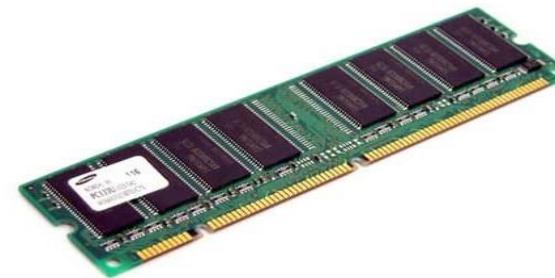
- **Introduction**
- Virtual Memory
- Contiguous Memory Allocation
- Paging

Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To learn the basic concept of virtual memory
- To discuss various memory-management techniques, including paging

Recall: Memory (Ch. 1)

- Address
 - relative position in memory
- Contents
 - the data stored in a memory cell
- The size of each memory cell is one BYTE



**1000 Memory Cells
in Main Memory**

Address	Memory Contents
0	-27.2
1	354
2	0.005
3	-26
4	H
.	.
.	.
998	X
999	75.62

Recall: Memory (Ch. 1)

Basics: Actually, memory looks like ...

Address

9278

9279

9280

9281

...

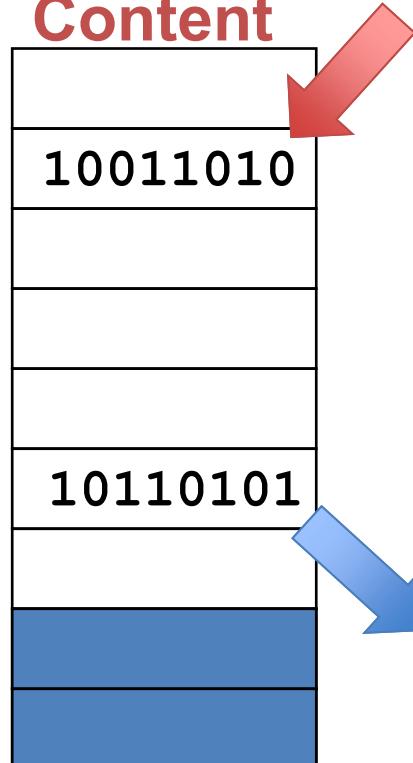
```
00001001001011100110011001100100101101100011001010000100100100010011011000110010101100011011101000111  
01010111001001100101001100010010111001100011001000010100110011101100011011000110011001001011111  
01100011011011110110101110000011010010110110001100101011100001011100011101000001010001011100111  
00110110010101100011011101000110101110110001100100001001001000100101110011001010111000010100010111100  
01110100001000100000101000000100100101110011000010100101110110010100101110011001010111100001000000001010000000  
1010000001001001011100110011101100011011101100010011000010110110001000000011011010110000101101001  
011011100000101000001001011100110011100001100101000011001010000101001000001101101011000010110  
100101100000010010111001100111011000110111011000100110000110110101100010011101011000010100  
0000100100101110011001110110001100100001001001001110011000010010010111001100101101011000010110  
1001011011000111010000010100000010010001001000000101000000101000000101000000101000000101000000101000000  
0101010101000101001000110010000001100000000010100000001010000000101000000010100000001010000000101000000  
01010111000001011101000110000001011101000100111000001011100000101110000010111000001011100000101110000  
000010100000010010001001000000101000000010100000001010000000101000000010100000001010000000101000000  
00011000000000010100000001001011101100111000000010010111011001110000000100101110011001010111001111  
00011000000000010100000001001011101100111000000010010111011001110000000100101110011001010111001111  
011001110000000101101001100110000001011010011000000101101001100000010110100110000001011010011000000  
011001110000000101101001100110000001011010011000000101101001100000010110100110000001011010011000000  
00110000000001010000000100101110011001110000000100101110011001110000000100101110011001010111001111  
00110000000001010000000100101110011001110000000100101110011001110000000100101110011001010111001111  
011001110000000101101001100110000001011010011000000101101001100000010110100110000001011010011000000  
00110000000001010000000100101110011001110000000100101110011001110000000100101110011001010111001111  
0010111000000001010000000100101110011001110000000100101110011001110000000100101110011001010111001111  
110001011011000100101100110011100000010110100110001101000101110100001010000010010110110001100100  
001000000010110110010010110011001110000001011010011000110000001011101001011000100101101110001111  
0000000000101000000010010110010010110011000000101101001001011001100000010110001001011001100000010100  
0101110100101100000100101101110011000100000010100000010100101100010010110010010000001001010111001111  
1111001100000001011000001001011011100110000001001011000100101101110011000000000101000001001011100111  
0111010000100000000100101101110011000000000101100000000101100000000101100000000101100000000101100000000  
100001011101000001001011011001011011100110000000001011000000001011000000001011000000001011000000001010000  
0000010100000001001011000100100000001011000000001011000000001011000000001011000000001011000000001010000  
00000000001010000000100101110010010111000000010010111000000010010111000000010010111000000010010111000000  
000001001011000000100101110010010111000000010010111000000010010111000000010010111000000010010111000000  
0110011001010011000000100101110010010111000000010010111000000010010111000000010010111000000010010111000000  
011011010110000001011010010110111000000010110100101101110000000101101001011011100000001011010010110111000000  
110101100000010110100101101110000000101100000010110010110010110010110010110010110010110010110010110010110000  
0010000000101000000010010111001001011100000001001011100000001001011100000001001011100000001001011100000000  
0010000000101100000010010111001001011100000001001011100000001001011100000001001011100000001001011100000000
```

Recall: Memory (Ch. 1)

- **Basics: Storage and Retrieval of Information in Memory**

- Data storage (Write)
 - ▶ Setting the individual bits of a memory cell to 0 or 1, destroying its previous contents
- Data retrieval (Read)
 - ▶ Copying the contents of a particular memory cell to another storage area

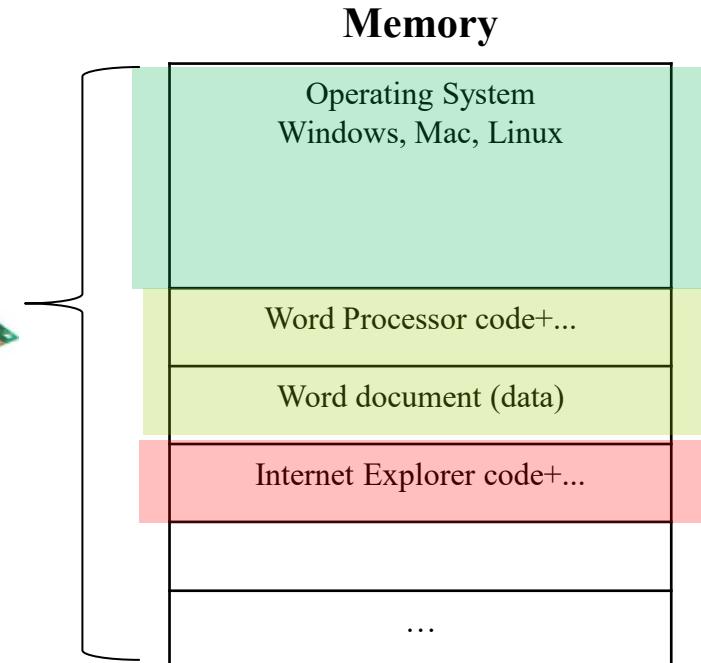
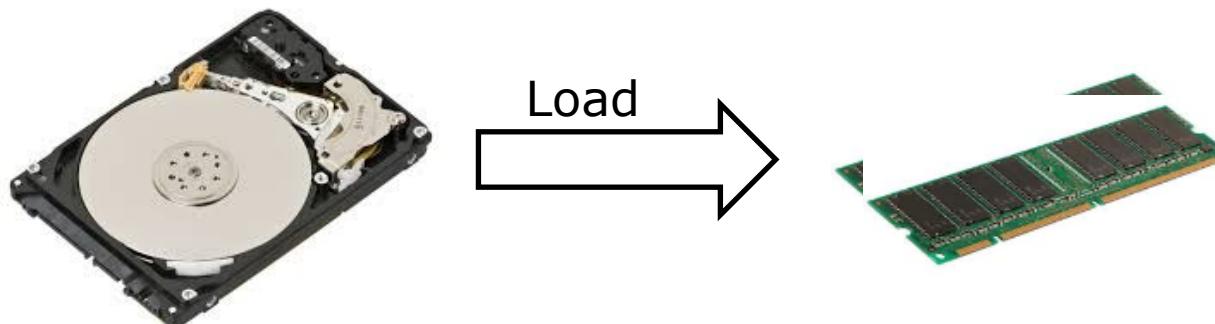
Address	Content
9278	
9279	10011010
9280	
9281	
9282	
9283	10110101
9284	
9285	
9286	



Recall: DRAM (main memory)

RAM 메모리

- DRAM (Dynamic random access memory)
 - **Programs** are usually stored in the **hard disk**
 - ▶ When programs are executed (= process), it is loaded on the **main memory** – faster than disk
 - Similarly, when **file** is **open**, it is loaded on the main memory



DRAM: main memory

- Why load program/data from disk to memory?
 - DRAM access time $\approx 50\text{ns}$
 - Disk access time $\approx 5,000,000\text{ns}$
 - If we use disks to run programs it will be $100,000\times$ slower
 - ▶ ~~SSD~~ access time $\approx 100,000\text{ns}$, so 2000x slower than DRAM



100,000x

SRAM: Cache Memory

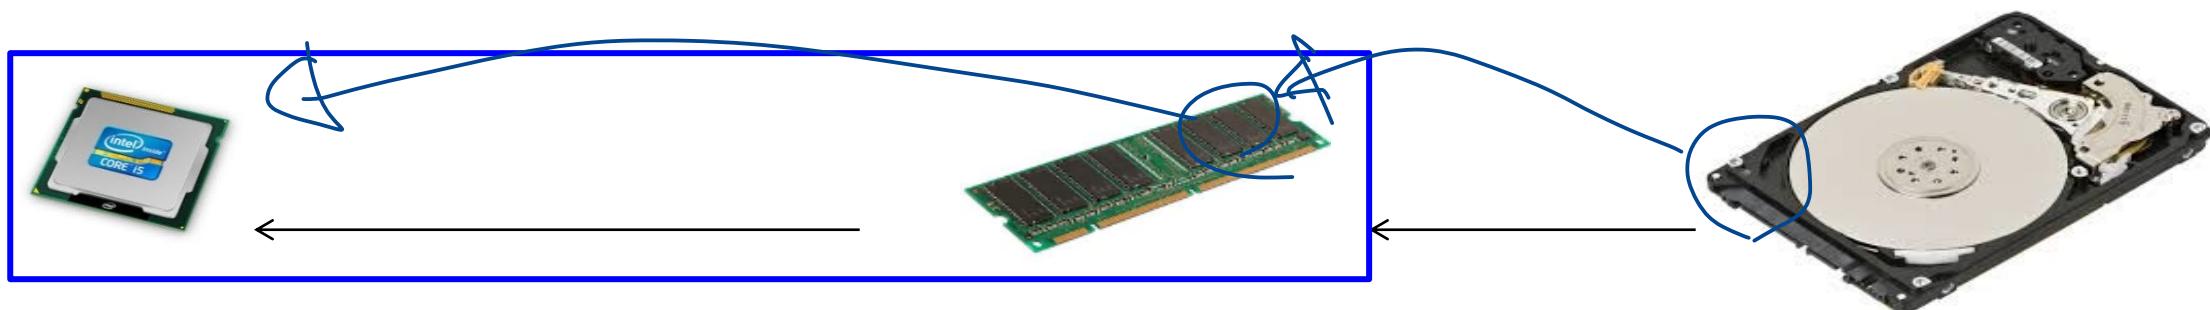
- But DRAM is still much slower than CPU
 - It takes about **100~400 CPU cycles** to access DRAM
 - Use a small but very fast memory = **Cache Memory!!**
 - SRAM access time: **only few ns**



- Computer hardware system takes care of cache
- OS view: A single memory. (does not know cache)

Introduction

- Memory is **central** to the operation of a computer system
- A typical **CPU instruction-execution cycle**
 - 1. The CPU **fetches** an **instruction** from **memory**
 - 2. The instruction is then decoded and may cause operands (data) to be fetched from memory
 - 3. the result may be **stored back in memory**



Summary

IMPORTANT

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are the only storage that CPU can access directly
 - CPU cannot directly access disk, network, printers, ... – **must** be copied from/to memory first
- Speed?
 - Register access in one CPU clock (or less)
 - Main memory can take many cycles, causing a **stall**
 - ▶ Note: Faster **cache memory** sits between main memory and CPU registers. **OS view of cache is same as main memory!!**
 - ▶ Thus we will not be dealing much with cache here (you will be learning more on cache in Computer Architecture)

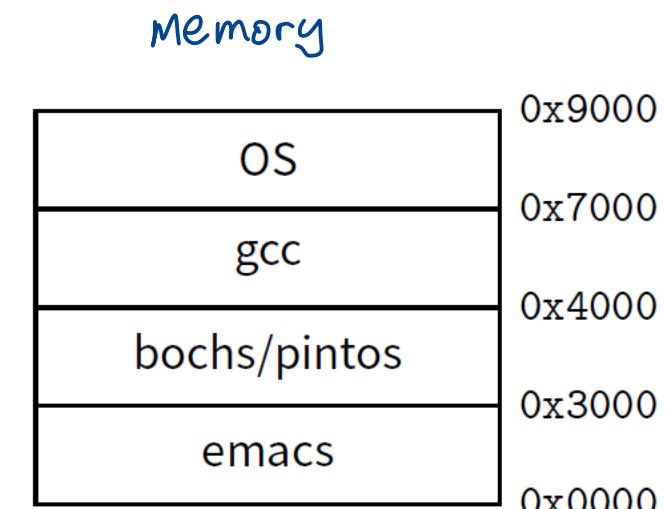
Chapter 9: Memory Management Strategies

- Introduction
- **Virtual Memory**
- Contiguous Memory Allocation
- Paging

Sharing/Protecting the Memory

- Kernel, System/user Processes, Data share the main memory
 - When does gcc have to know it will run at 0x4000? → **Binding** 어디로 가야 하는가? 어디의 어디로?
 - If pintos has an error and writes to address 0x7100? → **Protection** → 사용하지 않도록 한다. Protect!
 - What happens if pintos needs to expand? → **Memory Allocation/Paging**
 - If emacs needs more memory than is on the machine? → **Demand Paging (Ch. 10)**

물리적 공간이 더 작아도 실행시킬 수 있음



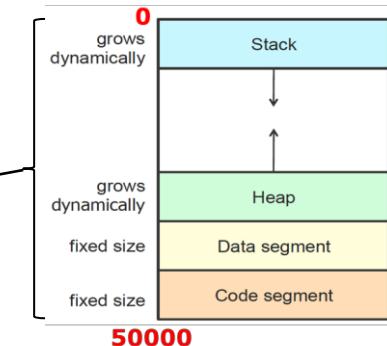
“가상 메모리”를 사용한다!!!

- Answer: Each process uses a separate **Logical (or Virtual) address** instead of actual physical memory address!

Virtual vs. Physical Address Space (Ch. 9.1.3)

■ Logical address (= Virtual address)

- Each user process has individual virtual address
(An address generated by **CPU**)
- Always starts from **address 0**
- **Logical Address space**
 - ▶ Set of all logical addresses generated by a process

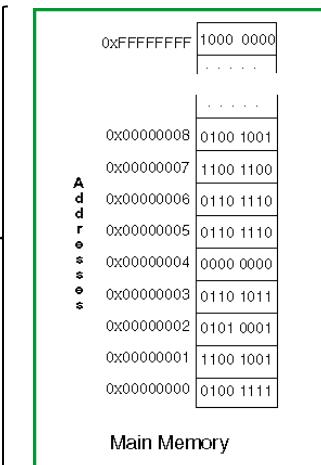


설계에서 주소는 0부터 시작합니다.

■ Physical address

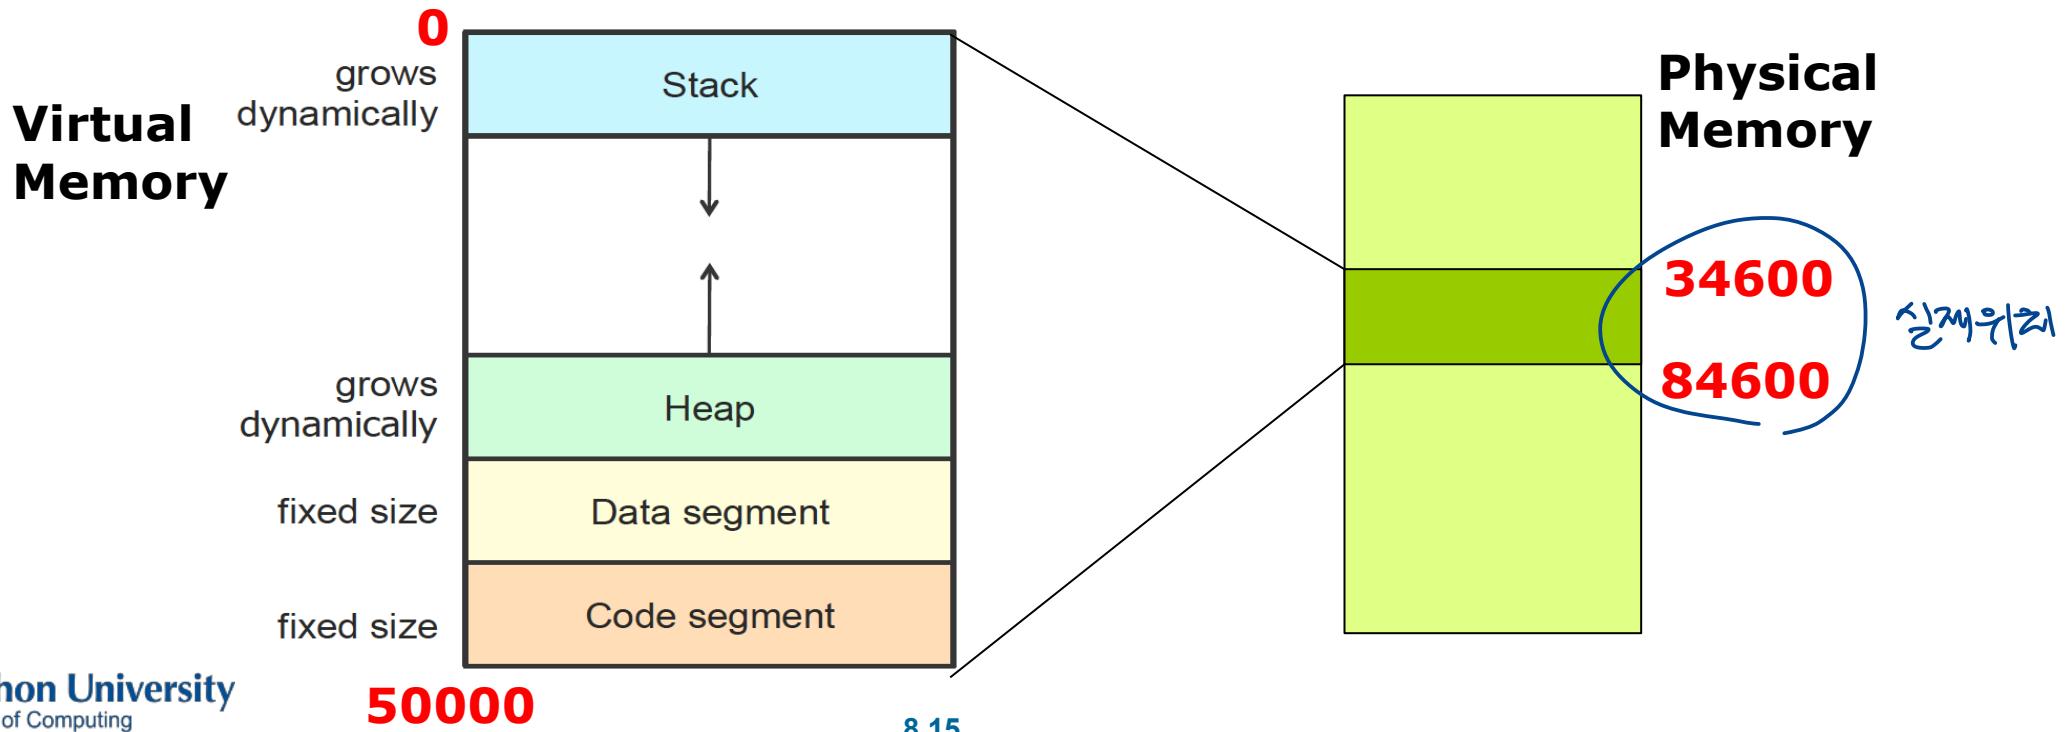
진짜 주소값. 물리적인 위치를 의미한다.

- Address seen by the **physical memory unit**
- **Physical Address space**
 - ▶ Set of all physical addresses



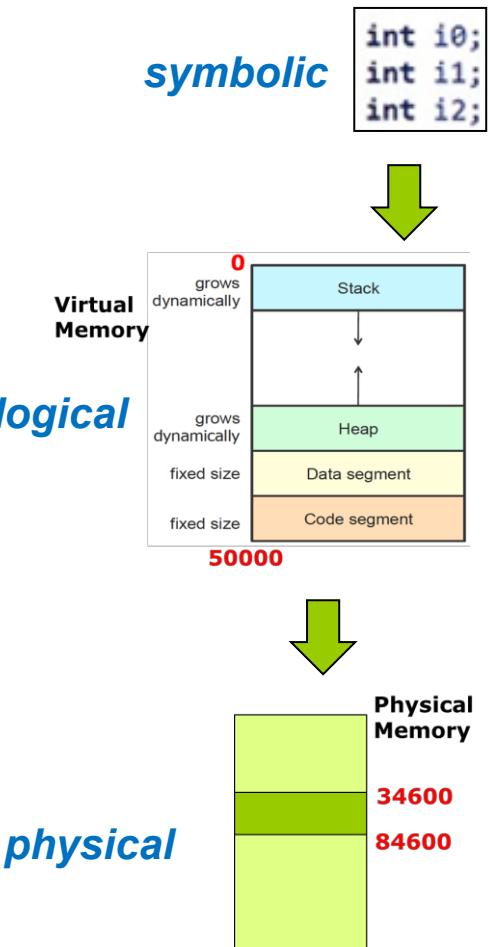
Address Space

- A **logical (virtual) address** is a memory address that a process uses to access its own memory
 - Logical address ≠ actual physical RAM address
 - When a process accesses a virtual address, the virtual address must be translated into a physical address = **Address binding**



Address Binding (Ch. 9.1.2)

- **Address binding:** mapping between *logical* and *physical* addresses:
가상주소와 — 실제주소와 위치를 맵하는 것
- When should we do address binding?
 - **Address binding** can happen at three different stages:
 - ▶ Compile time binding
 - ▶ Load time binding
 - ▶ Run time binding



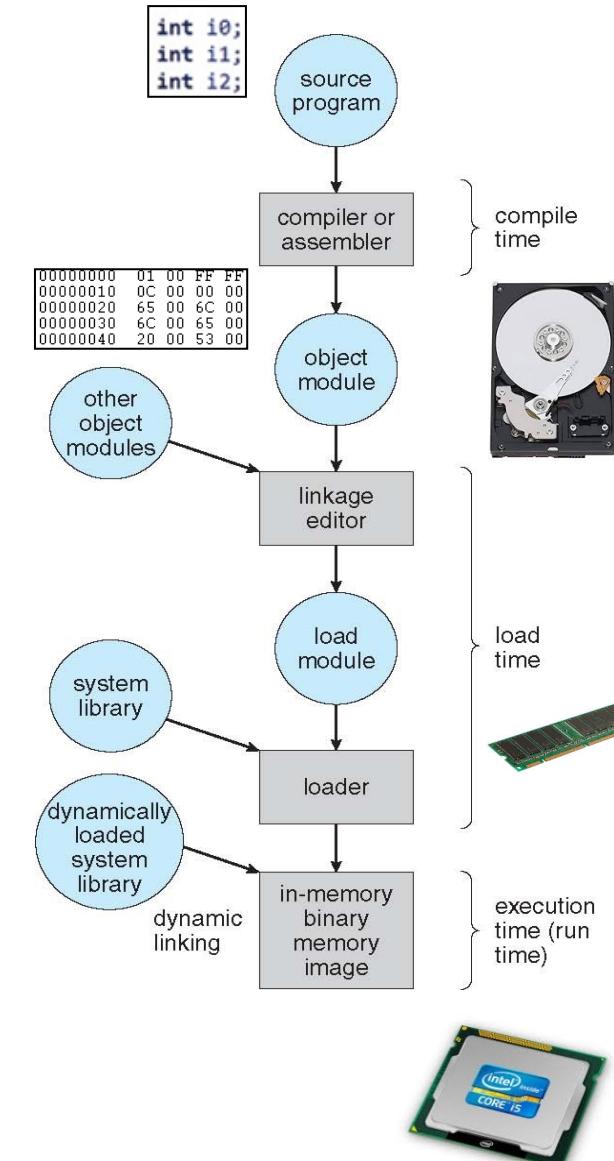
Address Binding

■ Compile time

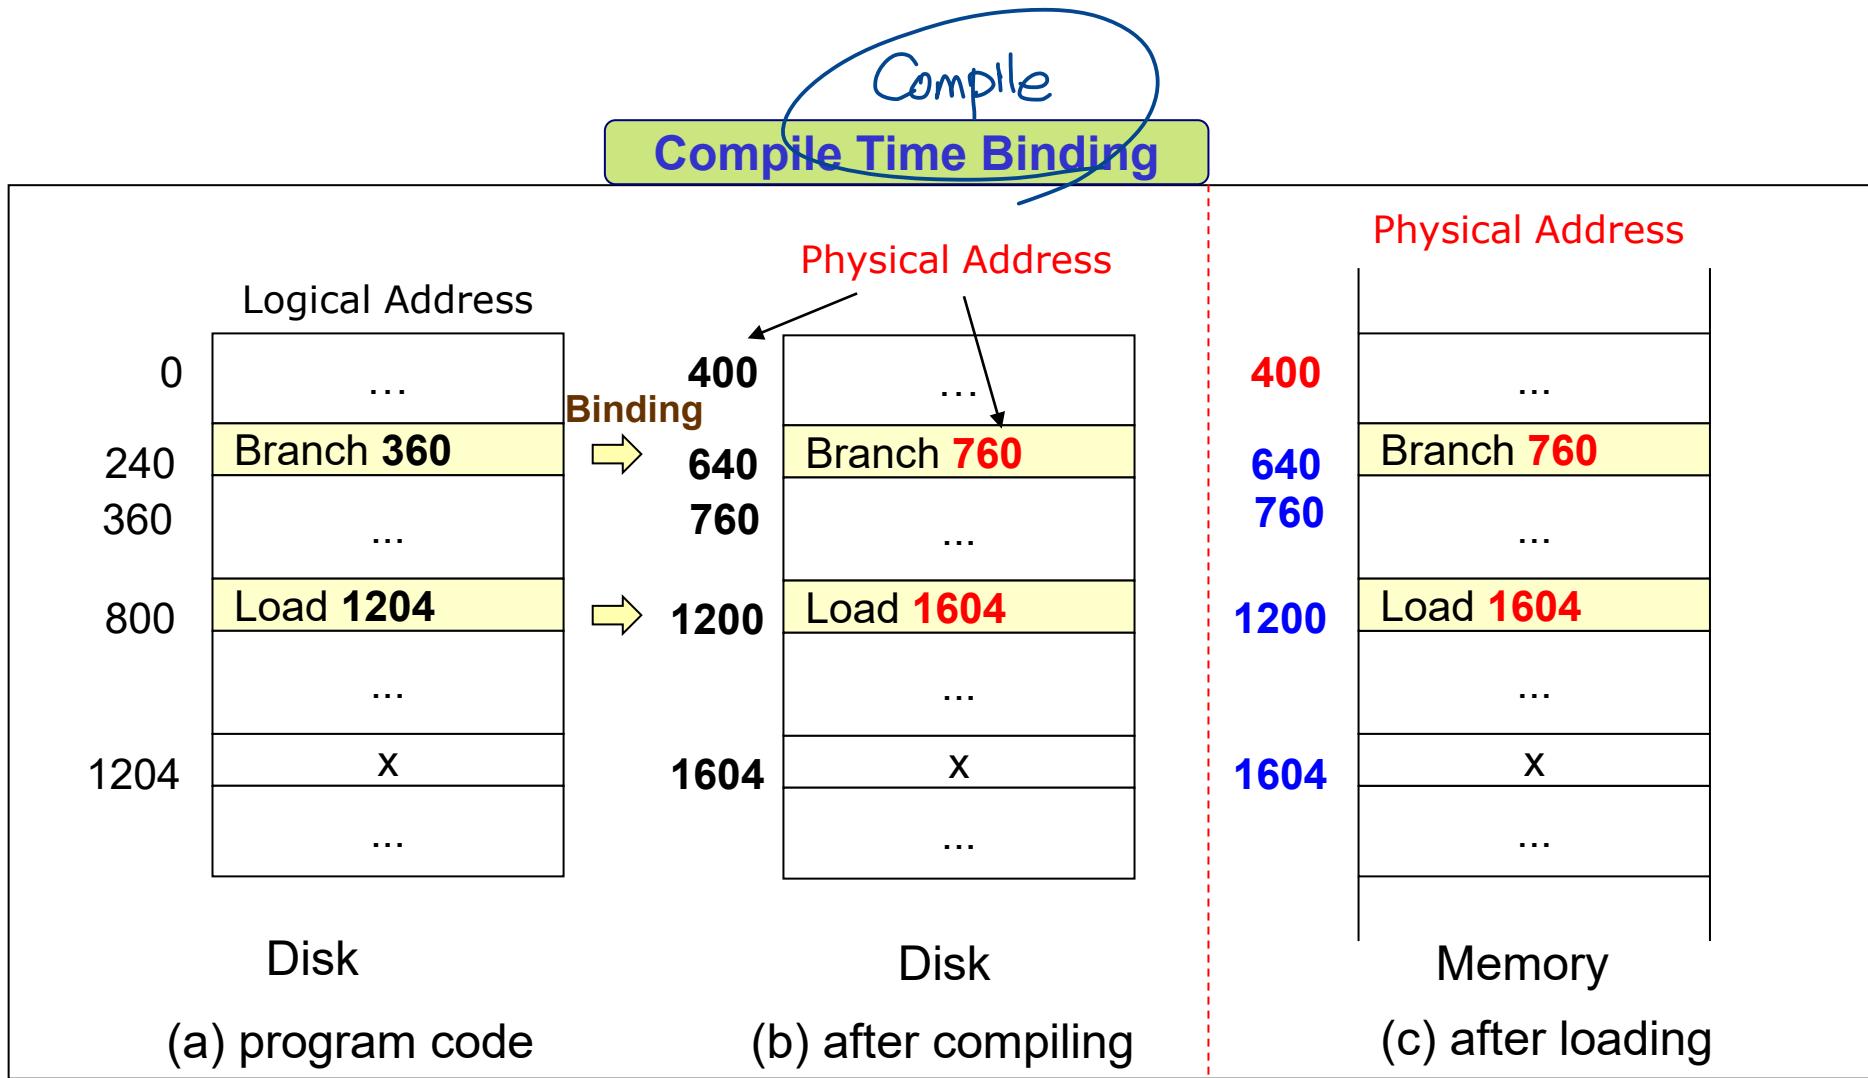
- Compiler: If memory location known at compile time, **absolute code** can be generated
- If starting physical address changes must **recompile** code

컴파일 시점에서 physical address가 지정됨

주소가 바뀌는 경우 recompile 해야 함



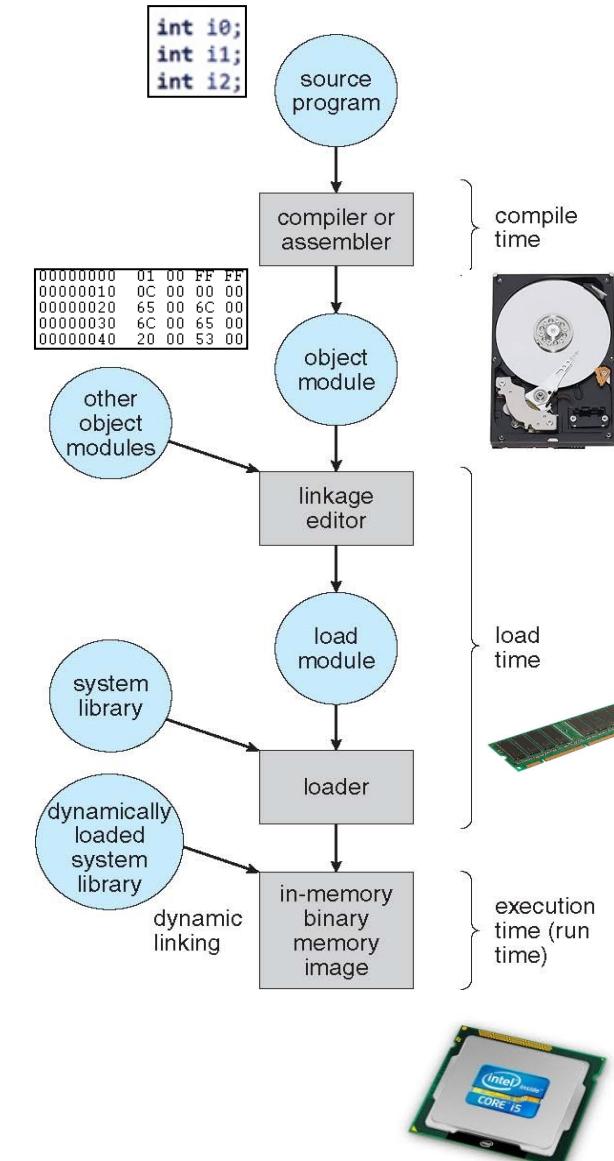
Example: Compile Time Binding



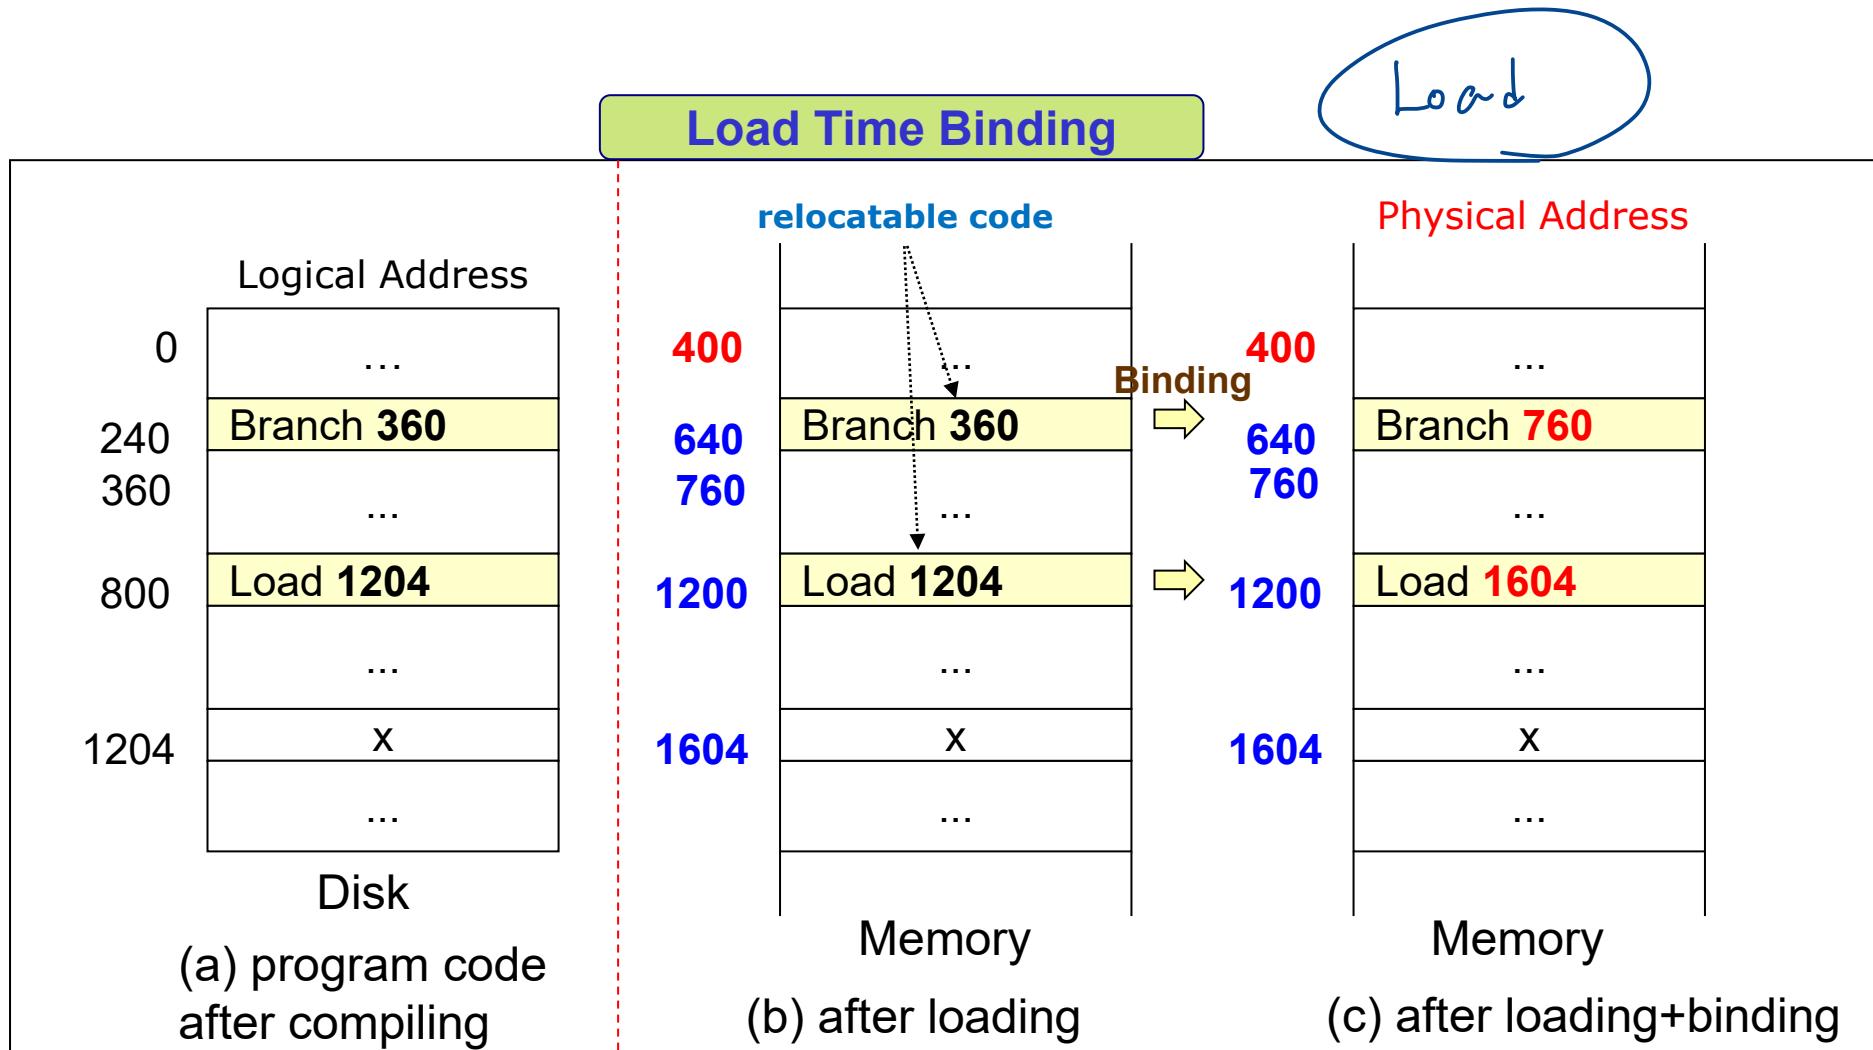
Address Binding

■ Load time

- **Loader:** If memory location is not known at compile time, compiler must generate **relocatable code**
 - ▶ can change code during load time
- Address binding is delayed until **load time**
- If starting address changes, need to **reload** user code



Example: Load Time Binding

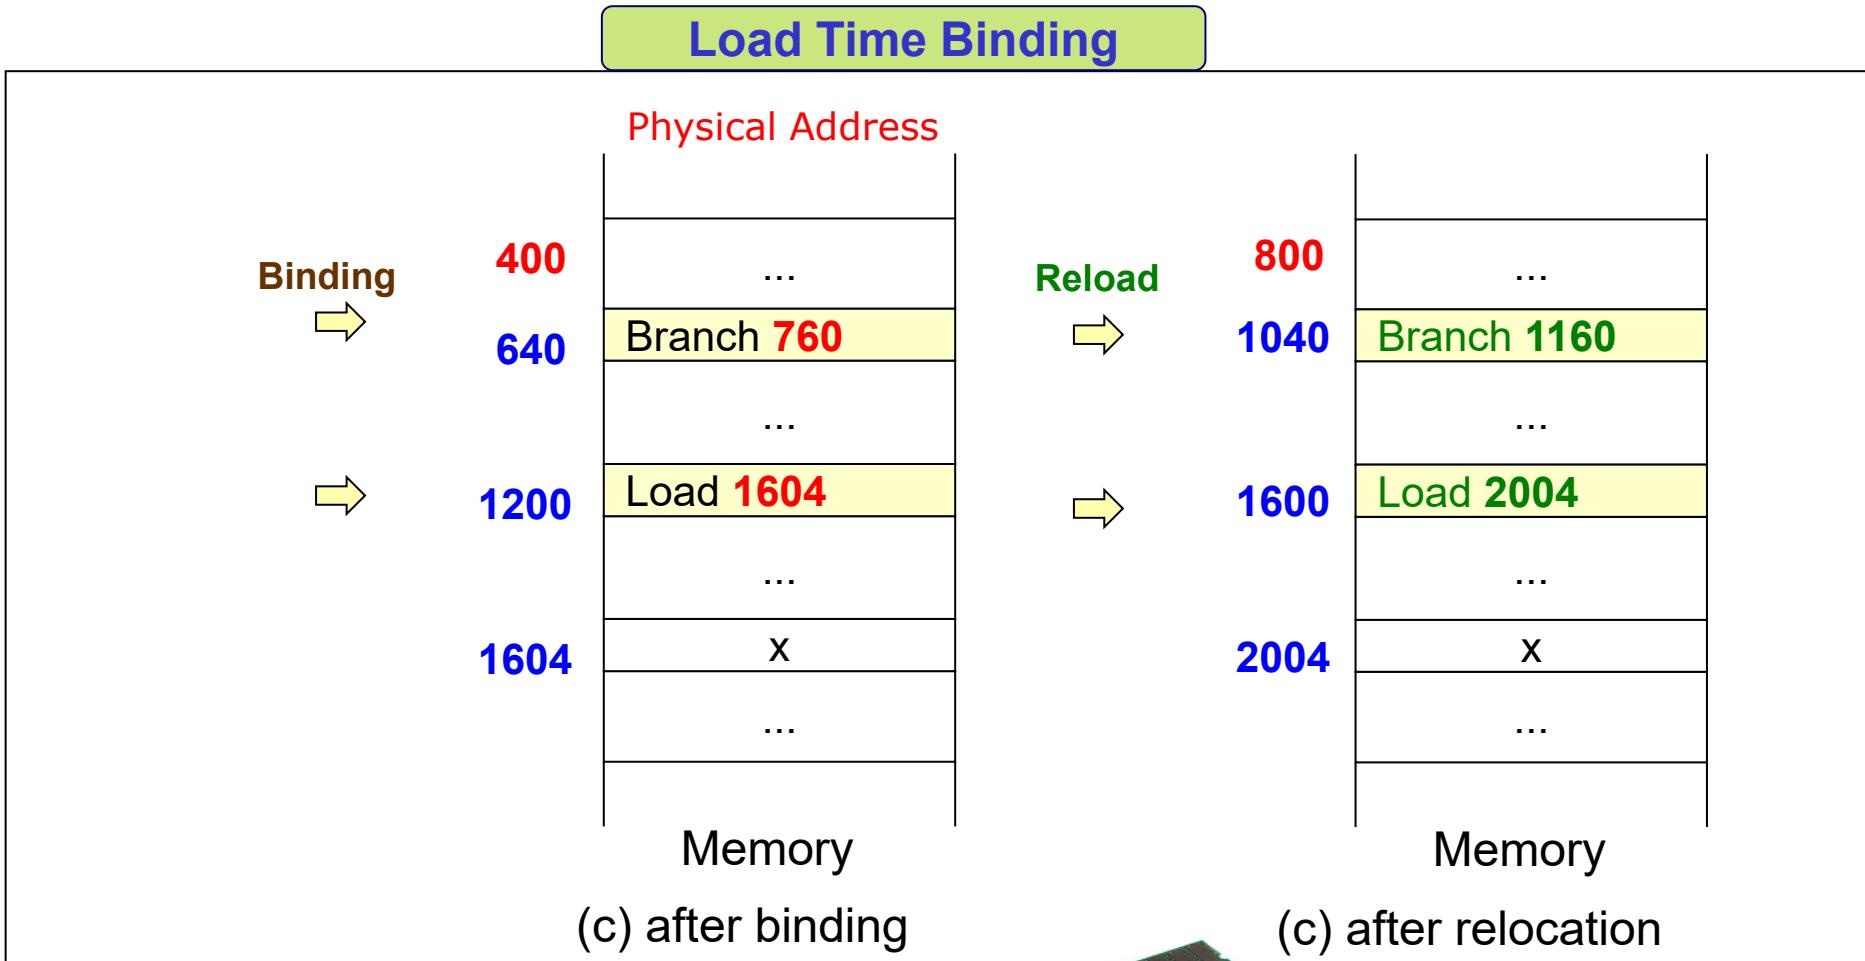


Address Binding

- Q: What if starting address changes to 800?
 - Compile time binding: need to recompile code
 - Load time binding: need to reload

Example: Load Time Binding

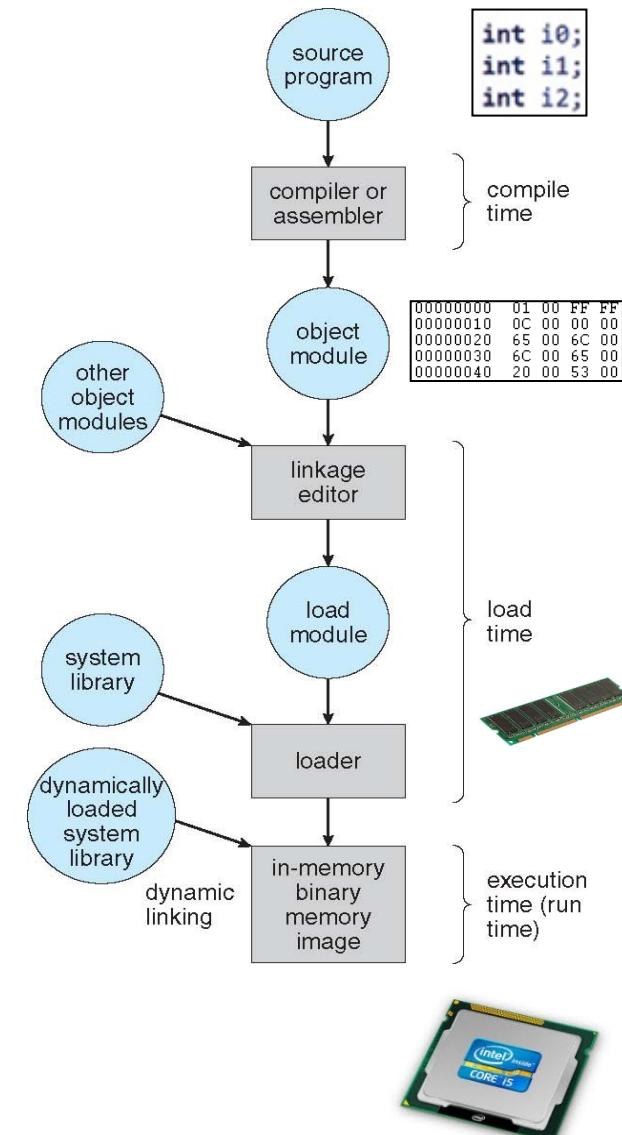
Q: What if starting address changes from 400 to 800? Need to reload



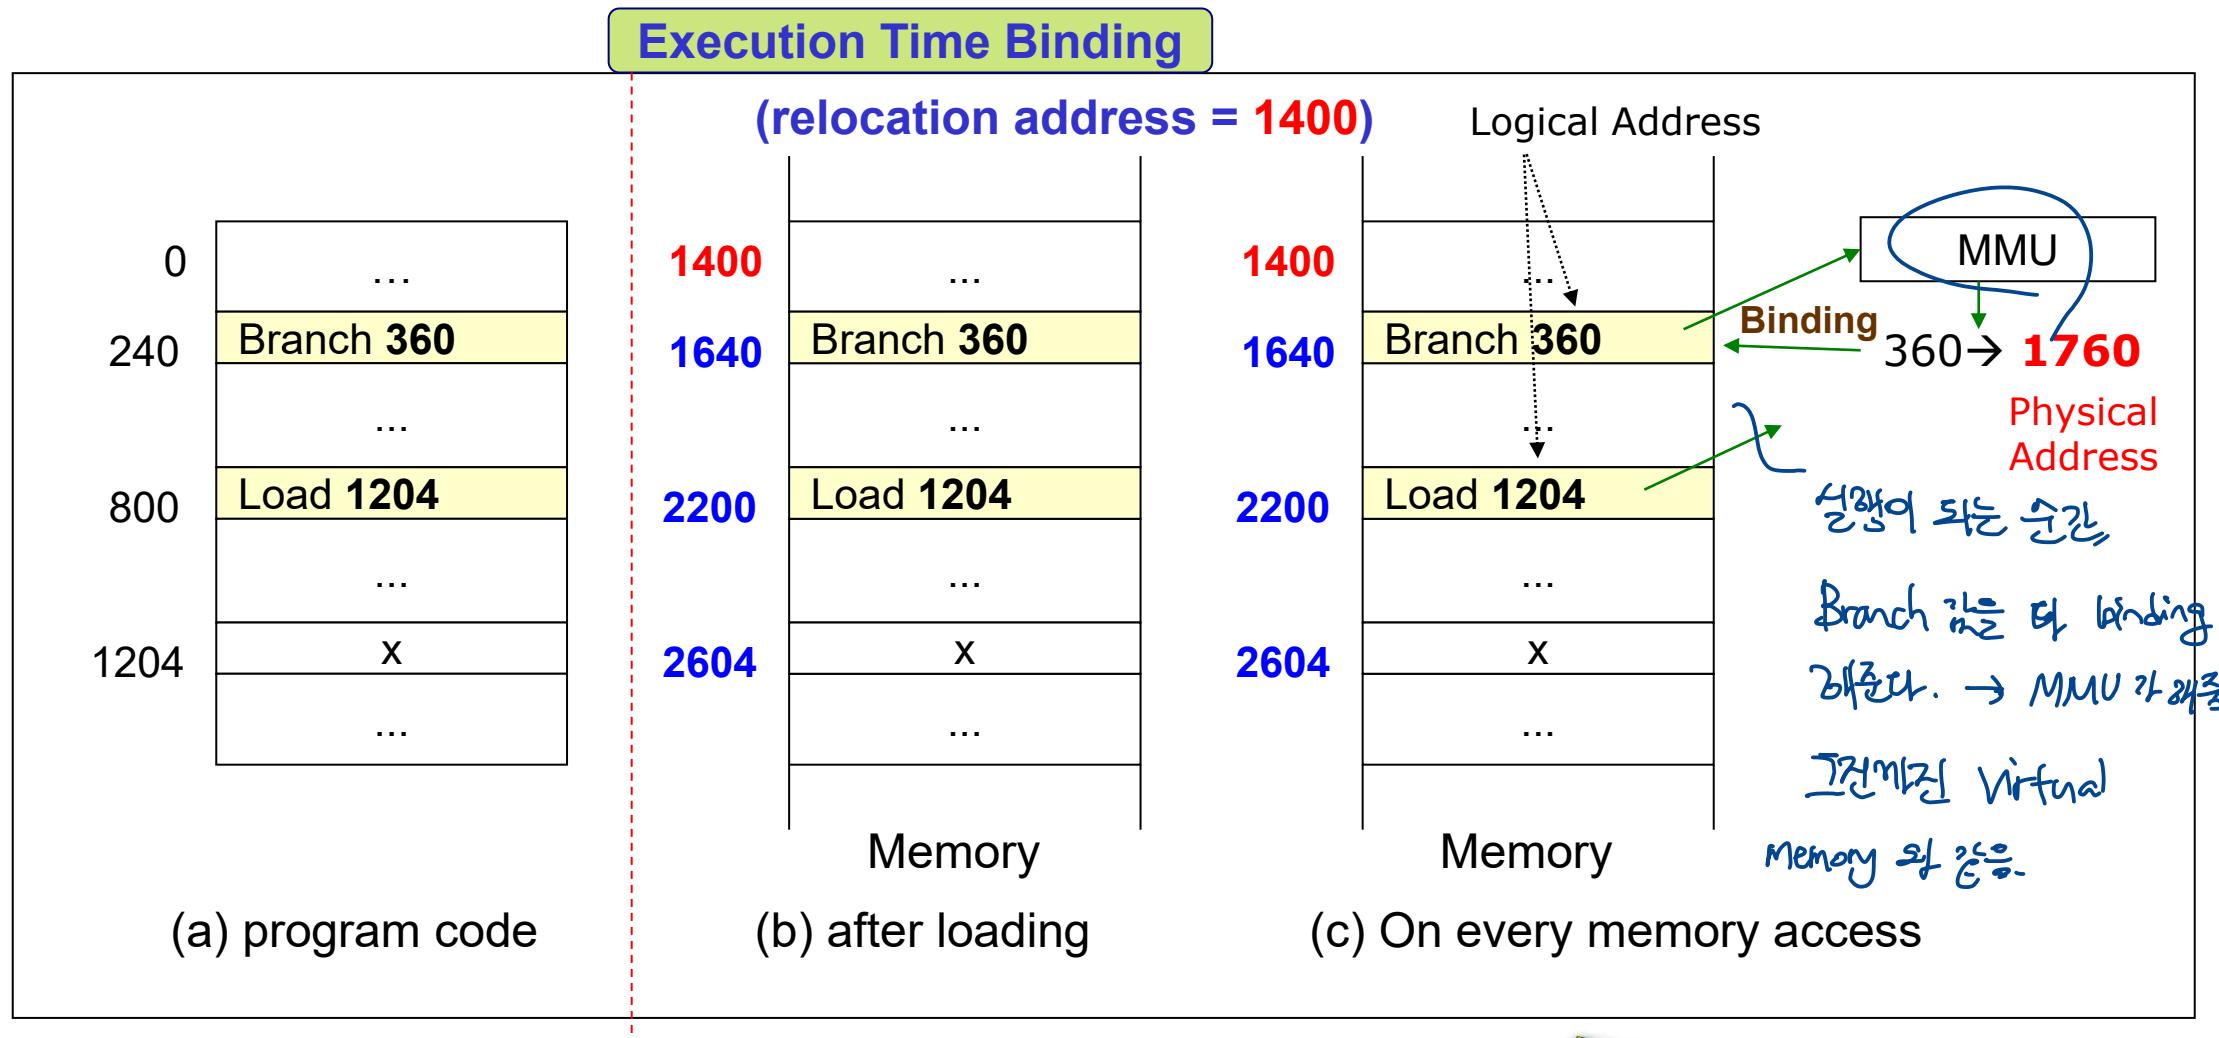
Run Time (Execution Time) Binding

Run time (Execution Time)

- process memory location can be moved during execution
- Address binding delayed until **run time**
- Need hardware support for address maps (e.g., MMU such as base and limit registers)
- Most OS's uses this method**



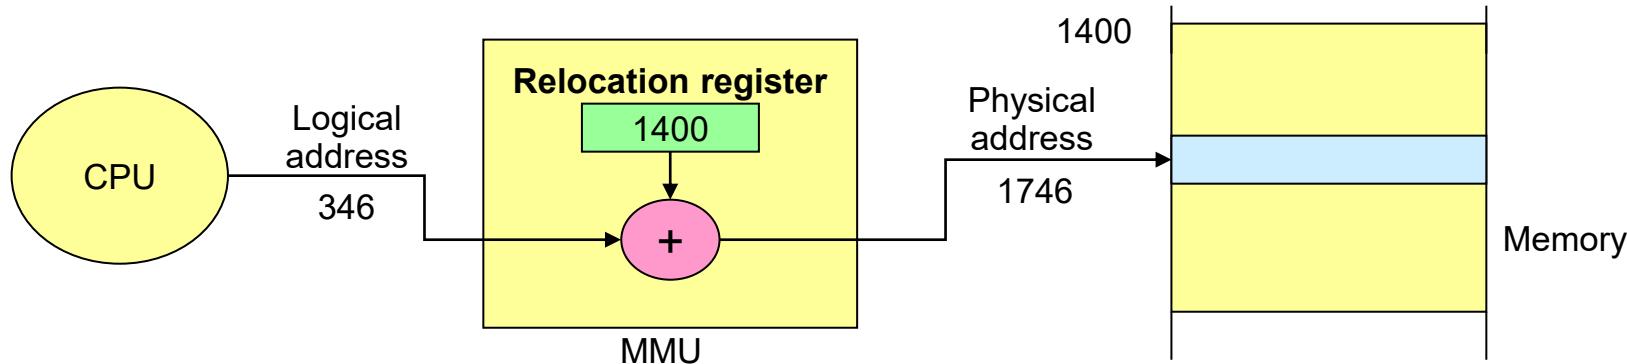
Example: Run Time binding



MMU (Memory Management Unit)

- **MMU** (Memory Management Unit)

- **Hardware** supported unit that maps from logical(virtual) address to physical address in **Run-time binding**



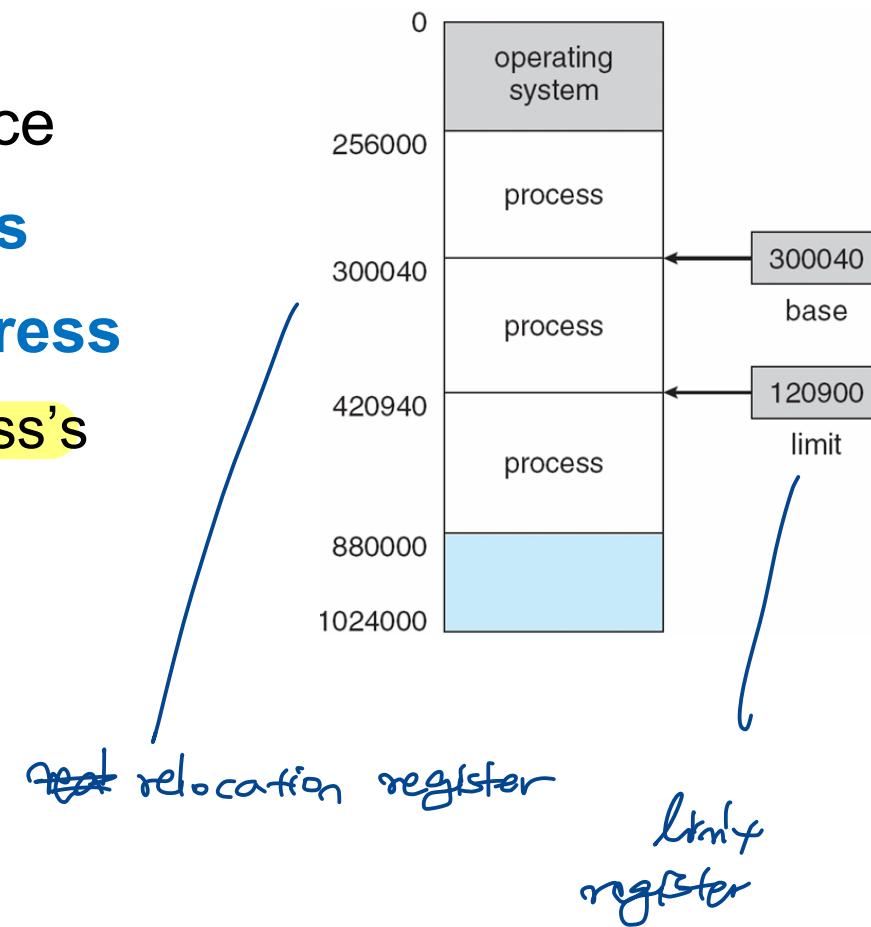
즉시 변환!!

Runtime 할 ..

Q: What if relocation address changes to 1800?

Relocation and Limit Registers (Ch. 9.1.1)

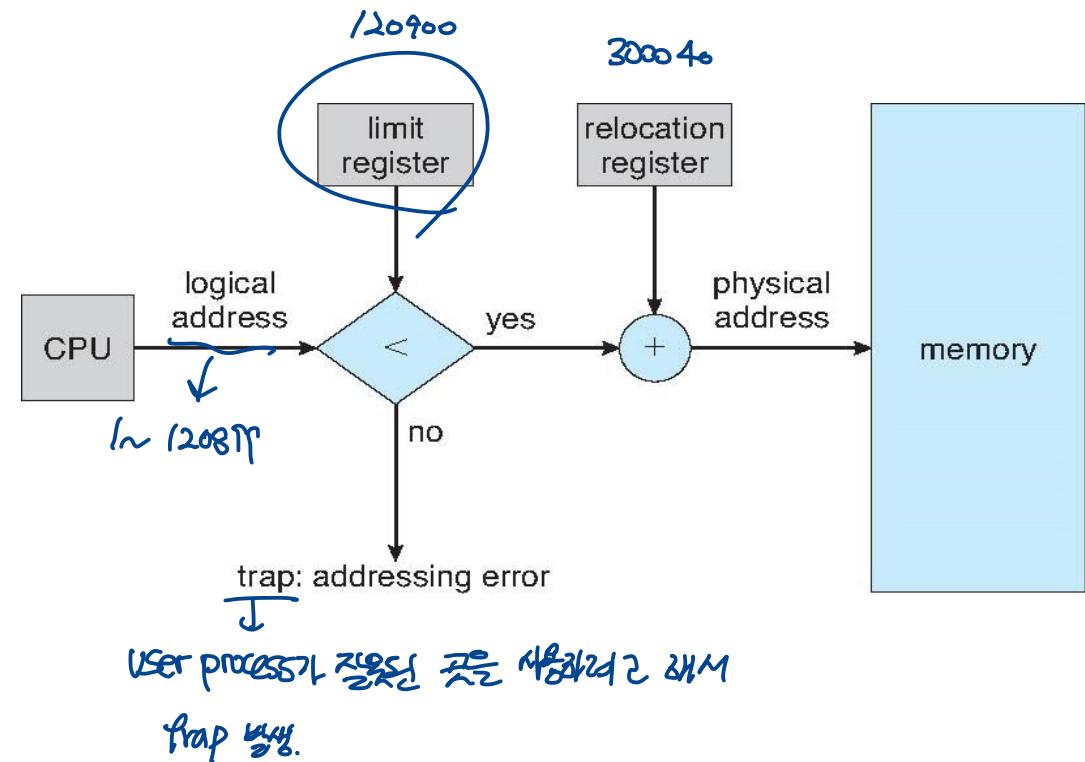
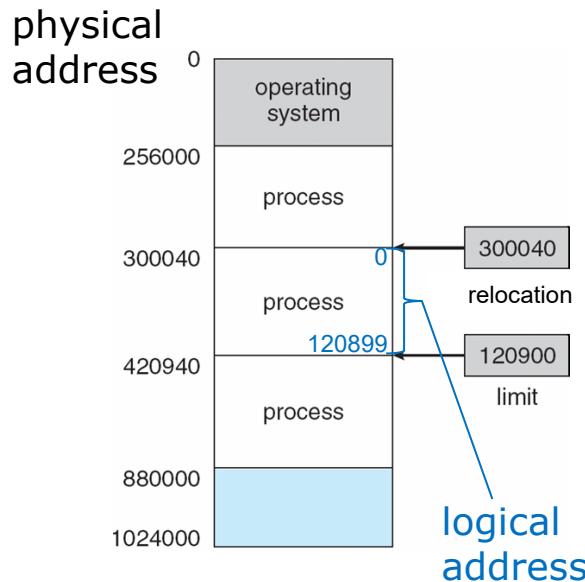
- A pair of **relocation** (= **base**) **register** and **limit register** define the **logical address space** of each process
- Each process has separate memory space
 - User process **only sees logical address**
 - User process **cannot see physical address**
 - User process **cannot access other process's memory space = provides protection**



Memory Protection

Memory Protection: Relocation-register scheme

- A relocation register and a limit register define a logical address space
- Address protection with relocation and limit registers



Chapter 9: Memory Management Strategies

- Introduction
- Virtual Memory
- **Contiguous Memory Allocation**
- Paging

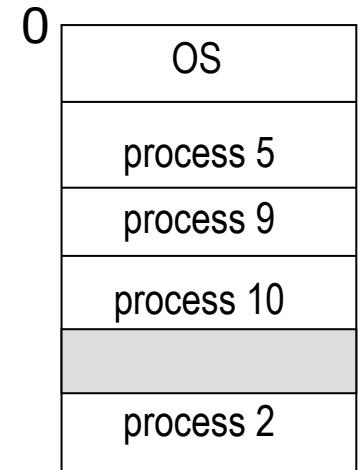
Memory Allocation

■ Memory Allocation

- Allocate memory to OS and user processes in an efficient way possible

■ Main memory usually into two partitions:

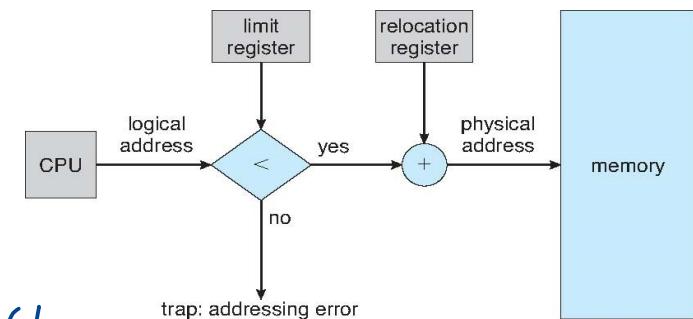
- Resident **operating system**
 - usually held in low memory with interrupt vector
- User **processes**
 - held in high memory



■ Contiguous Memory Allocation

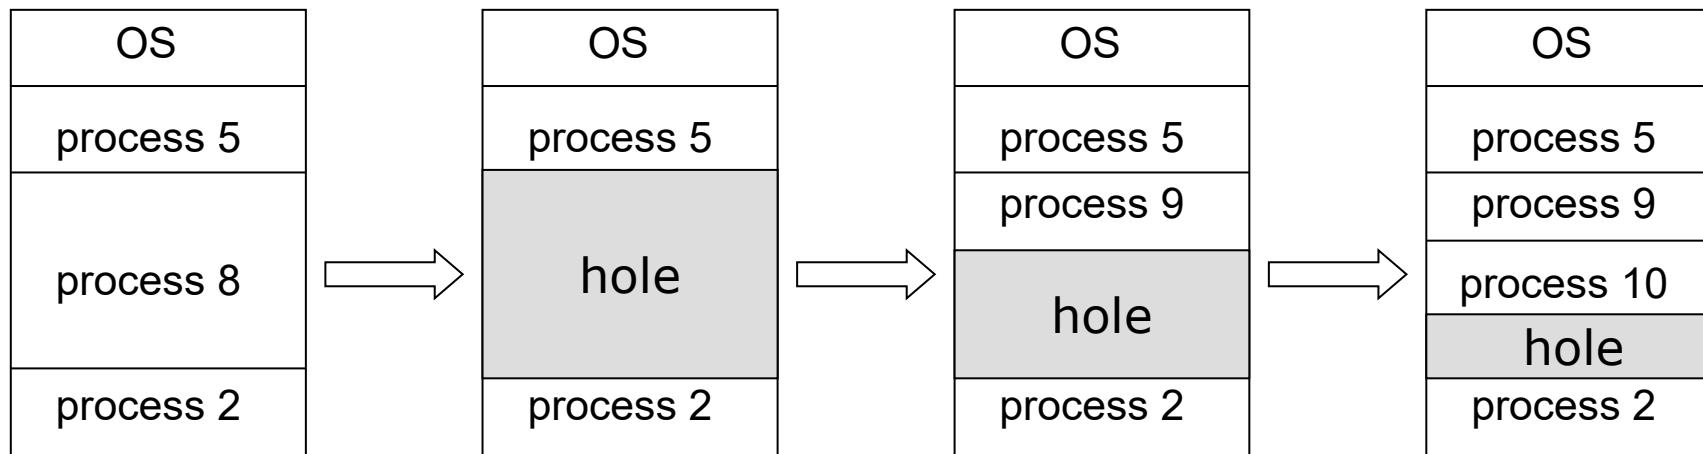
- Allocate available memory to processes
- Each process is contained in a single section of contiguous memory

연속적인 메모리 공간을 할당합니다!



Memory Allocation: Contiguous Allocation

- Multiple-partition allocation
 - Degree of multiprogramming limited by number of partitions
 - **Variable-partition** sizes for efficiency (sized to a given process' needs)
 - **Hole** – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions combined
 - Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)



Contiguous Allocation: Example

- Example

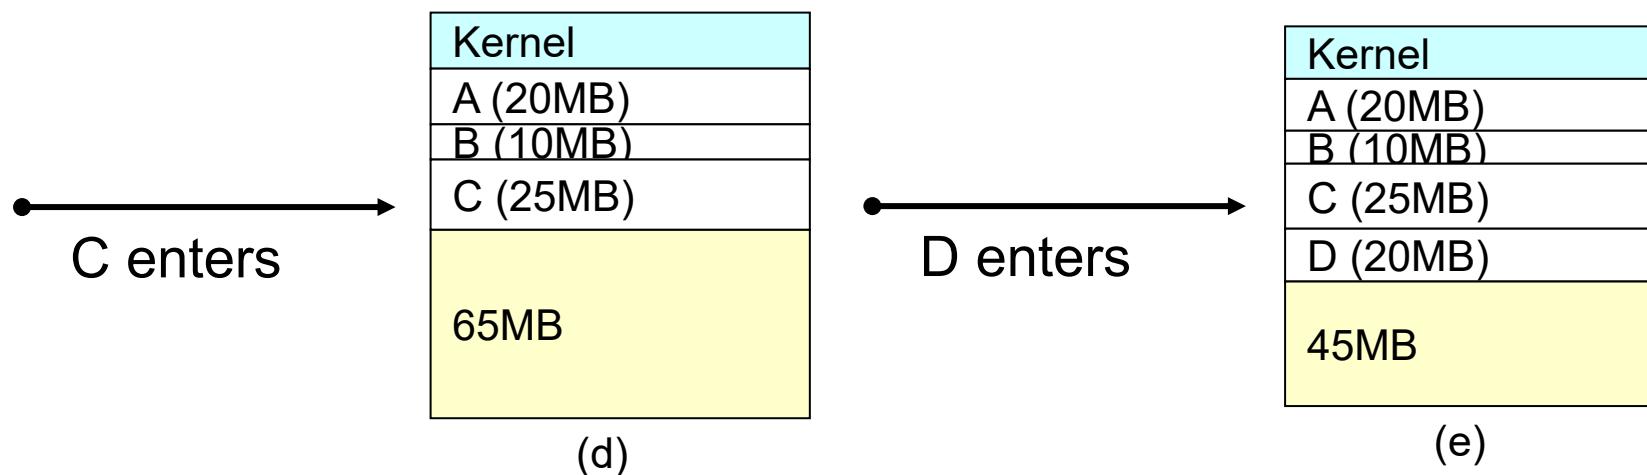
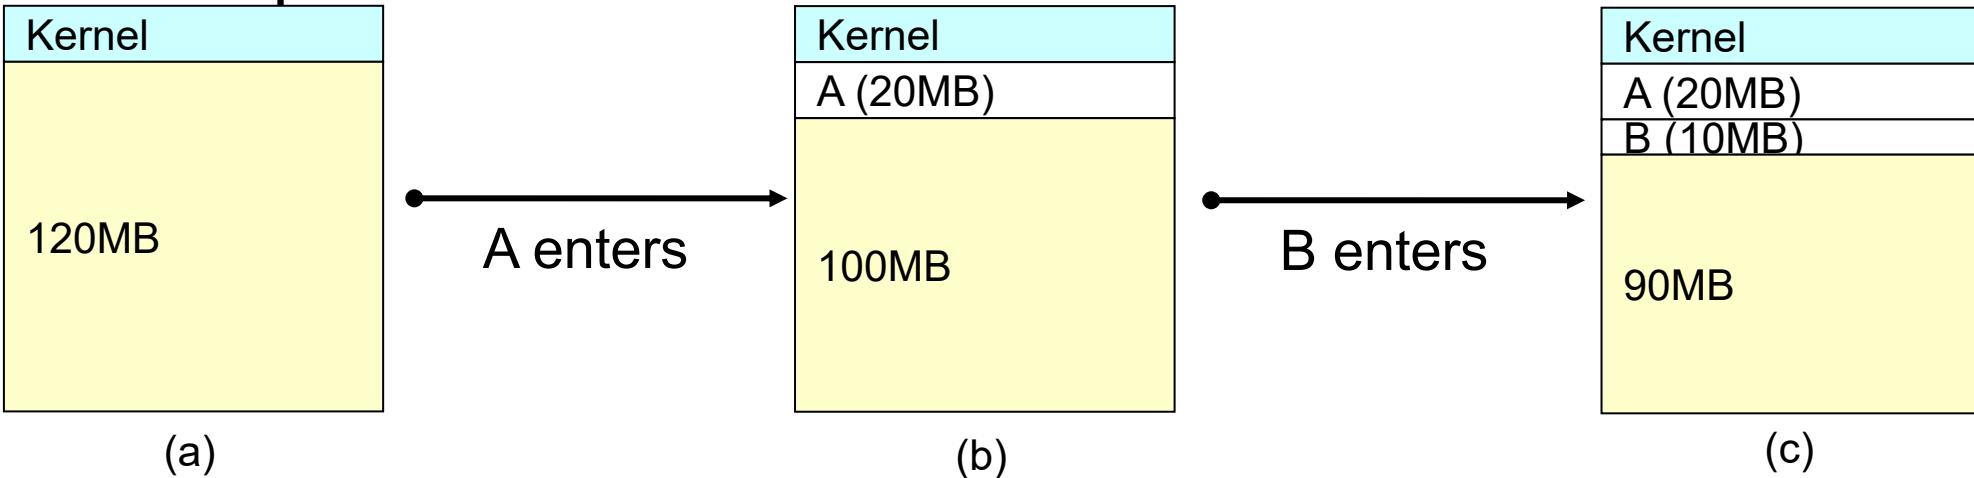
Memory allocation and partition scenario

- ◆ Assumption

- Memory space: 120 MB
- (a) : Initial state
- (b) : After loading process A(20MB)
- (c) : After loading process B(10MB)
- (d) : After loading process C(25MB)
- (e) : After loading process D(20MB)
- (f) : After process B releases memory
- (g) : After loading process E(15MB)
- (h) : After process D releases memory

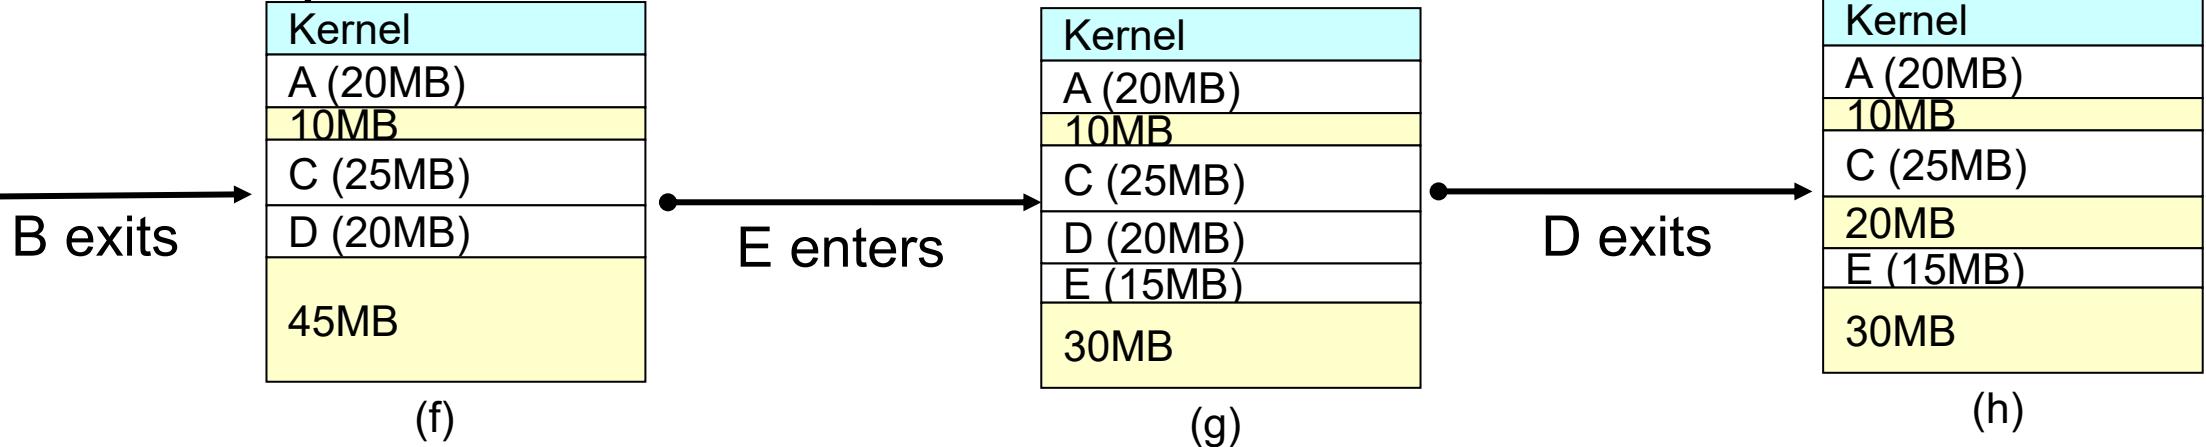
Contiguous Allocation: Example

- Example



Contiguous Allocation: Example

- Example



- **Hole** – block of available memory; **holes of various size are scattered throughout memory**

빠질 때 구멍이 송송...

Contiguous Allocation Pros/Cons

■ Pros

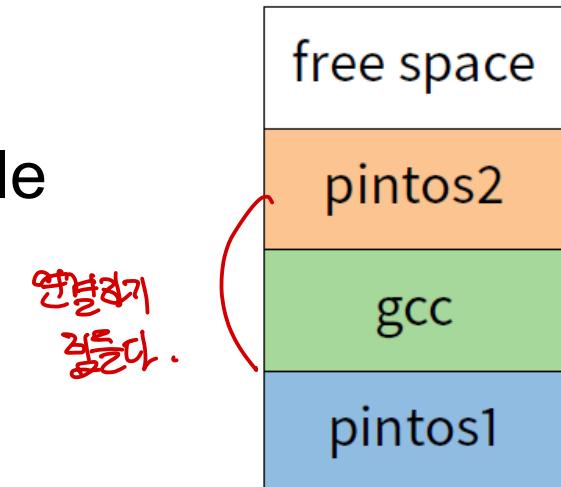
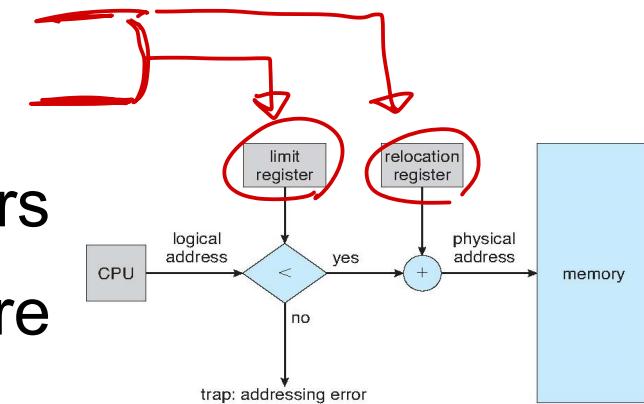
- Cheap in terms of hardware: only two registers
- Cheap in terms of cycles: do add and compare in parallel

빠르고, 연산 비용이 저렴하다. *logical address*

→ 각 process에獨立

■ Cons

- Growing a process is expensive or impossible
- No way to share code or data (E.g., two processes both running pintos)
- External Fragmentation (next two slides)



External Fragmentation

■ External Fragmentation

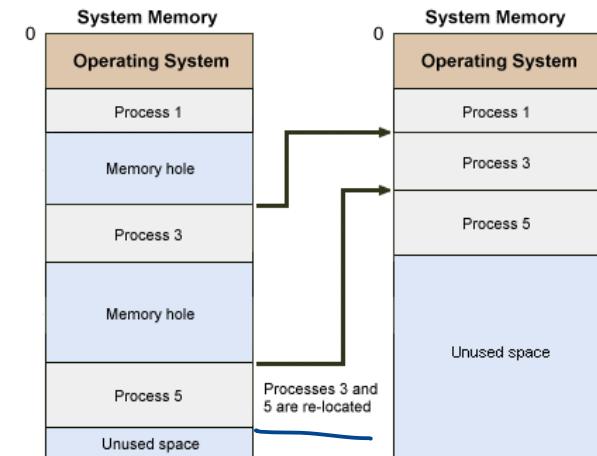
- Total memory space exists to satisfy a request, but it is not contiguous → *hole의 끝은 총Enough hole 각각은 부족한데*

F (40MB)?
Kernel
A (20MB)
10MB
C (25MB)
20MB
E (15MB)
30MB

■ Solution1: Compaction <hole 해결기>

- Shuffle memory contents to place all free memory together in one large block
- Compaction is possible *only* if relocation is dynamic, and is done at execution time

▶ **NOTE:** We cannot preempt memory space if logical memory is determined in load time or compile time



External Fragmentation

- Solution2: ***noncontiguous allocation***

- Permit the logical address space of the processes to be Allowing a process to be allocated physical memory wherever such memory is available
- e.g., **Paging**

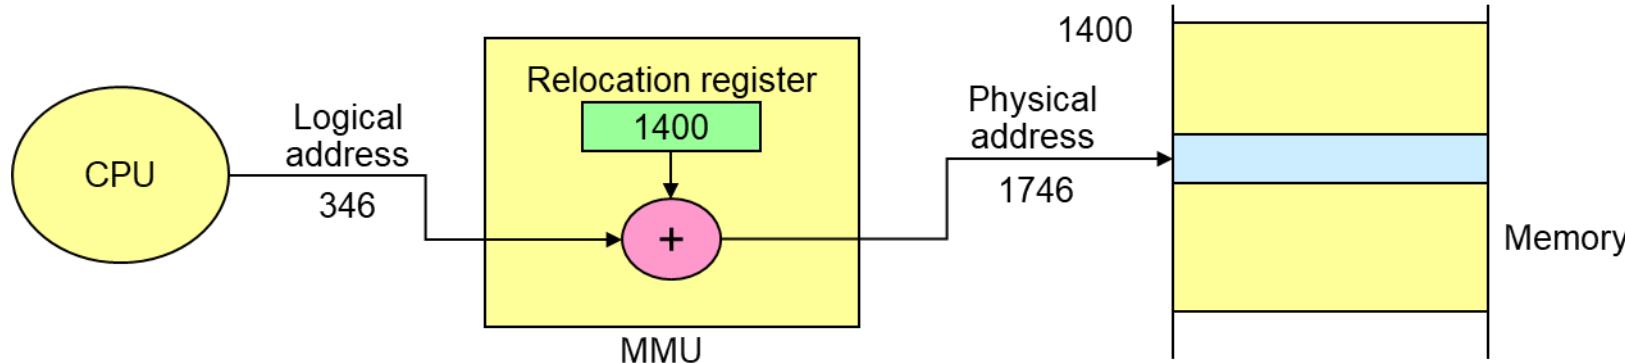
메모리를 조각시
할당한다.

Chapter 9: Memory Management Strategies

- Introduction
- Virtual Memory
- Contiguous Memory Allocation
- Paging 

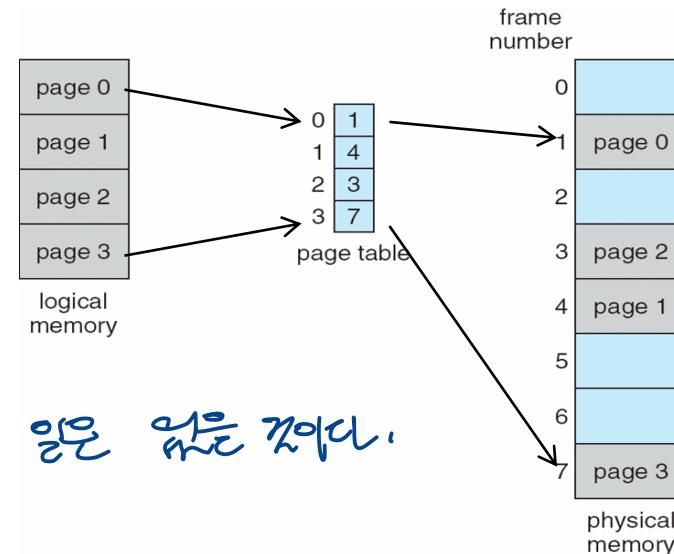
Contiguous vs. Non-contiguous Memory Allocation

- Contiguous 치고 치고 ...



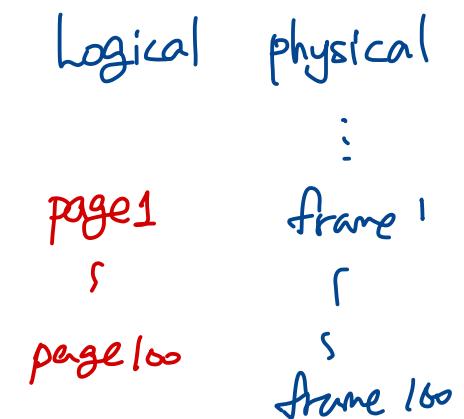
- Non-contiguous
 - No external fragmentation
 - no need for compaction

공간이 있는데 넣지 못하는 일은 없을 것이다.



Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes || same-size
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation



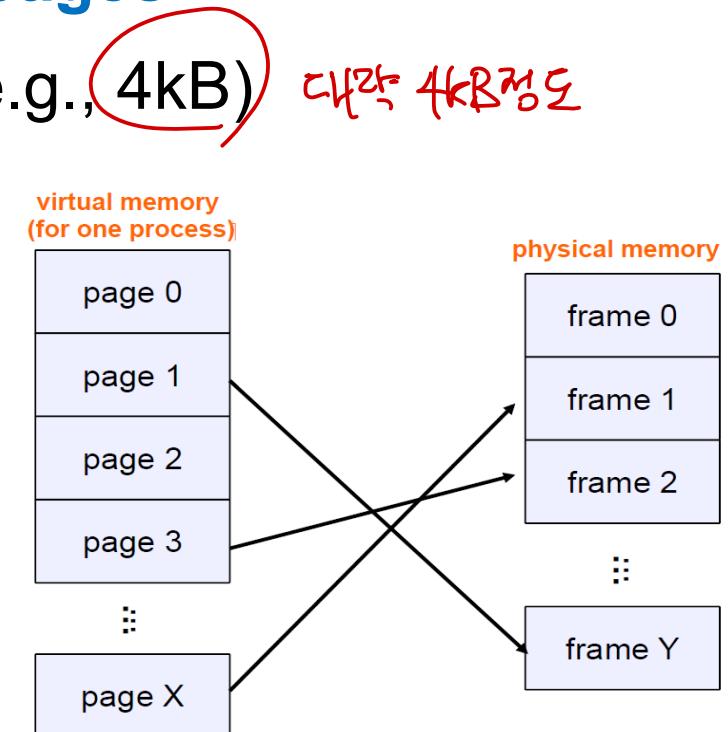
Paging : Basic Method

▪ Paging

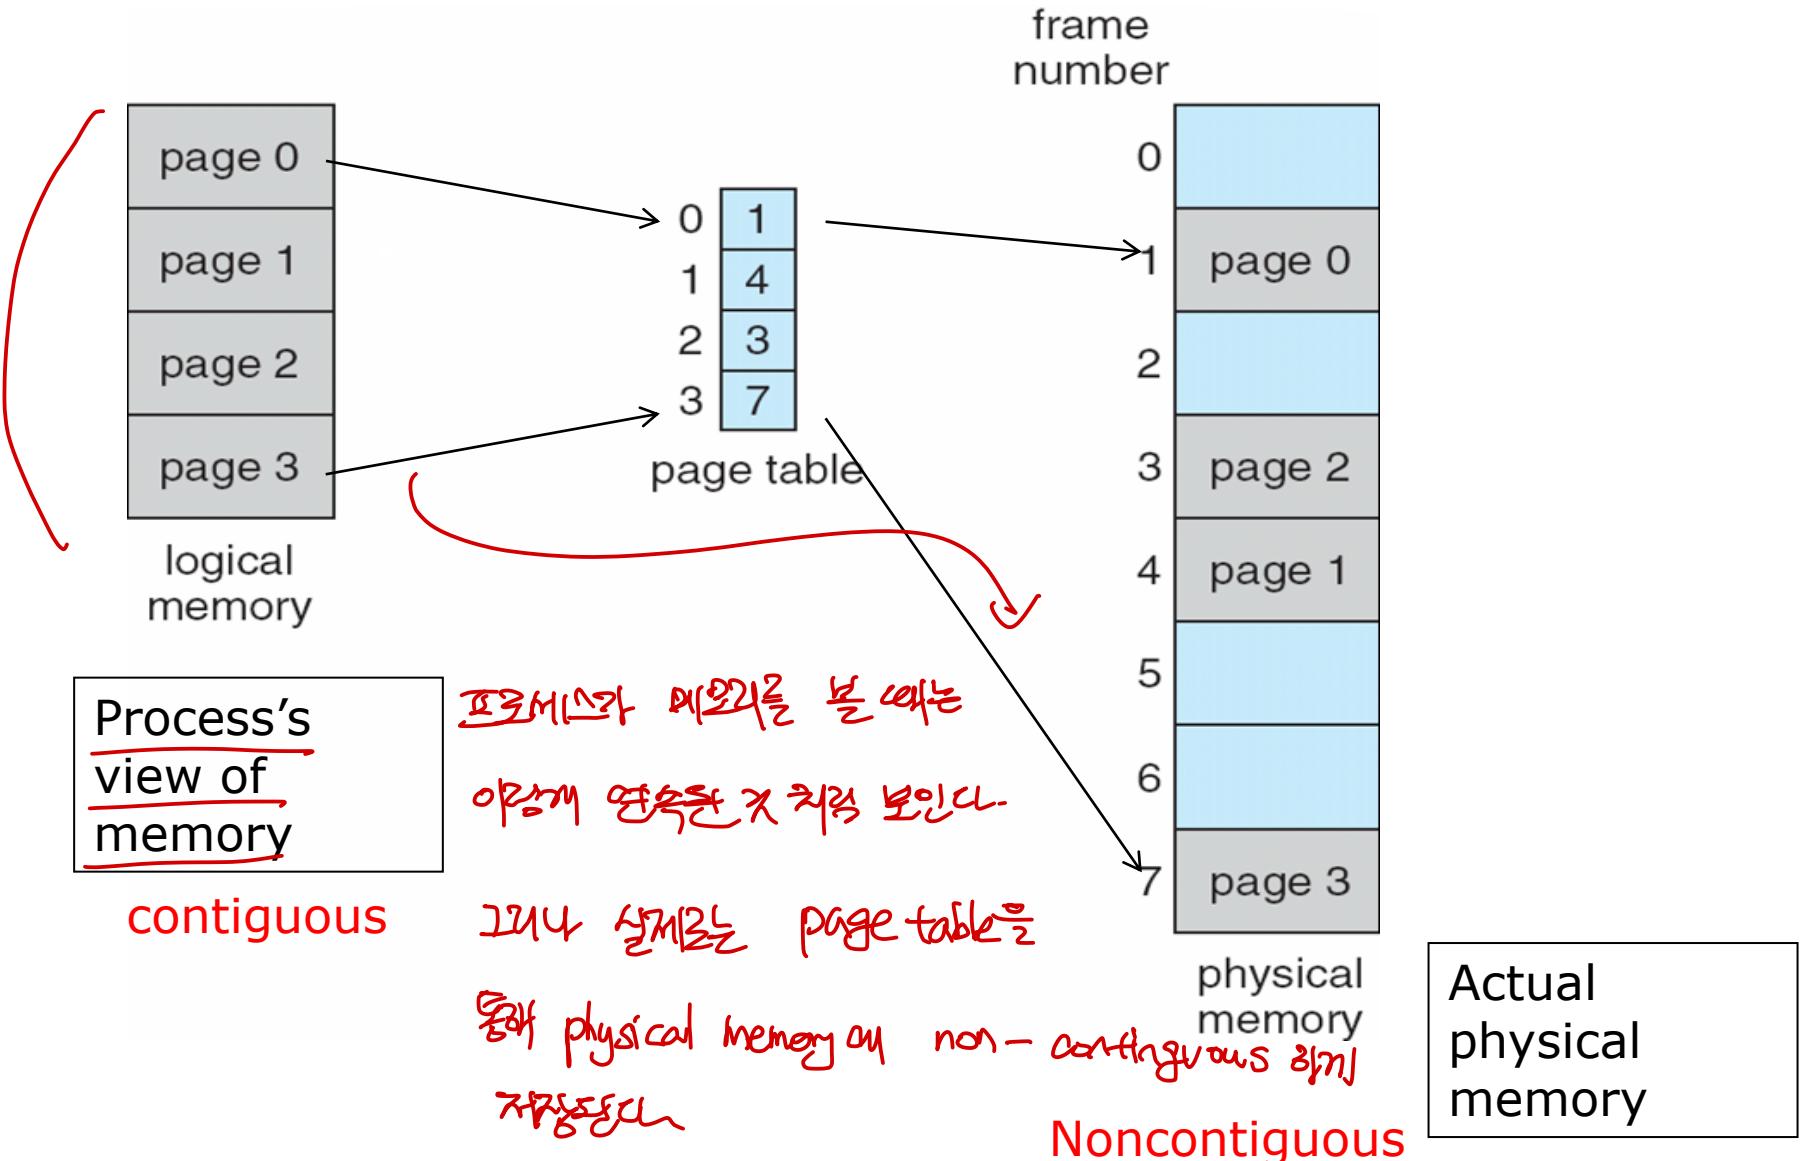
- **physical memory** → fixed-sized physical **frames**
- **logical memory** → fixed-sized virtual **pages**
- size of one frame = size of one page (e.g., 4kB) 대략 4KB정도

- Set up a per-process **page table** to translate logical to physical addresses

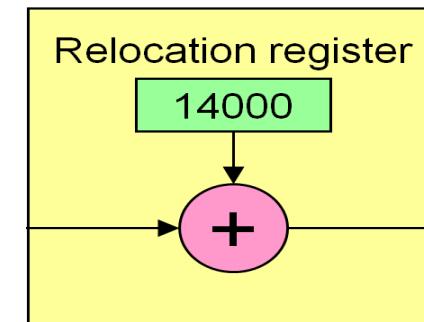
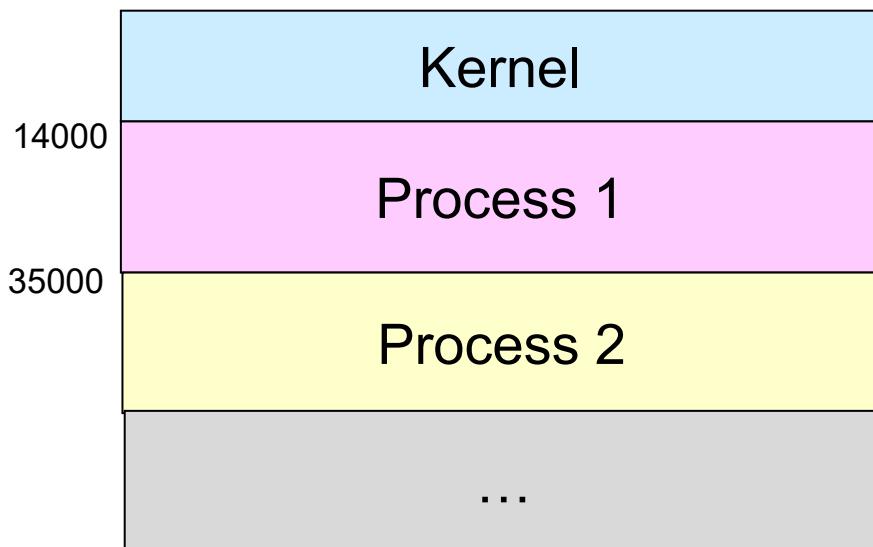
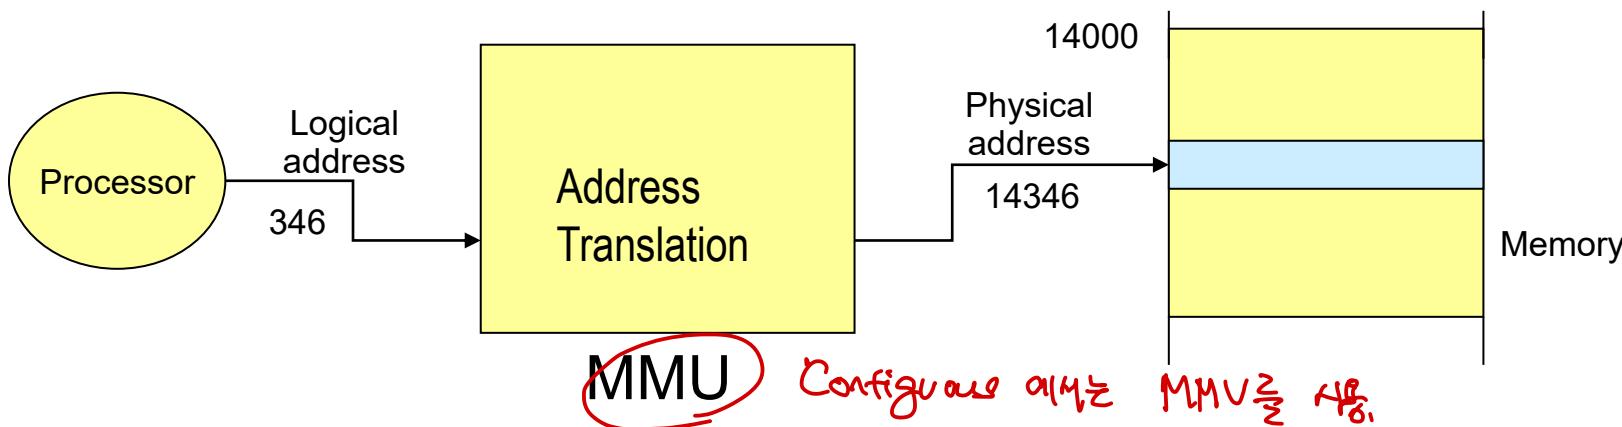
프로세스 마다 page table을 가지게 한다.



Paging Example

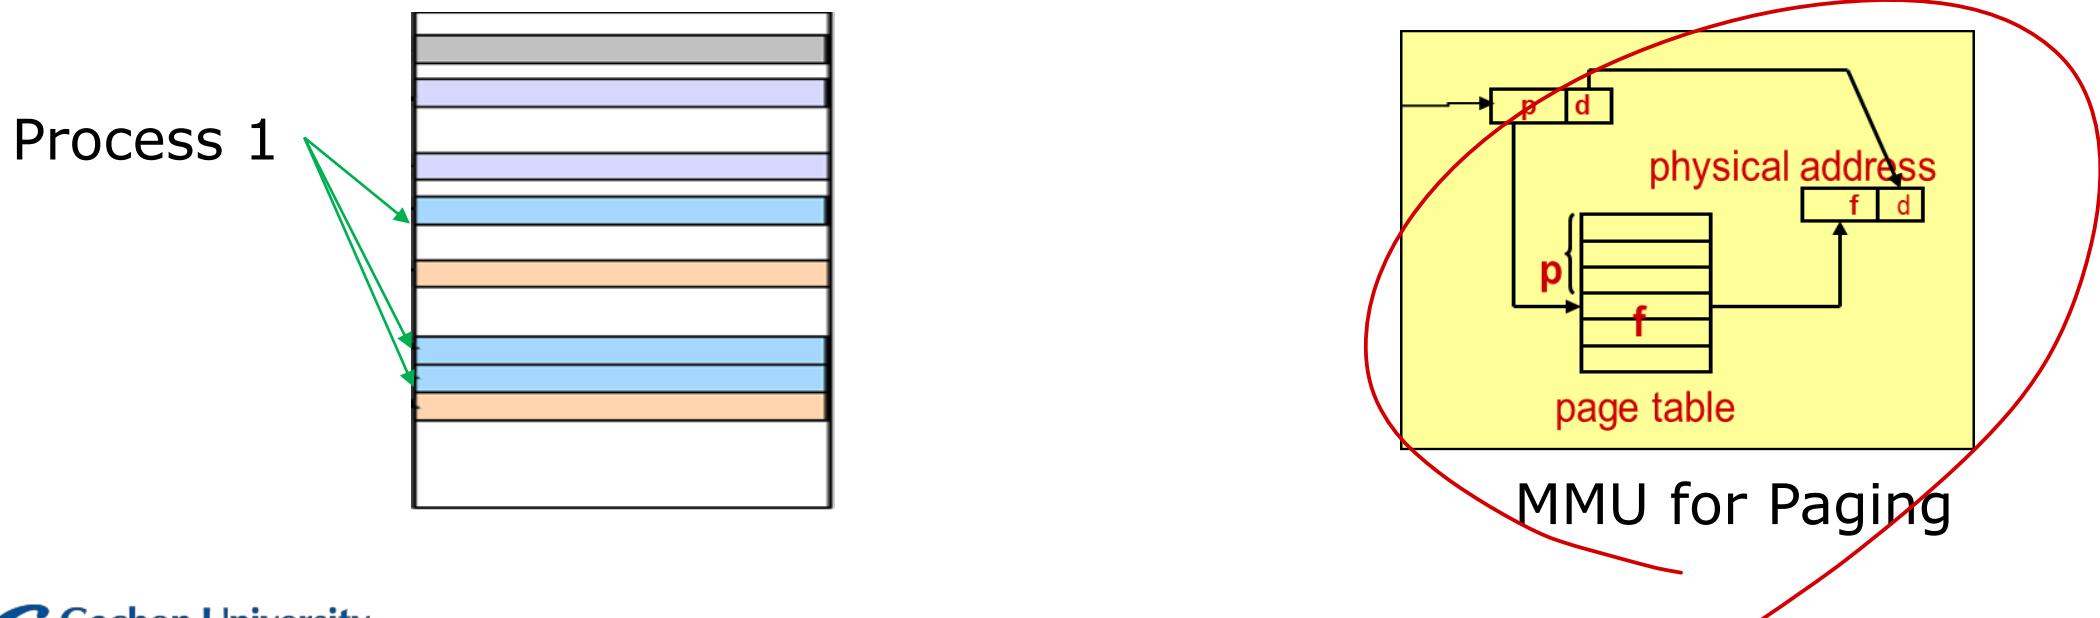
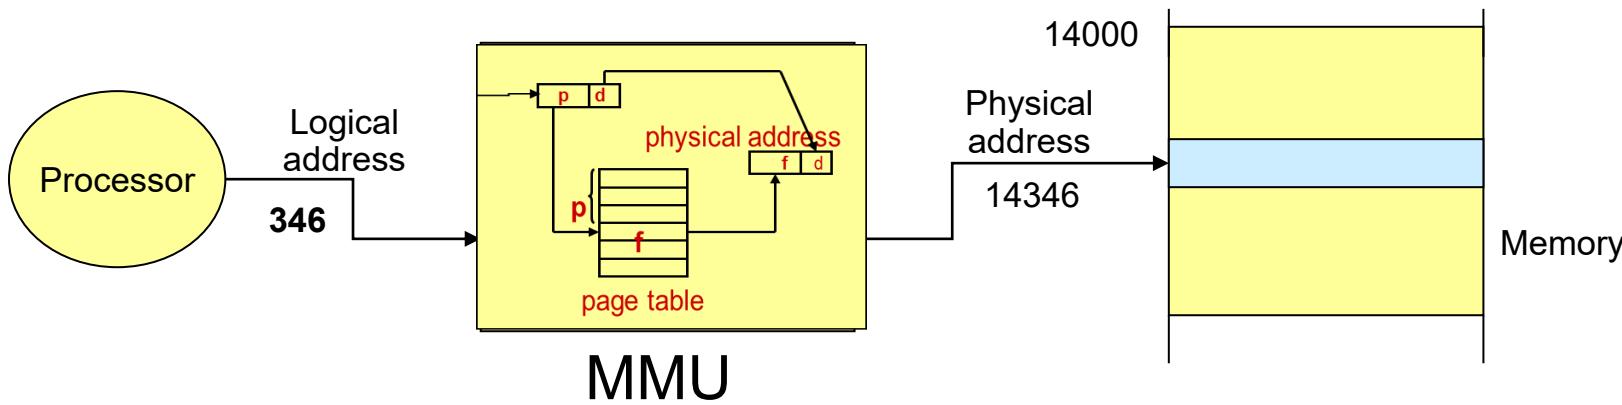


Contiguous vs. Paging



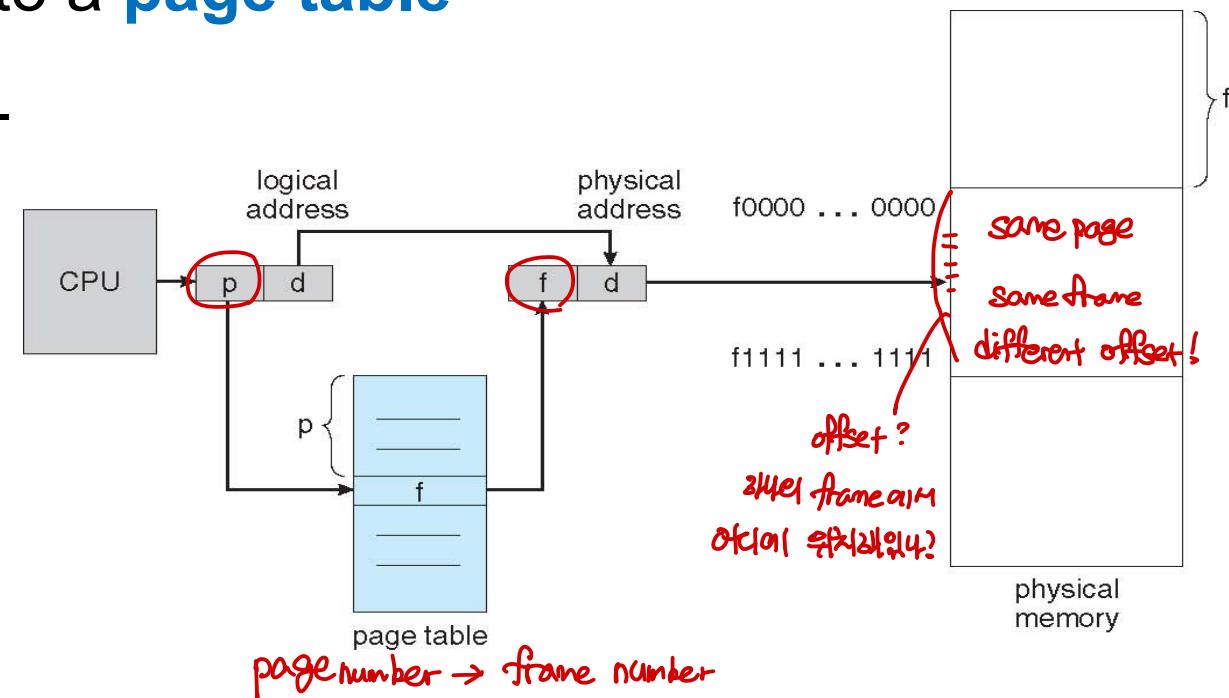
MMU for Contiguous Allocation

Contiguous vs. Paging



Page Number and Offset

- When a process is executed, its pages are loaded into any available memory frames
- Every address generated by CPU is divided into two parts: **page number (p)** and **page offset (d)**
 - page number (p)**: index into a **page table**
page table: contains page-frame mapping info.
 - page offset (d)**: page number and page offset combines to define the physical memory address



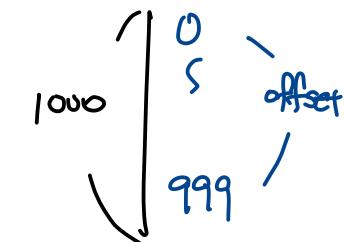
Paging Examples (From 2017 Quiz)

- Assuming a 1000B page size, what are the page numbers and offsets for the following address references (provided as decimal numbers)?

- (1) 42095: page number

42

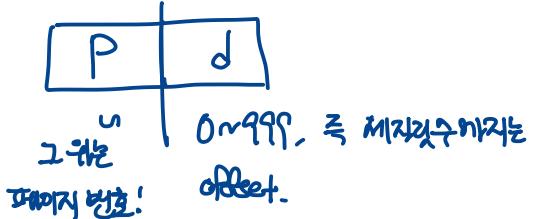
offset
95



- (2) 215201: page number

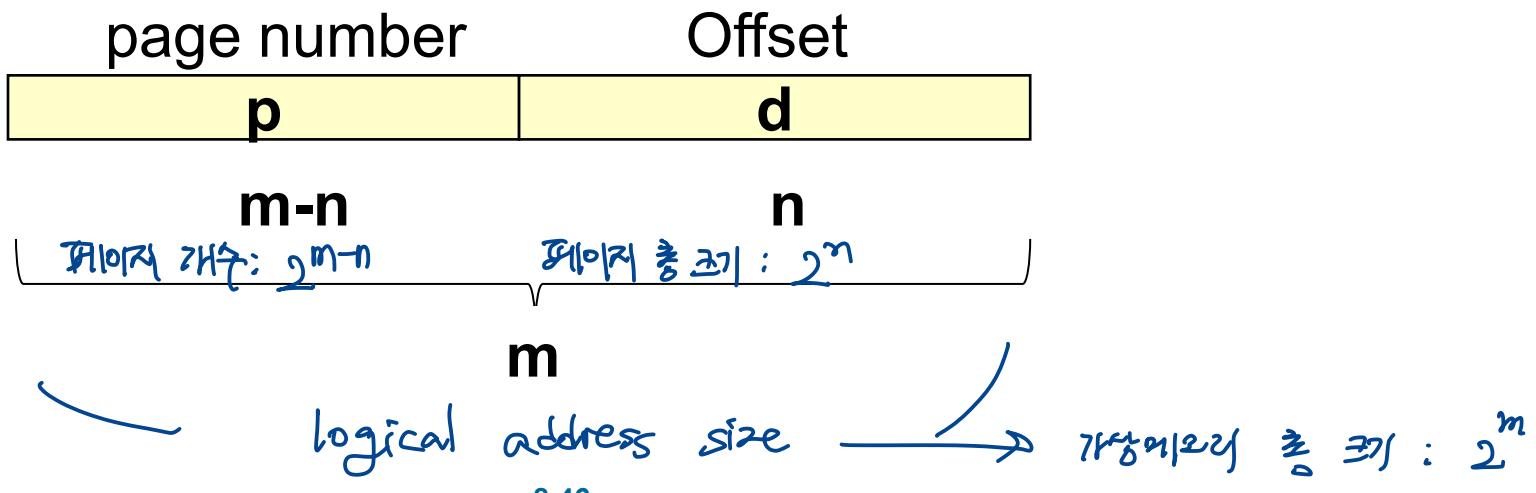
215

offset
201



Page Size

- Fixed sized chunk: **page size** (=frame size) defined by hardware
 - vary between 512B ($=2^9$ B) and 1GB ($=2^{30}$ B) per page – depending on computer architecture
→ 우리는 4KB 정도로 사용할 것.
- If the size of **logical address space** is 2^m and **page size** is 2^n
 - **page number(p)**: high-order $m-n$ bits of a logical address: **index** into a *page table*
 - **page offset(d)**: low-order n bits of a logical address



Paging Examples (From 2017 Quiz)

- Consider a logical address space of 128 ($=2^7$) pages with a 4kB ($=2^{12}$ B) page size, mapped onto a physical memory of 64 ($=2^6$) frames.

- (a) How many bits are required in the logical address?
 - 128 pages 2^7
 - 4kB page: $2^{12} \rightarrow 12\text{bits}$

$$\begin{array}{c} P \quad d \\ 2^7 \quad 2^{12} \\ \swarrow \quad \searrow \\ 2^{19} \end{array}$$

- (b) How many bits are required in the physical address?

- 64 frames: 2^6
- 4k-B frame: $2^{12} \rightarrow 12\text{bits}$

Paging Example

page#	0	1	2	3
0	a			
1	b			
2	c			
3	d			
4	e			
5	f			
6	g			
7	h			
8	i			
9	j			
10	k			
11	l			
12	m			
13	n			
14	o			
15	p			

$$\begin{array}{l} n = 2 \\ m = 4 \end{array}$$

0	5
1	6
2	1
3	2

page table

Frame#

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

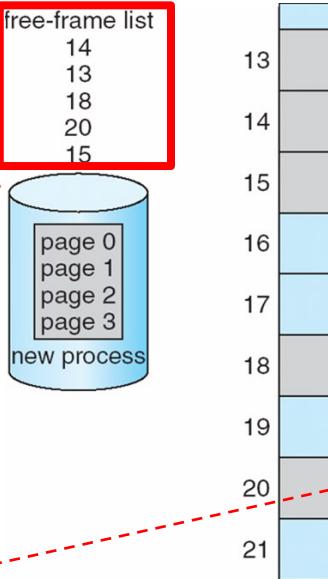
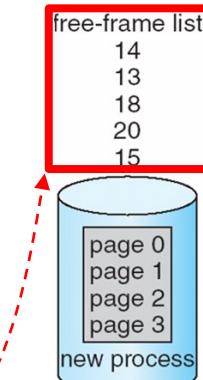
physical memory

- logical address space
 $= 2^4 = 16$
- page size is $2^2 = 4$

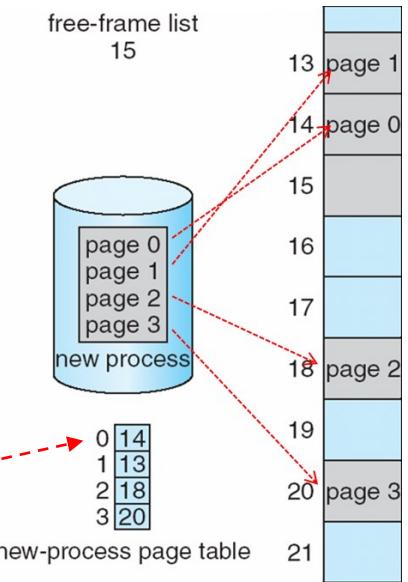
Free-frame management

- To execute a process whose size requires n pages, at least n free frames are required – How do we find the **free frames**?
 - The first page of the process is loaded into the first frame listed on **free-frame list**
 - and the allocated frame number is put into **page table**
 - The next page is loaded into another frame ...

비워 있는 frame의 list 를 가집고
그것을 업서 가집니다.



(a)
Before allocation

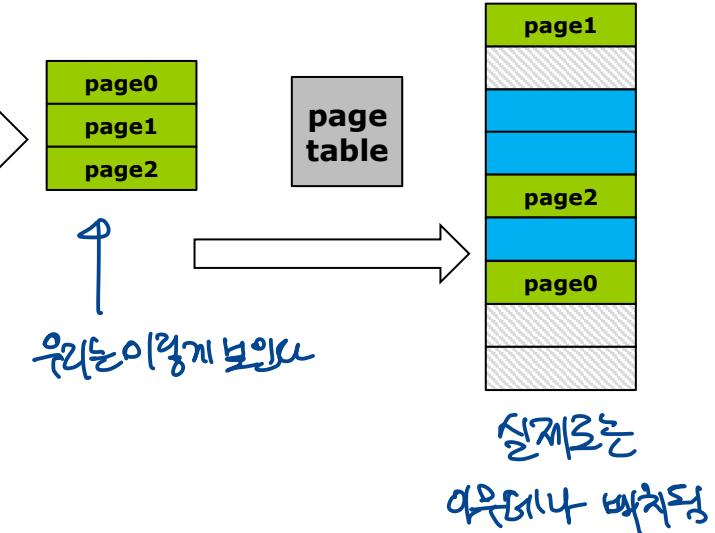


(b)
After allocation

A new process with 4 pages

Why Paging? (1)

- Clear separation between the programmer's view of memory (logical memory) and physical memory
 - Programmer: One contiguous single space
 - Actual physical memory: scattered throughout the physical memory
 - Needs **page table** for address translation



1. Provides **Transparency**

- Programmer only sees *logical address*, never sees actual *physical address* – *physical address* is transparent to programmer
- This mapping is *hidden* from programmer and controlled by OS

Why Paging? (2)

2. No external fragmentation

- Noncontiguous memory allocation - ANY free frame can be allocated to a process that needs it.
- However, we may have **internal fragmentation**
 - ▶ e.g.) if a page size is 2048 bytes, a process of 72766 bytes would need 35 pages plus 1086 bytes
 - ▶ In worst case, a process would need n pages plus 1 byte, and it would be allocated $n+1$ frames – internal fragmentation in last frame

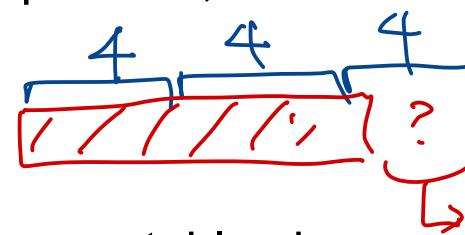
Internal Fragmentation

External Fragmentation

- Contiguous allocation: Total memory space exists to satisfy a request, but it is not contiguous

Internal Fragmentation

- Memory is allocated in fixed-sized **pages**
- Allocated **pages** may be larger than requested memory;
 - page size 4kB,
 - process needs 9kB – allocated 3 pages (12kB)
- this size difference is memory internal to a partition, but not being used

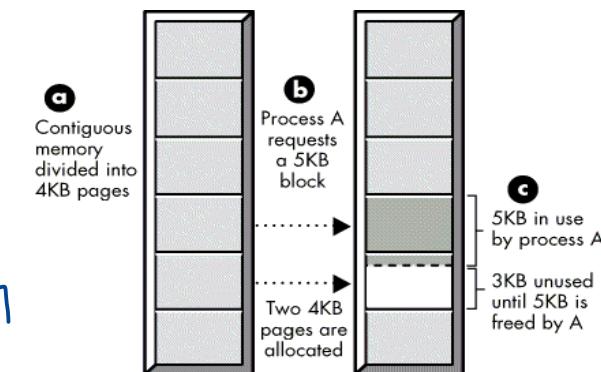
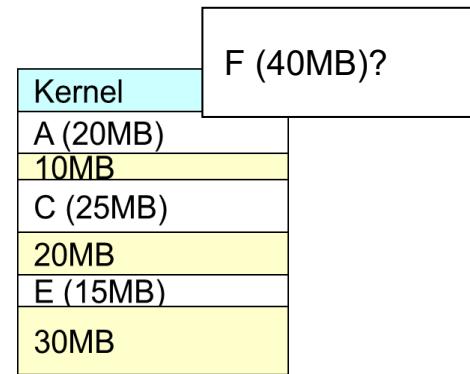


How do we determine the Page Size?

- Consider **tradeoff**: Internal fragmentation, page table size, disk I/O efficiency
- Practically 4 – 8KB in size

Page 크기가 적당 좋다

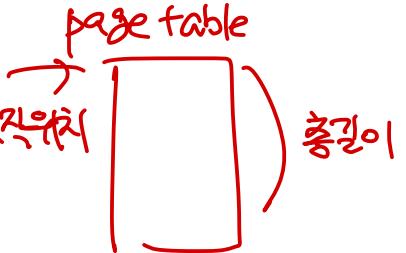
그러나 사용 줄일수는 없다!



Internal Fragmentation

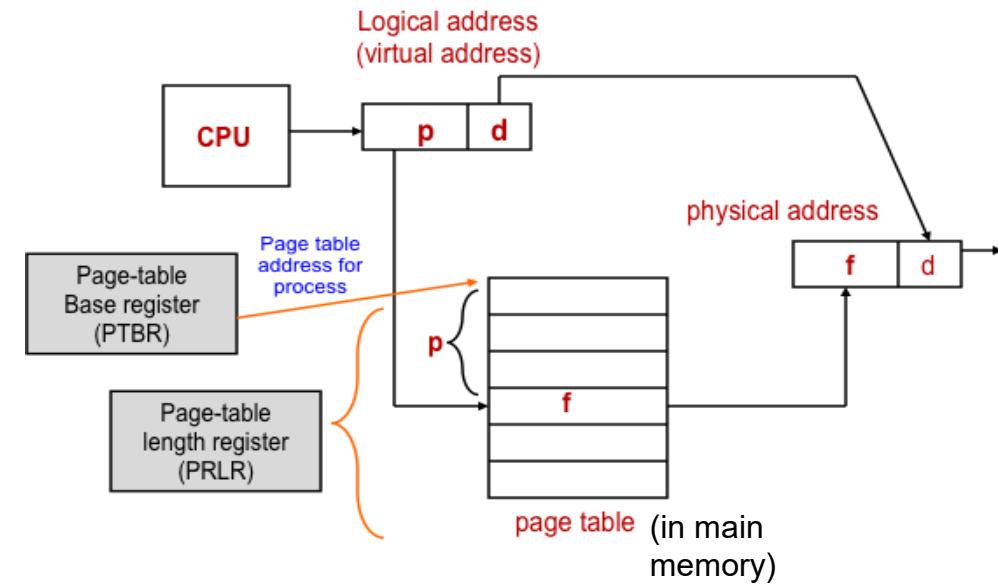
Implementation of Page Table

- There is a **page table** for each process!!
 - e.g., process P1's page table, P2's page table, ...
- **Option 1:** Keep in registers and make if fast
 - Implement page table as **registers**: translation very fast but expensive – need to keep it small (e.g., 256 entries)
 - But page table can be very large! (e.g., 1 million entries)
- **Option 2:** Keep in memory with register help
 - **Page table** is kept in **main memory**
 - Use **CPU registers** to store location and size of page table – hardware support



Implementation of Page Table

- **Page table** is kept in **main memory** with (register) **hardware support**:
 - **Page-table base register (PTBR)**
 - ▶ points to the page table for its process
 - **Page-table length register (PTLR)**
 - ▶ size of the page table (for the process)



- When context switch happens,
 - Need to change the page table (why? read the first line above)
 - ▶ changing page tables requires changing the **PTBR** and **PTLR**
 - ▶ so where should **PTBR**, **PTLR** values be stored when context switching happens?

process state
process number
program counter
registers
memory limits
list of open files
...

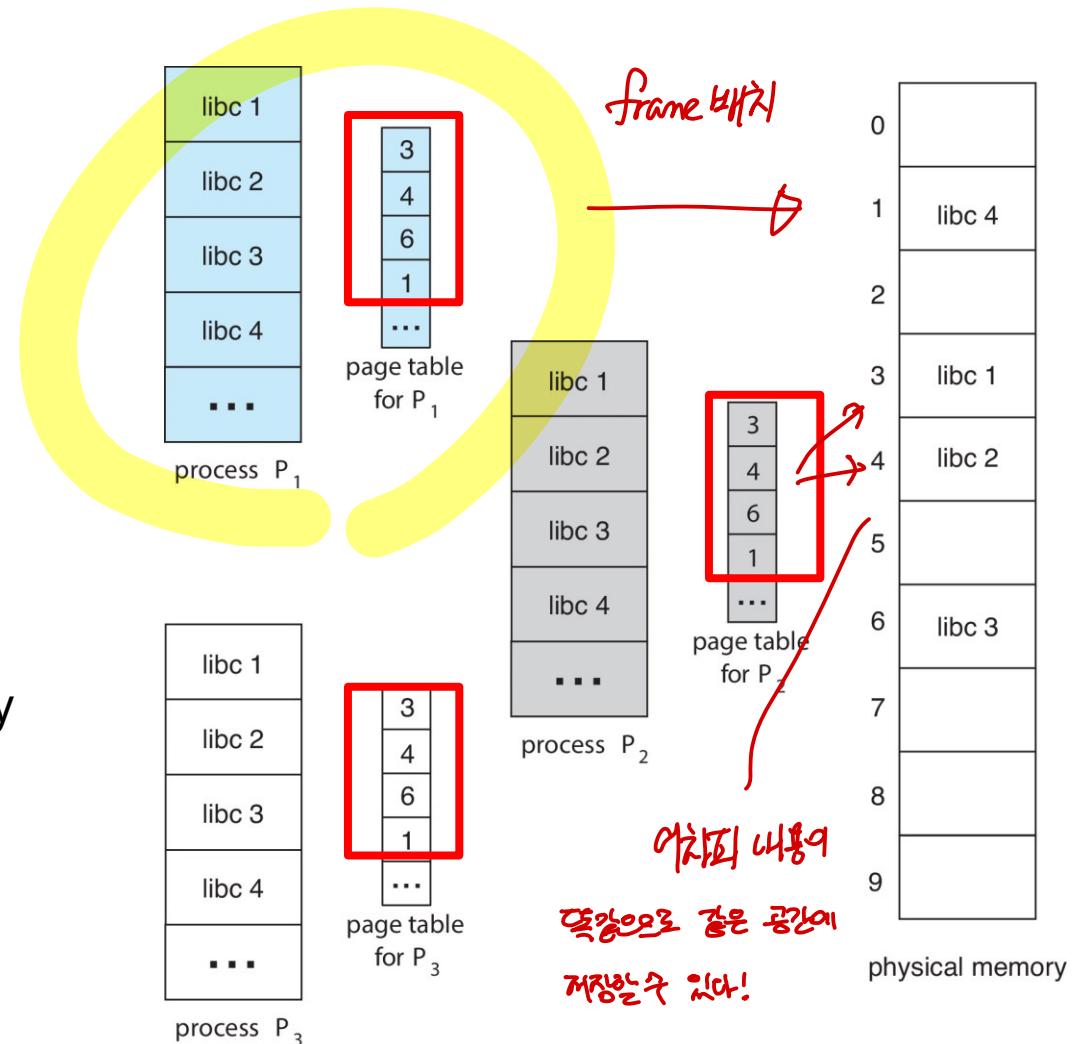
Shared Pages

- **Shared code**
 - One copy of **read-only** code shared among processes (i.e., libraries, text editors, compilers, database systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed
- **Private code and data**
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space

Shared Page (Shared Memory)

- Shared code

- One copy of library (libc) shared among processes
 - ▶ Example: 40 users run libc code (2MB)
 - need total 80MB physical memory
 - ▶ Share the code!
 - need only 2MB physical memory



Why Paging? (3)

- Provides **Protection** to memory
 - Programmer can never access memory outside the page table.
 - Programmer can never access memory outside the page table length defined in **page-table length register (PTLR)**
- Provides **Shared memory**
 - Use shared pages

Page Table in Main Memory

- **Page table is kept in main memory:** Problems?

- Every data/instruction memory access *actually* requires two memory accesses

- ▶ 1. page table 2. data/instruction

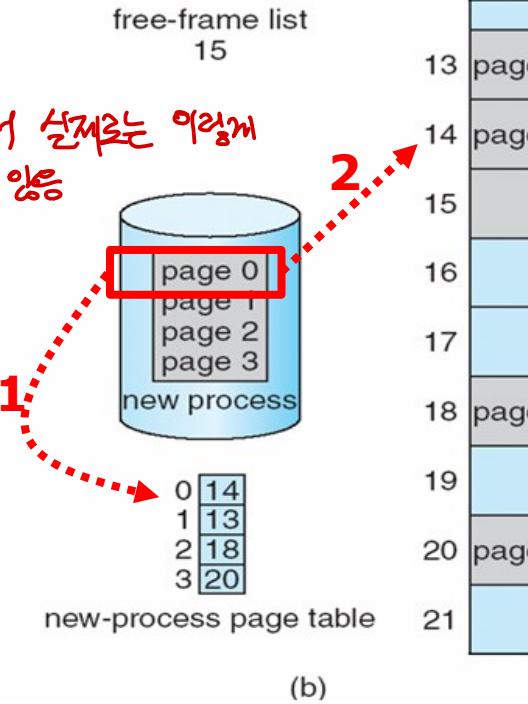
- Example
 - ▶ 어디 frame이
지정해야 하는지
알아오는 일
 - ▶ 지정해야 하는 위치에
데이터를 배치하는 일

- ▶ Process wants to read memory at logical address 010 (page 0, offset 10)

1. First read page table to look for physical frame number of page 0 (=14)
2. Read frame number 14

속도가 느림!

그래서 실제로는 이렇게
하지 않음



Page Table in Main Memory

- Isn't it slow to have to go to memory twice every time?
 - It takes about 100~400 CPU cycles to access DRAM
 - Yes, it would be... so, real MMUs don't
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
 - **TLB** is like a cache to the page table
 - ▶ e.g., cache memory is a cache to the main memory (RAM)
 - Store recently referenced page table entries in the TLB

Cache 랙은 예술!

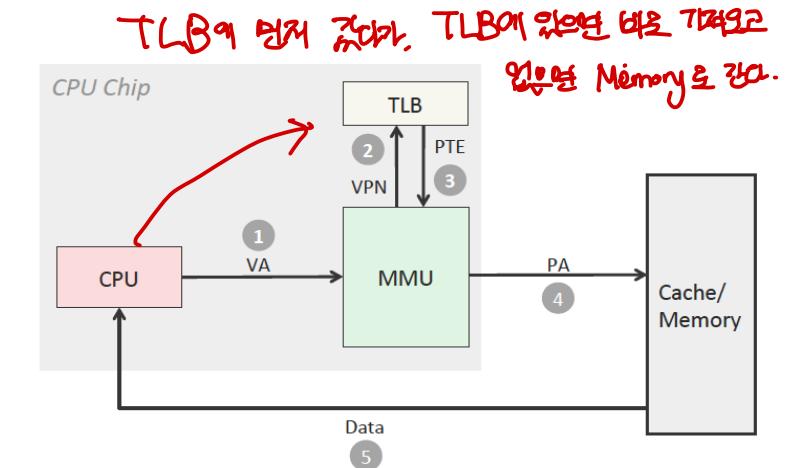


TLB (Translation Look-aside Buffer)

무조건 사용중인
TLB

- Small, dedicated, super-fast hardware cache in MMU
 - High speed memory technology, approaching register speeds
 - ▶ Implement using the dedicated registers or cache memories
 - ▶ Parallel search
 - Each TLB entry contains a page number and the corresponding page-table entry
 - ▶ Search key = logical address, result = frame number

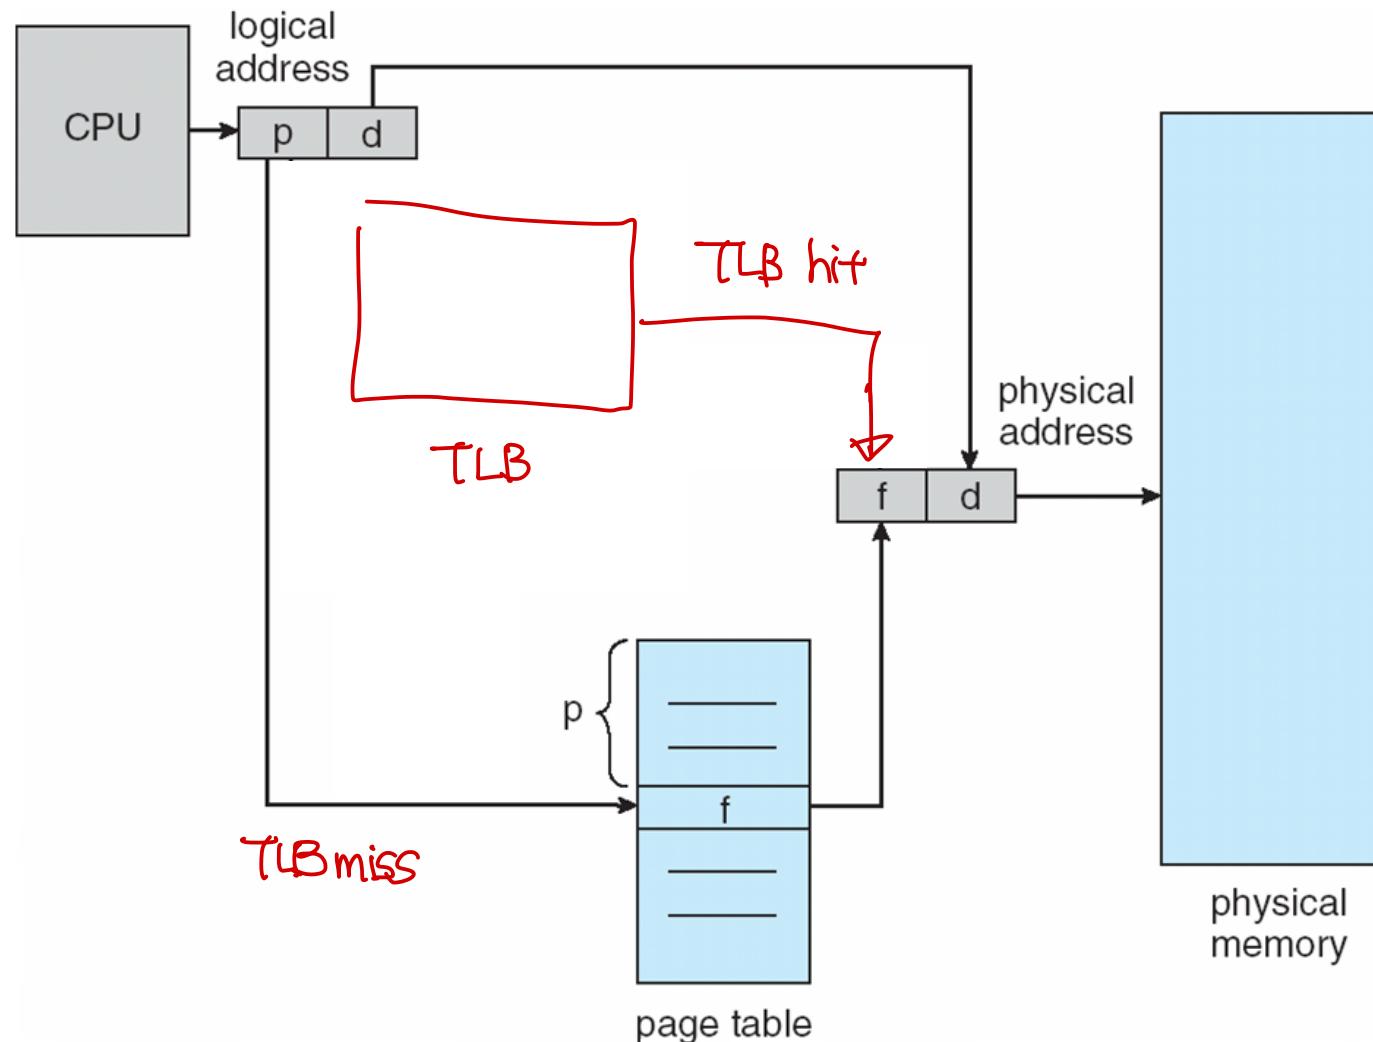
- Memory references that can be translated using TLB entries are much faster to resolve
 - high speed, but Expensive hardware



TLB

- **Recently-referenced** Page Table Entries (PTE's) are stored in the **TLB**
 - only small part of the page table is kept in TLB, typically 32 ~ 1,024 entries
- **TLB hit:** memory reference is in TLB
- **TLB miss:** memory reference is not in TLB
 - A TLB miss incurs an additional memory access to the page table
 - In this case, the page number/frame number is *added* to TLB for next page reference

TLB



Effective Access Time (EAT)

- Hit ratio

- percentage of times that a page number is found in TLB
- e.g., 80% hit ratio: find the page in TLB 80% of the time

- Assume

- memory access time = 100 ns, TLB access time = 1ns
- Hit ratio = α

만약 TLB를 쓰지 않을 경우 고정 200ns 발생

- Effective Access Time (EAT)

$$EAT = \alpha \times (1 + \underbrace{100}_{\substack{\text{TLB} \\ \text{data/instruction}}}) \text{ns} + (1 - \alpha) \times (1 + \underbrace{1}_{\substack{\text{TLB} \\ \text{page table}}} + \underbrace{100}_{\substack{\text{data/instruction}}}) \text{ns} = 201 - 100\alpha$$

if $\alpha = 0.8$ then EAT = 121ns, if $\alpha = 0.99$ then EAT = 102ns

TLB with cache

Memory access order

Logical (virtual) address

Check **TLB**

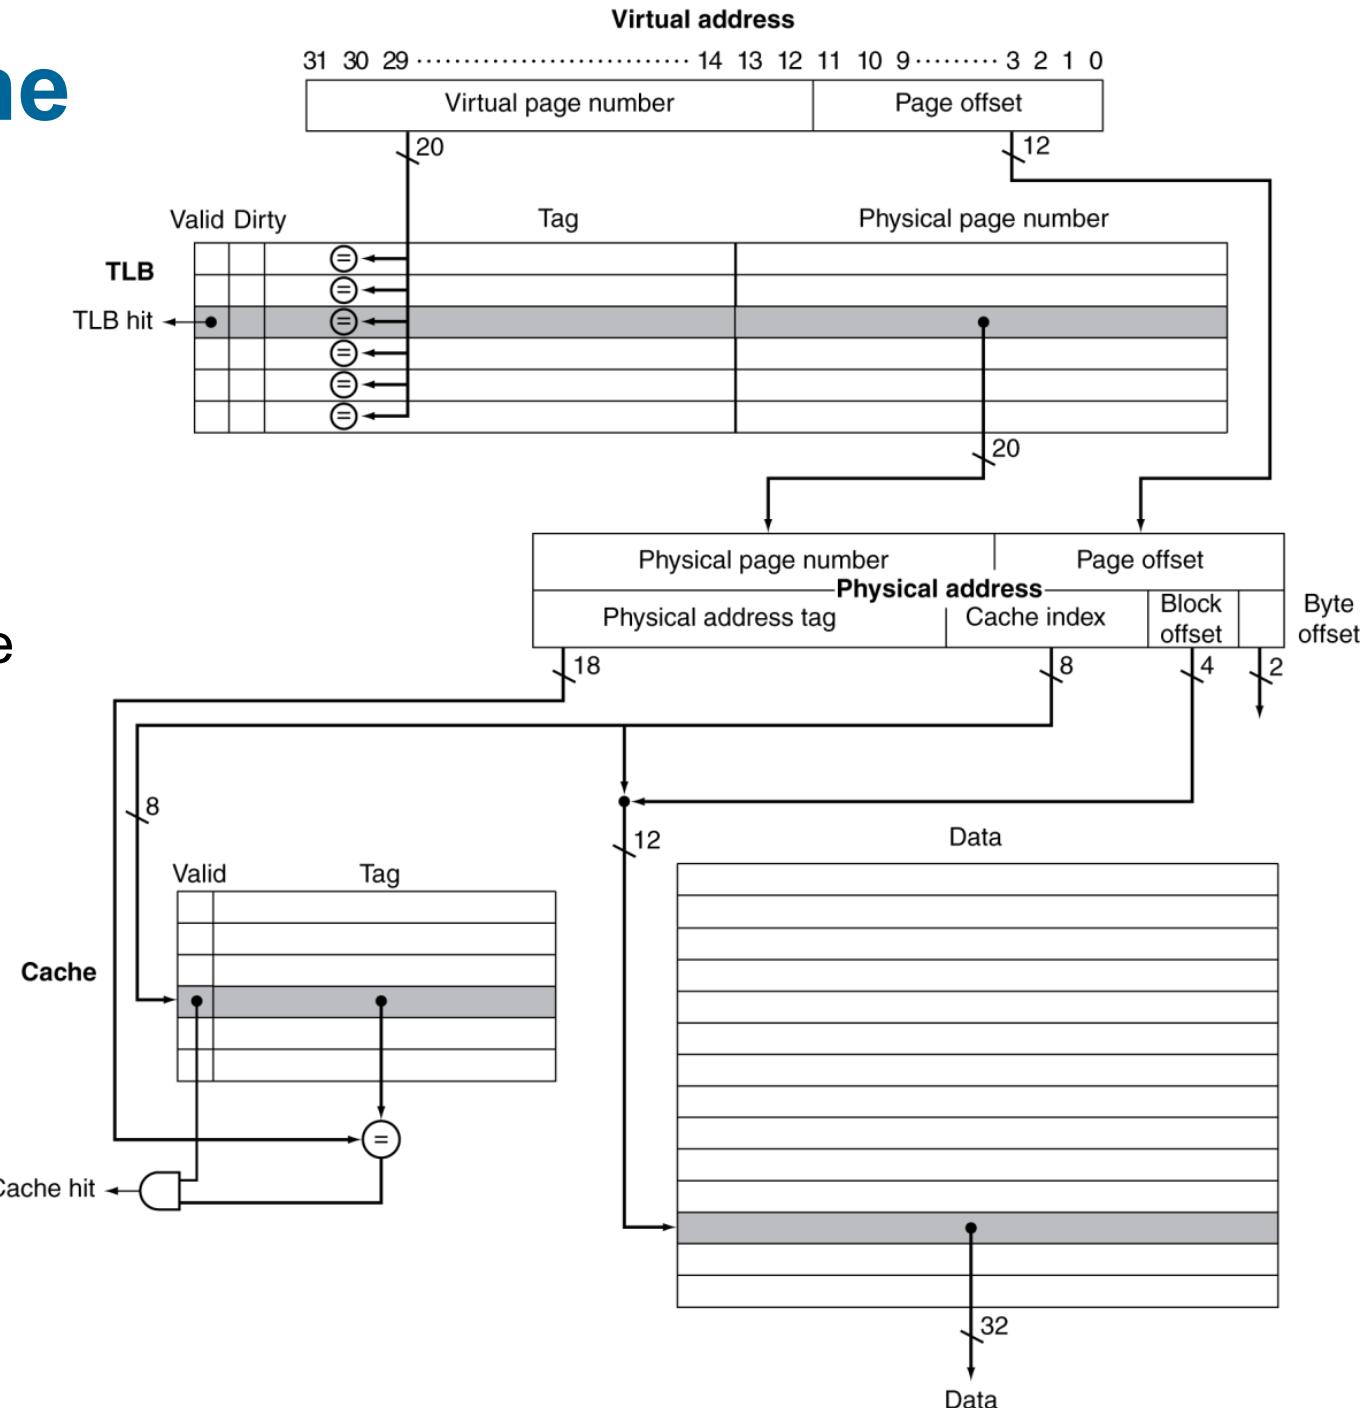
- ▶ **TLB hit**: get address from TLB
- ▶ **TLB miss**: check page table in memory

Acquire **physical address**

Check **cache**

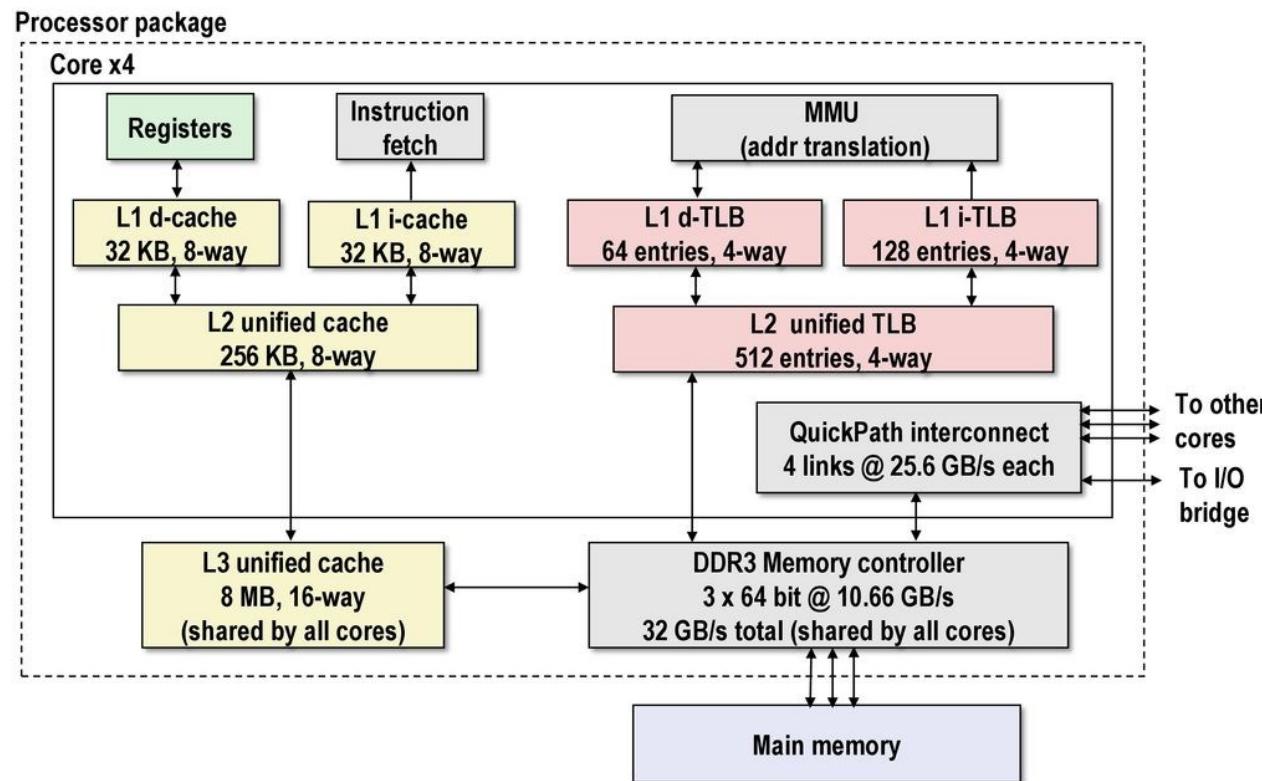
- ▶ **cache hit**: get data from cache
- ▶ **cache miss**: get data from main memory

Acquire data



Example: Intel Core i7 Memory System

Intel Core i7 Memory System



<https://slideplayer.com/slide/13277282/>

Chapter 8: Memory Management Strategies

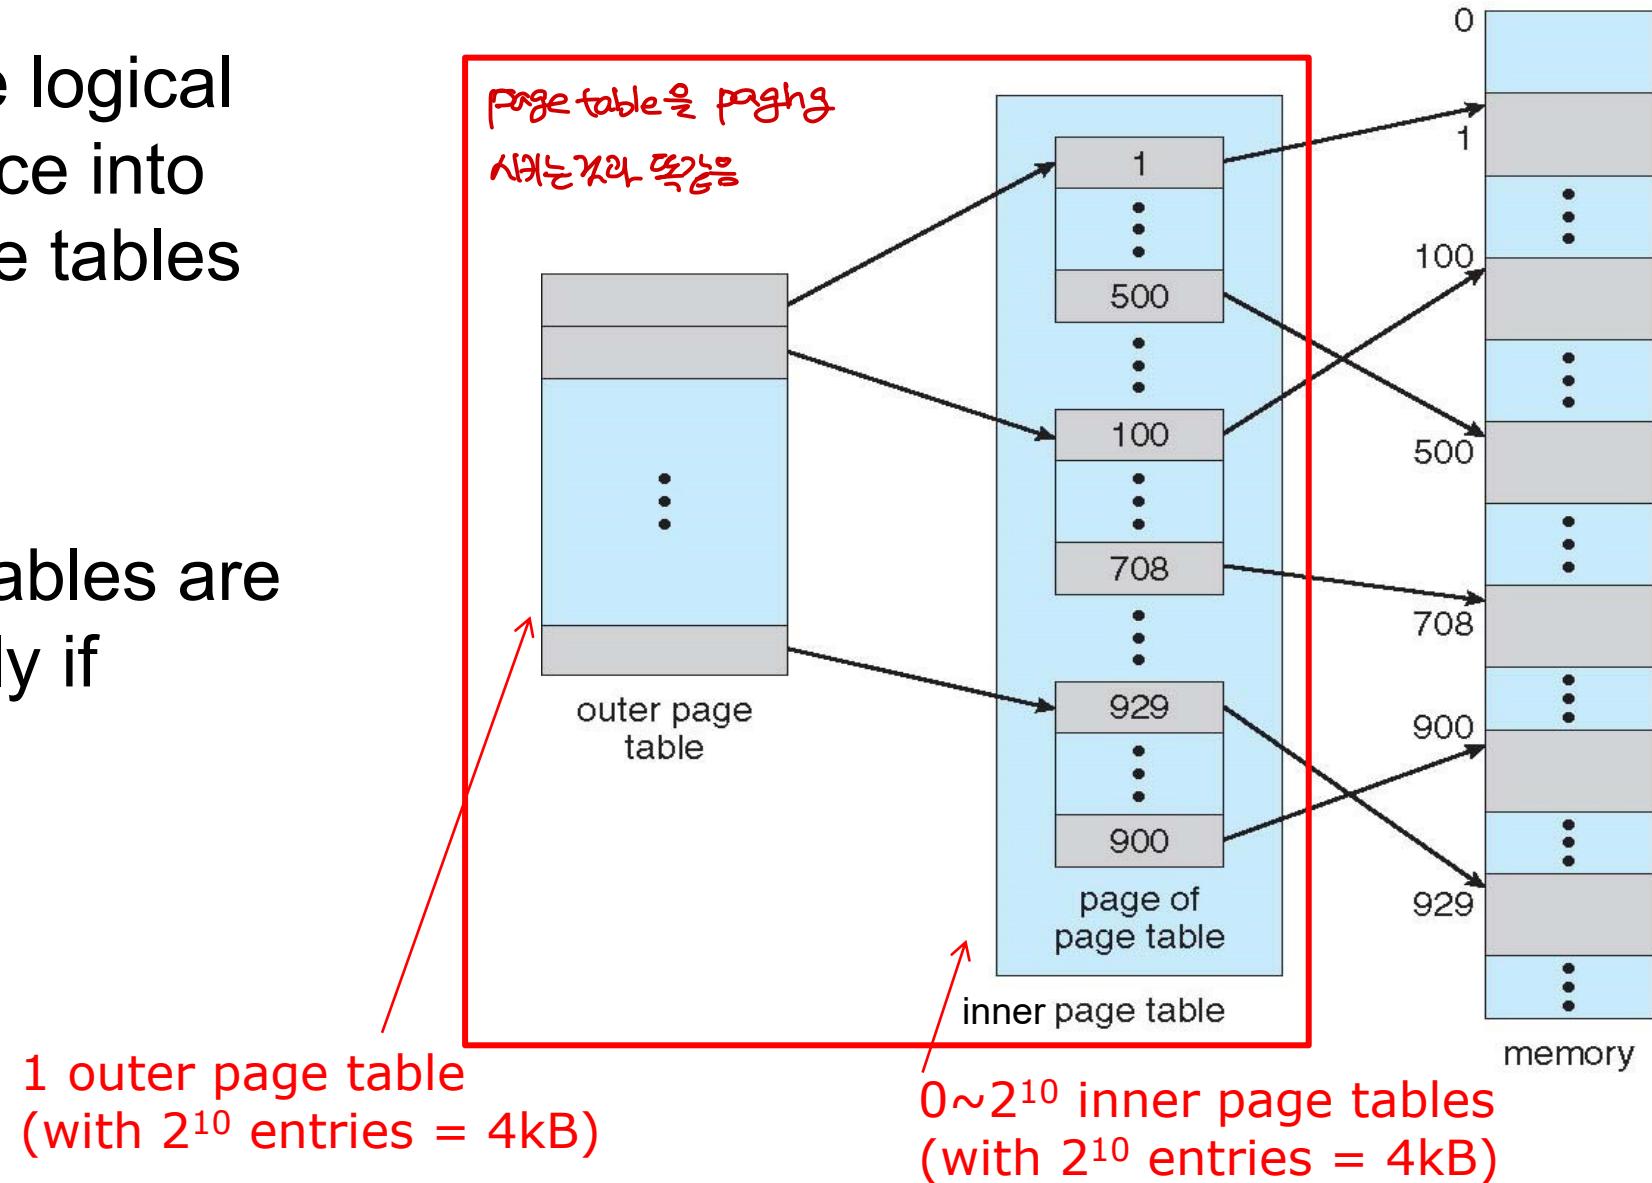
- Introduction
- Virtual Memory
- Contiguous Memory Allocation
- Paging
- **Structure of the Page Table**
 - Hierarchical Paging (Two-level page table)
 - Inverted Page Tables

How to Store Page Table

- Consider a system with 32-bit logical address space
 - 32-bit logical address space: 4GB ($=2^{32}B$)
 - If page size: 4KB ($=2^{12}B$)
 - Each page table needs up to $4GB/4KB = 2^{20} \approx 10^6 = 1$ million (2^{20}) 4B page entries
→ 4MB **contiguous physical memory space** per each page table
 - ▶ Note: each process needs its own page table!
- Solution? **Hierarchical paging**
 - Divide the page tables into smaller pieces
 - Multi-level paging, in which the page table is also paged

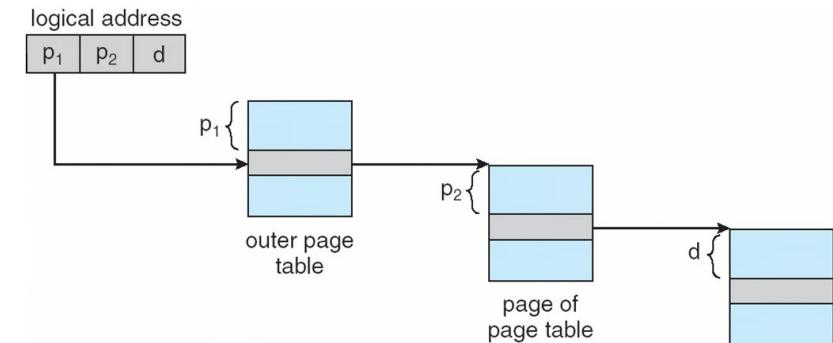
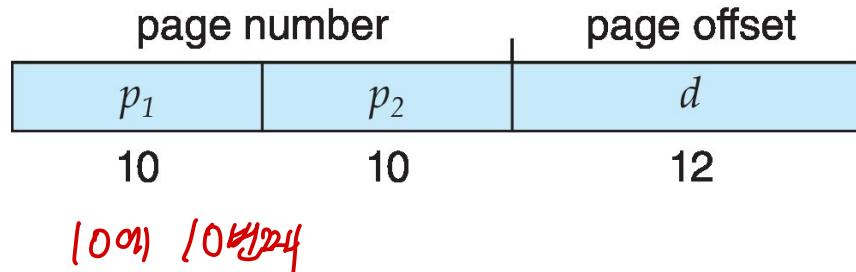
Two-Level Paging (32-bit address, 4kB page)

- Break up the logical address space into multiple page tables (logically)
- Inner page tables are allocated only if needed



Two-Level Paging (32-bit address, 4kB page)

- Since the page table is paged, the page number is further divided into:
 $\frac{4KB}{2^{12}}$
 - a 10-bit outer-page number, a 10-bit inner-page number, a 12-bit page offset.

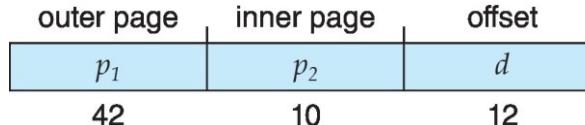


- where p_1 is an index into the **outer page table** ($2^{10} = 1,024$ entries: 4KB), and p_2 is the displacement within the page of the **inner page table** ($2^{10} = 1,024$ entries: 4KB).

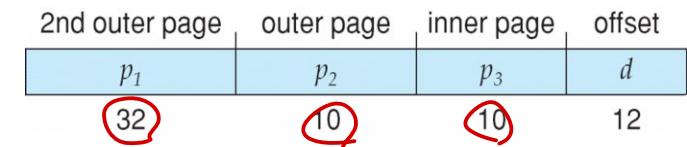


Three- Four-Level Paging (64-bit address)

- Even two-level paging scheme not sufficient
- If page size is 4 KB (2^{12})
 - Then page table has 2^{52} entries
 - If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - Address would look like



- Outer page table has 2^{42} entries or 2^{44} bytes ($2^{42} \text{ entries} \times 4B/\text{entry} = 2^{44} B$)
- One solution is to add a 2nd outer page table
 - But in this example the 2nd outer page table is still 2^{34} bytes in size
 - ▶ And possibly 4 memory access to get to one physical memory location!
 - ▶ Hierarchical paging is not appropriate for 64-bit architectures

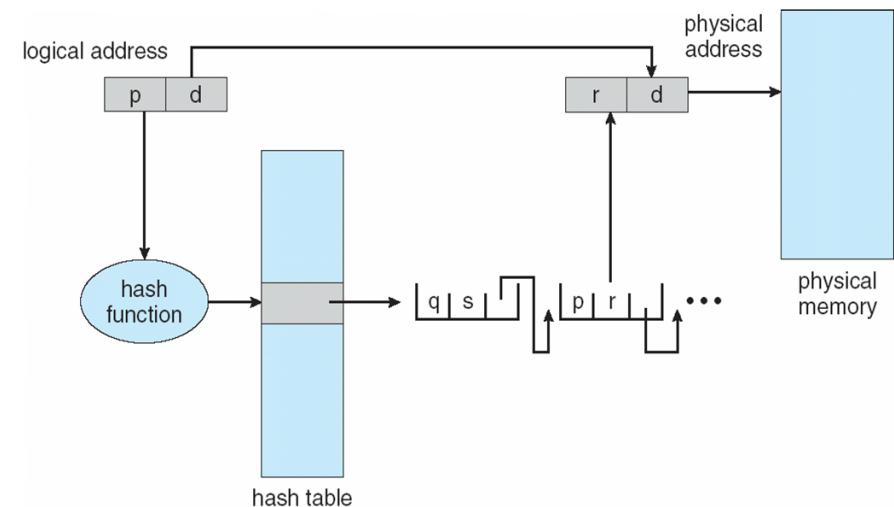


↳ 64bit에서는 ... 사용 불가능

Hashed Page Tables

- Common in address spaces larger than 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location (to handle collisions)
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted

검색 테이블



Inverted Page Table

- Usually, each process has a page table
 - Translate logical address (page) into physical address (frame)
 - Each page table can be very large! (as seen in last few slides)

역주

- Inverted page table**

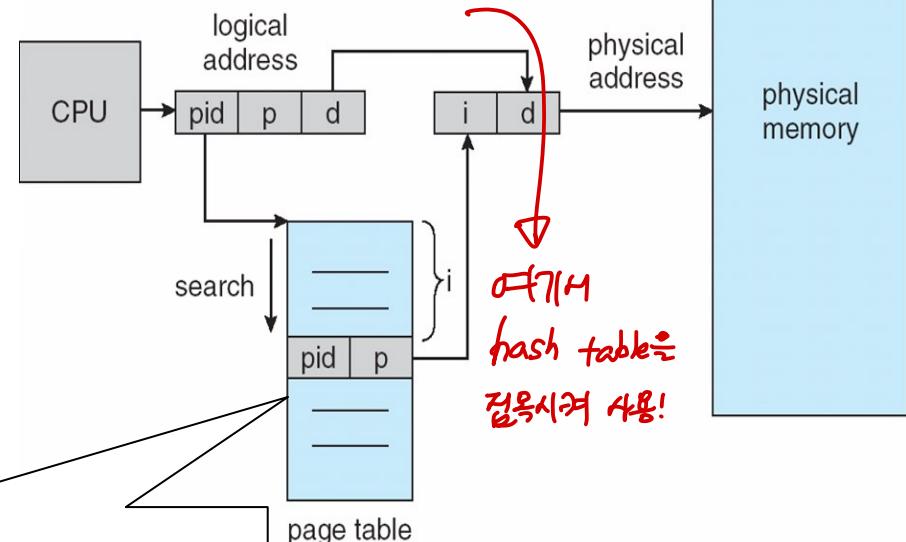
- Only one page table in OS!
- One entry for each real frame of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page

- Decreases memory needed to store each page table
 - Total table size = 1M entries (4MB size)

- Any disadvantages?

PT: page# → frame#
inverted PT: frame# → pid, page#

frame 번호를 넣고 그에 맞는 page 찾기
찾는 것. → linear search

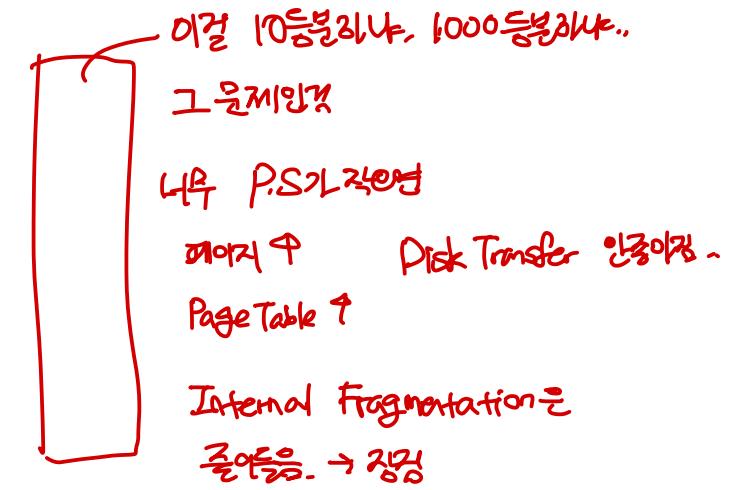


Inverted Page Table

- Advantage: Decreases amount of memory needed
- Disadvantages:
 - Increases time needed to search
 - ▶ Inverted page table is sorted by physical address, but the lookups occur on virtual addresses: If there are N entries, may need up to N searches!
 - ▶ Solution: Use **hash table** to limit the search page-table entries
 - Difficult to implement shared memory *shared frame → shared page or... Share X*
 - ▶ One mapping of a virtual address to the shared physical address
 - One physical page cannot have two (or more) shared virtual addresses

Other Considerations: Page Size (Ch. 10.9)

- Sometimes OS designers have a choice
 - Especially if running on custom-built CPU
- Smaller Page size gives:
 - More pages
 - Larger Page Table
 - Inefficient Disk Transfer
 - ▶ More I/O overhead
 - ▶ More page faults
 - Less Internal Fragmentation
- Trend
 - Larger page sizes



Obtaining the Page Size on Linux Systems

On a Linux system, the page size varies according to the architecture, and there are several ways of obtaining the page size. One approach is to use the system call `getpagesize()`. Another strategy is to enter the following command on the command line:

```
getconf PAGESIZE
```

Each of these techniques returns the page size as a number of bytes

Summary

- Various memory management schemes
 - Contiguous memory allocation
 - Noncontiguous memory allocation – paging
- H/W support for memory protection
 - Base and limit registers
 - TLB – associative mapping table
- Fragmentation problems
 - Internal and external fragmentation → 2(2)yo1R-
- Paging system
- More examples of Intel 32/64-bit and ARMv8 Paging architecture in your textbook



<https://www.istockphoto.com/kr/%EC%9D%BC%EB%9F%AC%EC%8A%A4%ED%8A%B8/qa>