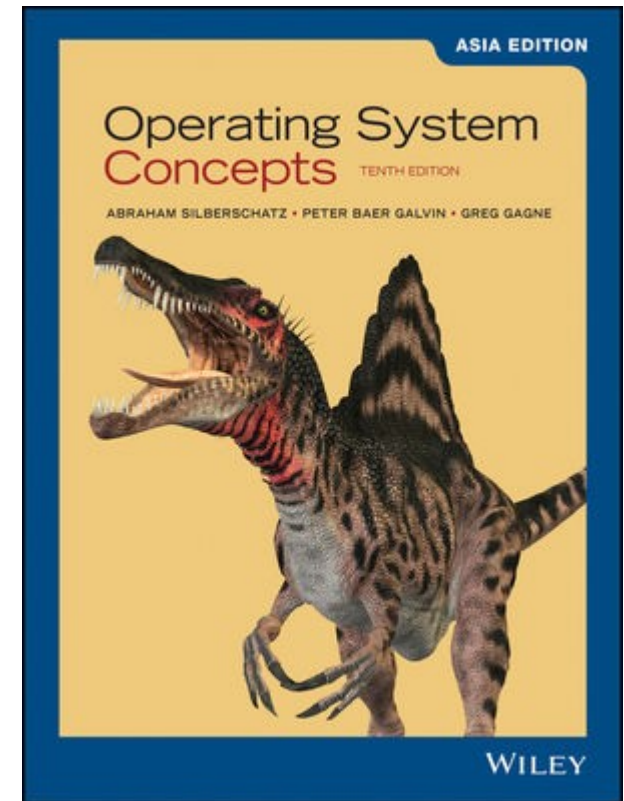


Chapter 10: Virtual-Memory Management

Dept. of Software, Gachon Univ.
Joon Yoo



Objectives

- Illustrate how pages are loaded into memory using demand paging.
- Explain concept of page faults
- Apply the FIFO, optimal, and LRU page-replacement algorithms.
- Describe concept of thrashing and working set algorithms

Chapter 10: Virtual-Memory Management

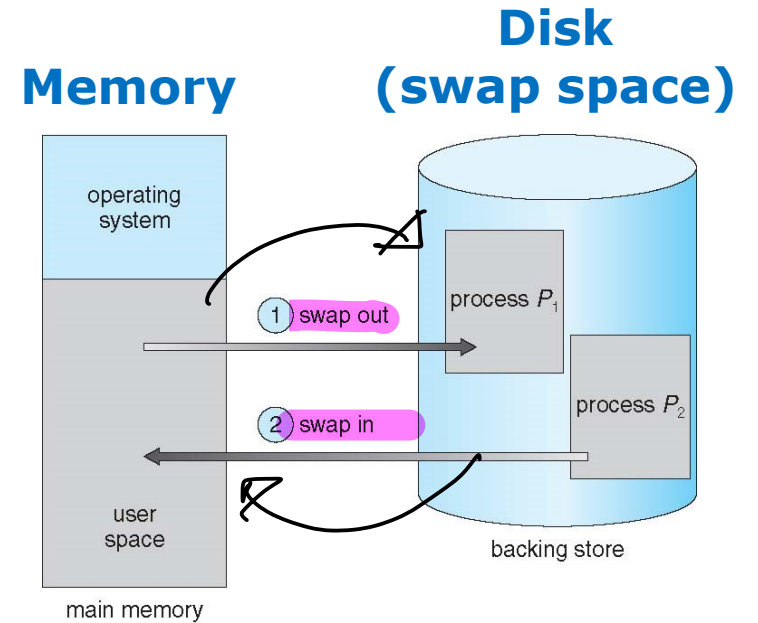
- **Background**
- Demand Paging
- Page Replacement
- Page Replacement Algorithms
- Thrashing
- Swapping on Mobile Systems

Chapter 9: Swapping

- Q: What if memory needed by process exceeds physical memory size?
 - In other words, can process size > physical memory?

- **Swapping**

- A whole process can temporarily **swap out** of memory to a backing store
- then **swap in** into memory for continued execution
 - ▶ swap device – usually HDD, SSD, or Flash



Chapter 9: Swapping

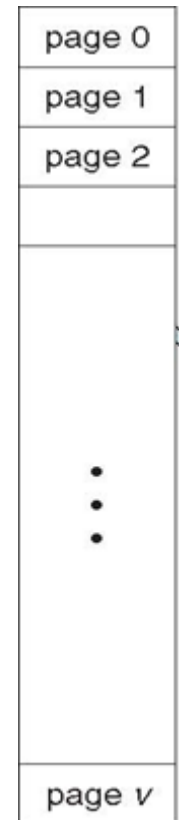
- Context switching time (=swap time) of swapping system is **very high**
 - e.g., user process 100MB, hard disk 50MB per second: then, $100\text{MB} / 50\text{MB per second} = 2 \text{ seconds}$. Swap-in + swap-out = $2 \text{ seconds} \times 2 = 4 \text{ seconds}$
- **Swapping takes too much time!** Modern OSes try to avoid swapping – only enable if **available physical memory is small**
 - ↳ 물리적인 메모리가 작을 경우만 허용
- Can we do better?
 - Swap only **portions of process** (rather than entire process) to decrease swap time: **Demand Paging**
 - 프로세스 전체가 아니라 "일부" 만 swapping 하는 것.

Chapter 10: Virtual-Memory Management

- Background
- **Demand Paging**
 - Basic Concepts
 - Page Fault
 - Performance of Demand Paging
 - Examples
- Page Replacement
- Page Replacement Algorithms
- Thrashing
- Swapping on Mobile Systems

Demand Paging Concept

- We can view a process as
 - a sequence of pages (rather than large contiguous address spaces)
- How can an executable program be loaded from disk into memory?
 - **Option 1**
 - ▶ Load the entire program into memory
 - **Option 2**
 - ▶ Load pages only as they needed – Demand Paging



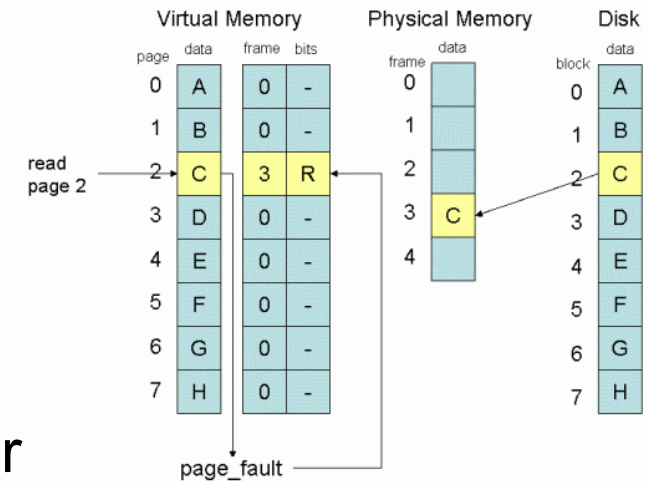
사용하는 page만 메모리에 올려놓자!

Demand Paging

- Motivation: Code needs to be in memory to execute, but entire program rarely used
 - Error code, unusual routines, large data structures (arrays)
 - e.g., Total 500MB program. Partially load only 100MB and start program. The rest can be loaded on-demand!

- Why? (Advantages)

- Less I/O needed from disk to load memory
 - ▶ Faster response to load programs
- Less memory needed
 - ▶ Larger programs can be partially loaded
- Better multi-program – more users for server



Valid-Invalid Bit

- With each page table entry a **valid-invalid bit** is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
 - or it can be represented as **1 / 0** for valid/invalid.
- Initially valid-invalid bit is set to **i** (or **0**) on all entries
- Example of a page table snapshot:
- During address translation, if valid-invalid bit in page table entry is **i** (or **0**) \Rightarrow **page fault**

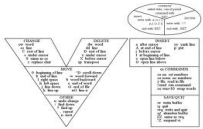
Frame #	valid-invalid bit
	v
	i
	v
	v
	i
....	
	v
	i

page table



처음에는 전부 다 invalid. 메모리에 올라와야 그때 valid로 변경됨.

Page Table When Some Pages Are Not in Main Memory



0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

logical memory

frame	valid-invalid bit	
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

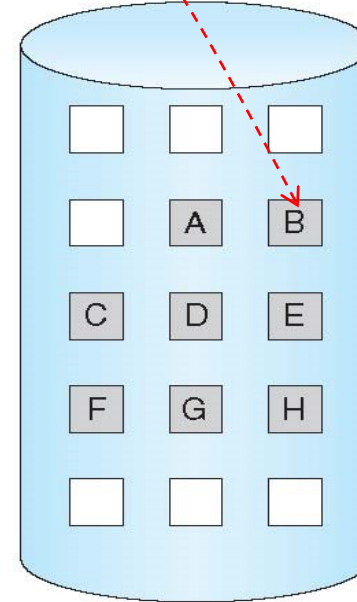
page table

0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15

physical memory



Disk (swap space)



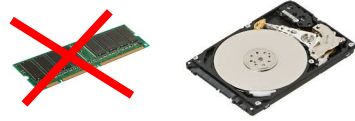
page fault

원하는 page가 memory에 없을 경우, disk에 그 page가 있다면 그것을 page-fault 라고 한다.



Page Fault

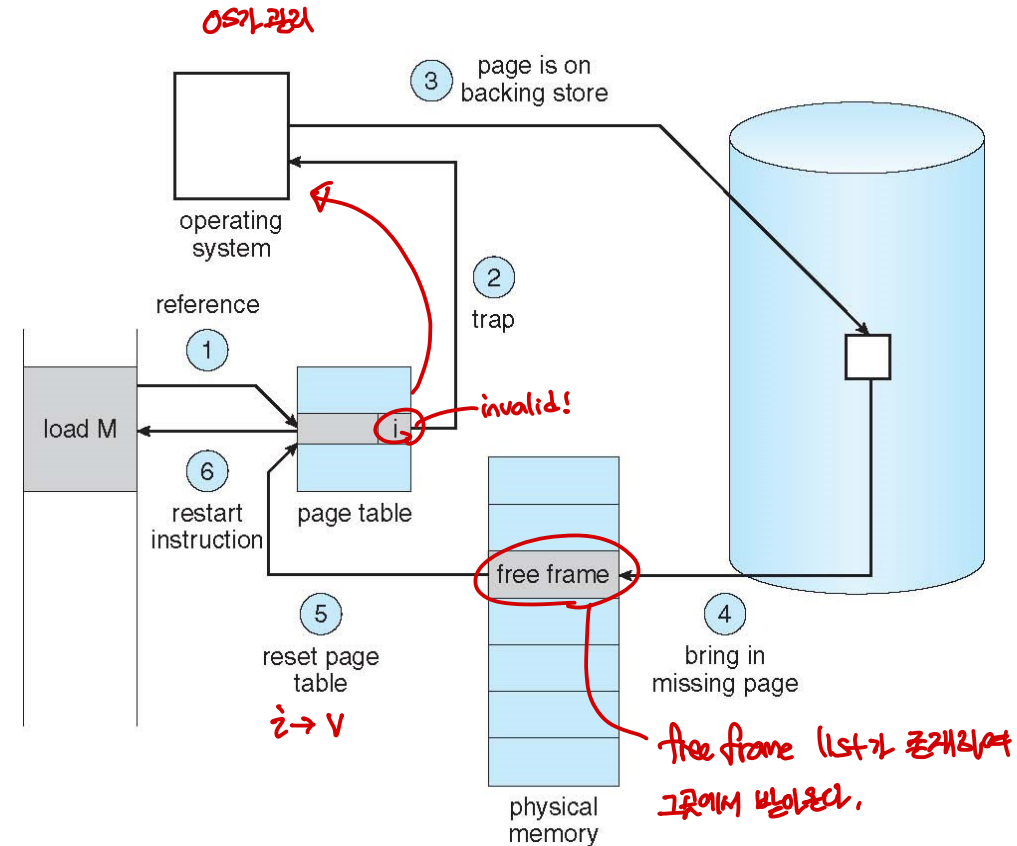
■ Page fault



- If there is a reference to an **invalid** page
- Reference to that page will **trap (software interrupt)** to operating system

■ OS handling **page fault**

1. Operating system looks at page table to decide
2. **PT: validation bit = i** \Rightarrow **not** in memory (trap)
3. Page is in swap-space - get an empty frame from **physical memory**
4. Bring in page from **disk** into **physical memory frame**
5. Reset **page table** to indicate **page** now in memory
Set **validation bit = v**
6. **Restart** the instruction that caused the page fault



Performance of Demand Paging

- Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- **Effective Access Time (EAT)** for *demand paging*

$$\text{EAT} = (1 - p) \times \text{memory access} + p \times \text{page fault time}$$

page fault time = page fault overhead

memory access time <<< page fault time

크게 고려하지 않음

오래 시간이 걸림.

→ disk I/O 발생






+ swap page out [if needed]

+ swap page in

+ restart overhead

page replacement

Demand Paging Example

- Memory access time = **200 ns**  \leftrightarrow 
- Average page-fault service time = **8 ms** (=8,000,000 ns)  \leftrightarrow  \leftrightarrow 
- $$\begin{aligned} \text{EAT} &= (1 - p) \times 200 \text{ ns} + p (8 \text{ ms}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \text{ (ns)} \\ &= 200 + p \times 7,999,800 \text{ (ns)} \end{aligned}$$
- If one access out of 1,000 causes a page fault ($p=0.001$), then
$$\text{EAT} = 8,200\text{ns} (= 200\text{ns} \times 41)$$

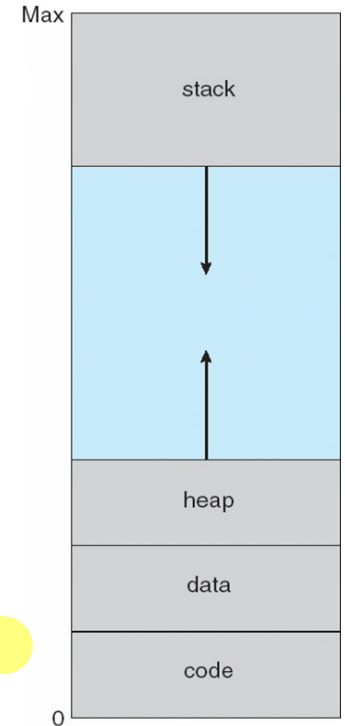
This is a slowdown by a factor of 41x!! (or 97.6% degradation)
- If we want performance degradation < 10 % (or $\text{EAT} < 220\text{ns}$)
 - $220 > \text{EAT} = 200 + 7,999,800 \times p$
 - $20 > 7,999,800 \times p$
 - $p < .0000025$
 - < one page fault in every 400,000 memory accesses
- It is important to keep the page-fault rate low in a demand-paging system!

page fault가 훨씬만 일어나도
엄청난 시간낭비가 발생한다.

10% 성능 향상을 위해 p의 최댓값이 ...
= 0.0000025 ... ㅋㅋ!

Demand paging usage: Heap/Stack

- Heap: grow upward in memory as it is used for dynamic memory allocation
- Stack: grow downward in memory through successive function calls
- Do not need to allocated physical frames for heap/stack at the beginning - will require actual physical frames only if the heap and stack grows
- Ideally implemented via demand paging
- Easy to extend memory space of process!



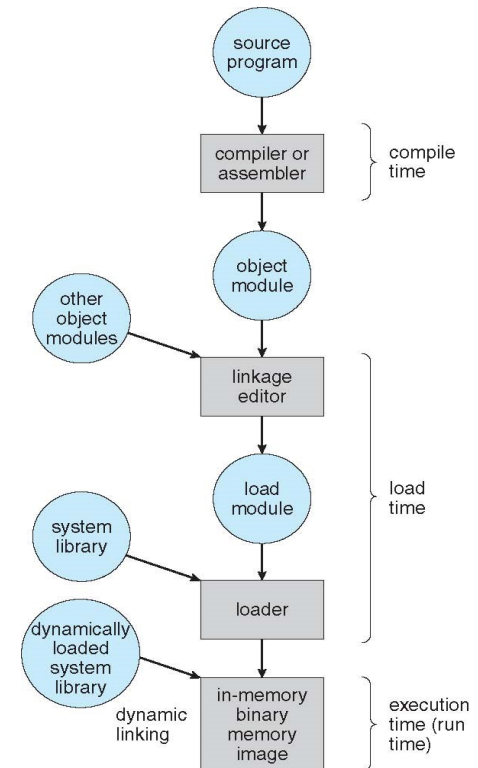
Demand paging usage: Dynamically linked libraries (DLL)

■ Dynamically linked libraries (DLL) – Chapter 9.1.5

DLL이 사용되는 library는 해당하는 코드가 실행될 때 그제서야 실행됨!

- Library linking is postponed until execution time
- A **stub** is included in binary image for each DLL reference

▶ **stub**: piece of small code that indicates how to locate appropriate routine.

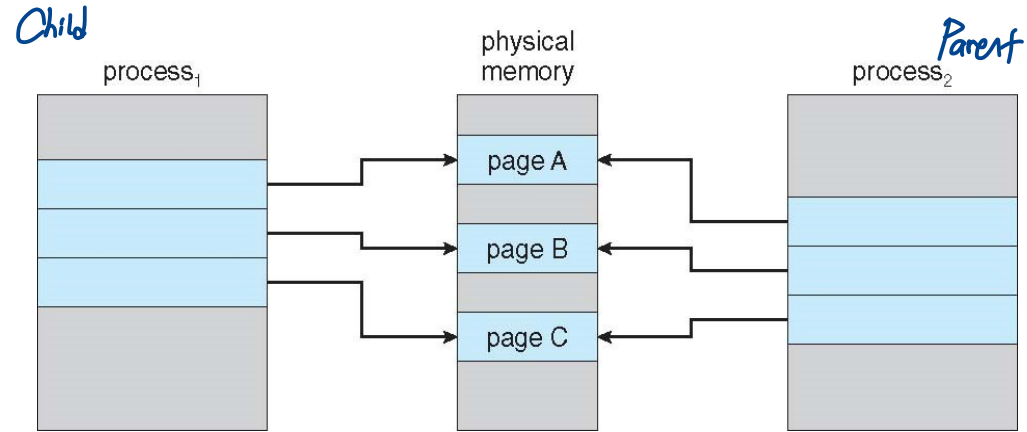


fork(): Copy-on-Write (Ch. 10.3)

- **Copy-on-Write (COW)** allows both parent and child processes to **initially share** the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- In general, free pages are allocated
 - when there are **copy-on-write** pages
 - or when child executes **exec()**

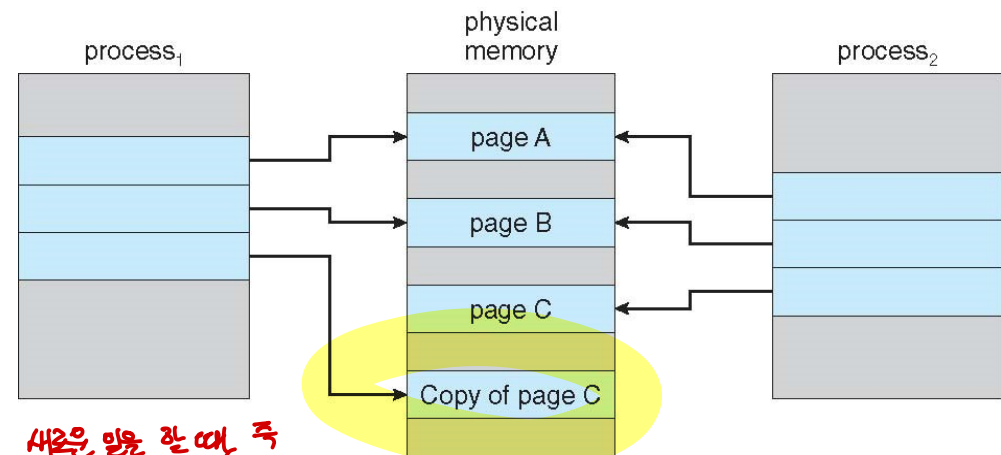
↪ demand paging!

Before Process 1 Modifies Page C



아직 fork(); 하면, child = parent 이므로 그냥 child를 위한 physical memory를 따로 주지 않고 parent가 쓰던 메모리 공간을 share 한다

After Process 1 Modifies Page C



새로운 일을 할 때 즉 'write' 를 할 때 그때마다 copy된 공간을 따로 주게 된다

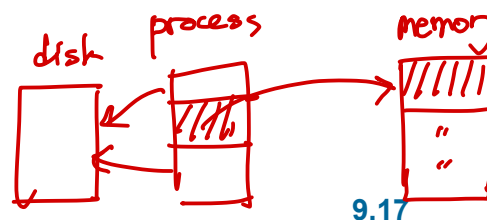
Benefits of Virtual Memory (using Demand Paging)

■ Protection, Transparency, Sharing

- Protection: A process can only access the virtual address (pages); A bug in one process can corrupt memory in another
- Transparency: A process sees a contiguous virtual memory space (pages) cannot access physical memory space directly

■ Resource exhaustion

- Sum of sizes of all processes often greater than physical memory
- Demand paging
 - ▶ e.g., If a process is 20 pages, we can execute with 10 frames and the rest in the swap space (disk) – use demand paging and page-replacement



Chapter 10: Virtual-Memory Management

- Background
- Demand Paging
- **Page Replacement**
 - Basic Page Replacement
 - Modify (Dirty) Bit
- Page Replacement Algorithms
- Thrashing
- Swapping on Mobile Systems

What Happens if There are no Free Frames?

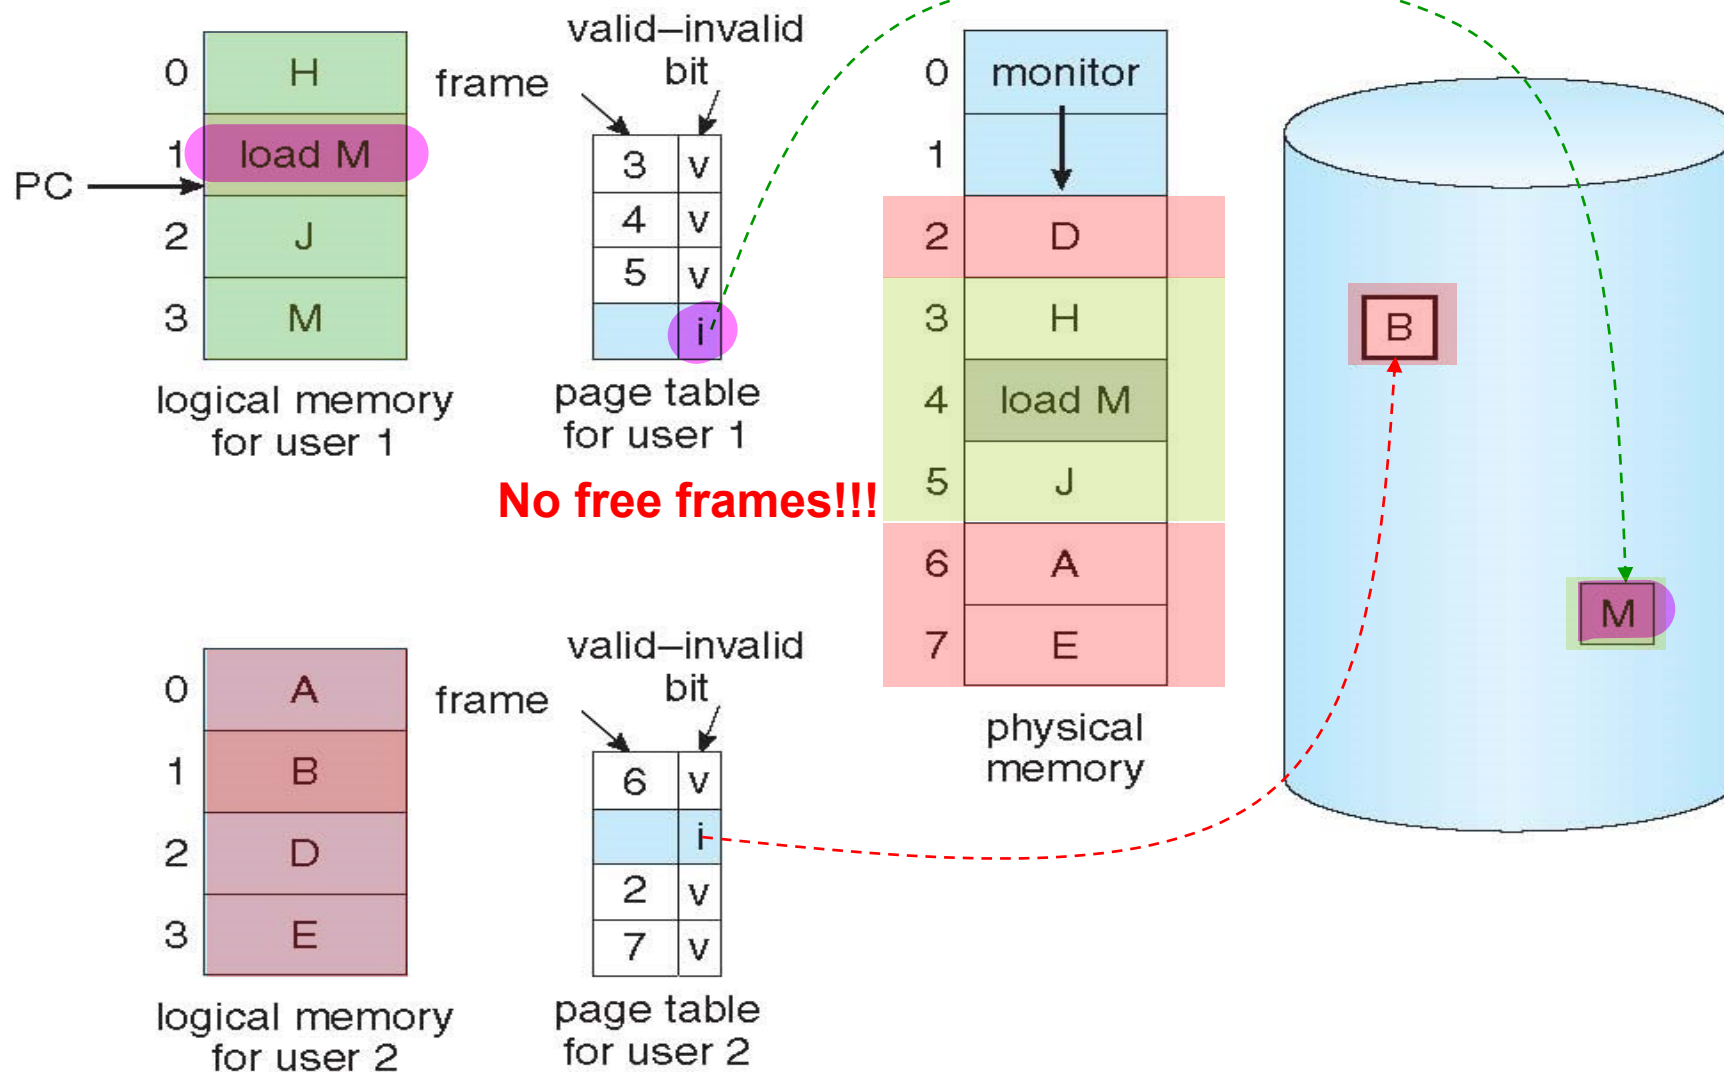
- All memory may be in use
 - Multiprogramming – multiple processes share the memory
 - Page fault → OS finds the desired page on the disk → find the free-frame
 - If NO free frames on the **free-frame list** ? – i.e., physical memory is full!



Swapping을 사용해도 되지만...

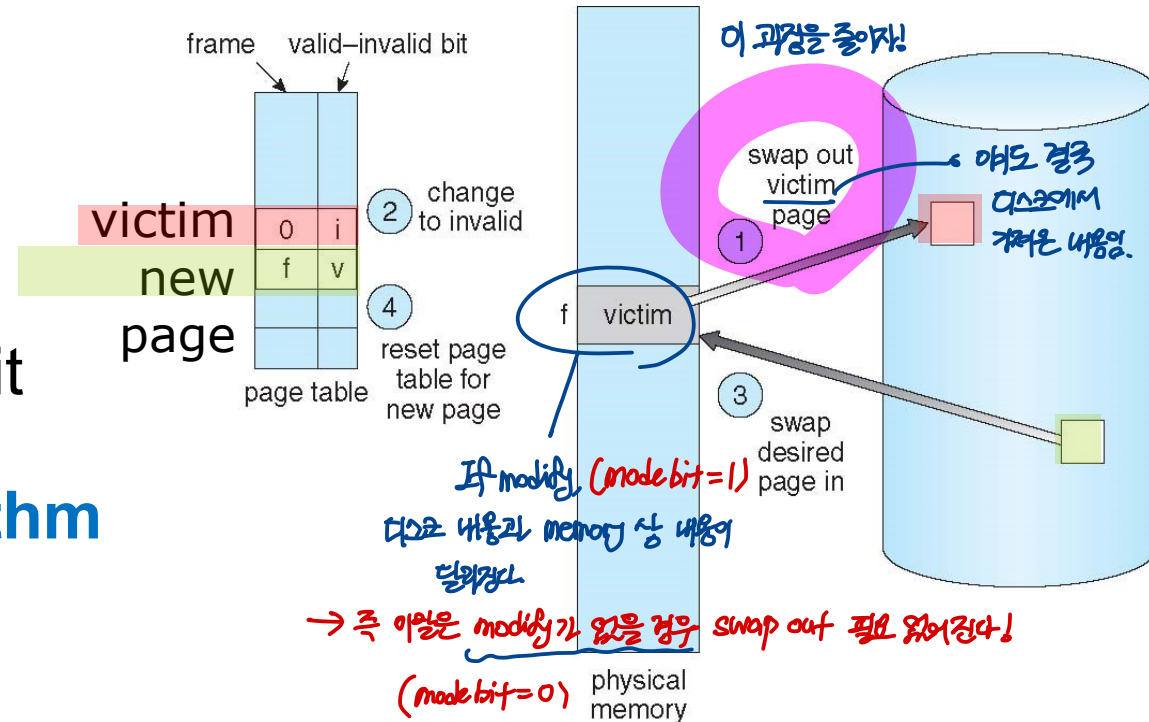
- **Page replacement**
 - Find some page in memory, but not really in use, **page out**
 - Performance – want an algorithm which will result in minimum number of **page faults**
 - ▶ **Page replacement Algorithms**

Need For Page Replacement



Basic Page Replacement

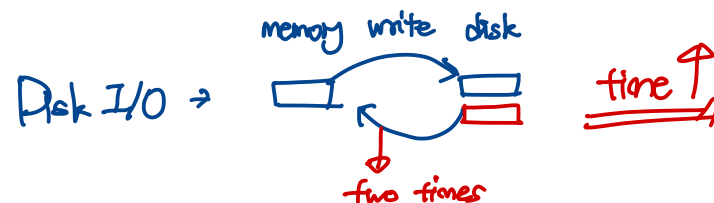
1. Find the location of the desired page on disk
2. Find a **free frame**:
 - If there is a free frame, use it
 - If there is no free frame, use a **page replacement algorithm** to select a **victim frame**



3. Read the desired page into the (newly) free frame. Update the page and frame tables.

새로운 frame 공간을 위해 page table을 제거하는데 - 그것이 victim frame.
그곳에 page를 넣는다.

4. Restart the process



Modify Bit (= Dirty Bit)

- If no frames are free → **TWO** page transfers are required
 - One page out and one page in (see previous slide)
 - Can we reduce the number of disk writes (page-outs)?
 - Solution: **Modify bit (dirty bit)**
 - Initially set to zero
 - Set the bit (to 1) whenever any word or byte in the page is modified
 - ▶ only modified pages are written to disk – unmodified pages can just be deleted without writing on disk (it has **not** changed!)
- ➔ can reduce the page-transfer overhead

Chapter 10: Virtual-Memory Management

- Background
- Demand Paging
- Page Replacement
- **Page Replacement Algorithms**
 - FIFO Page Replacement
 - Optimal Page Replacement
 - LRU Page Replacement
- Thrashing
- Swapping on Mobile Systems

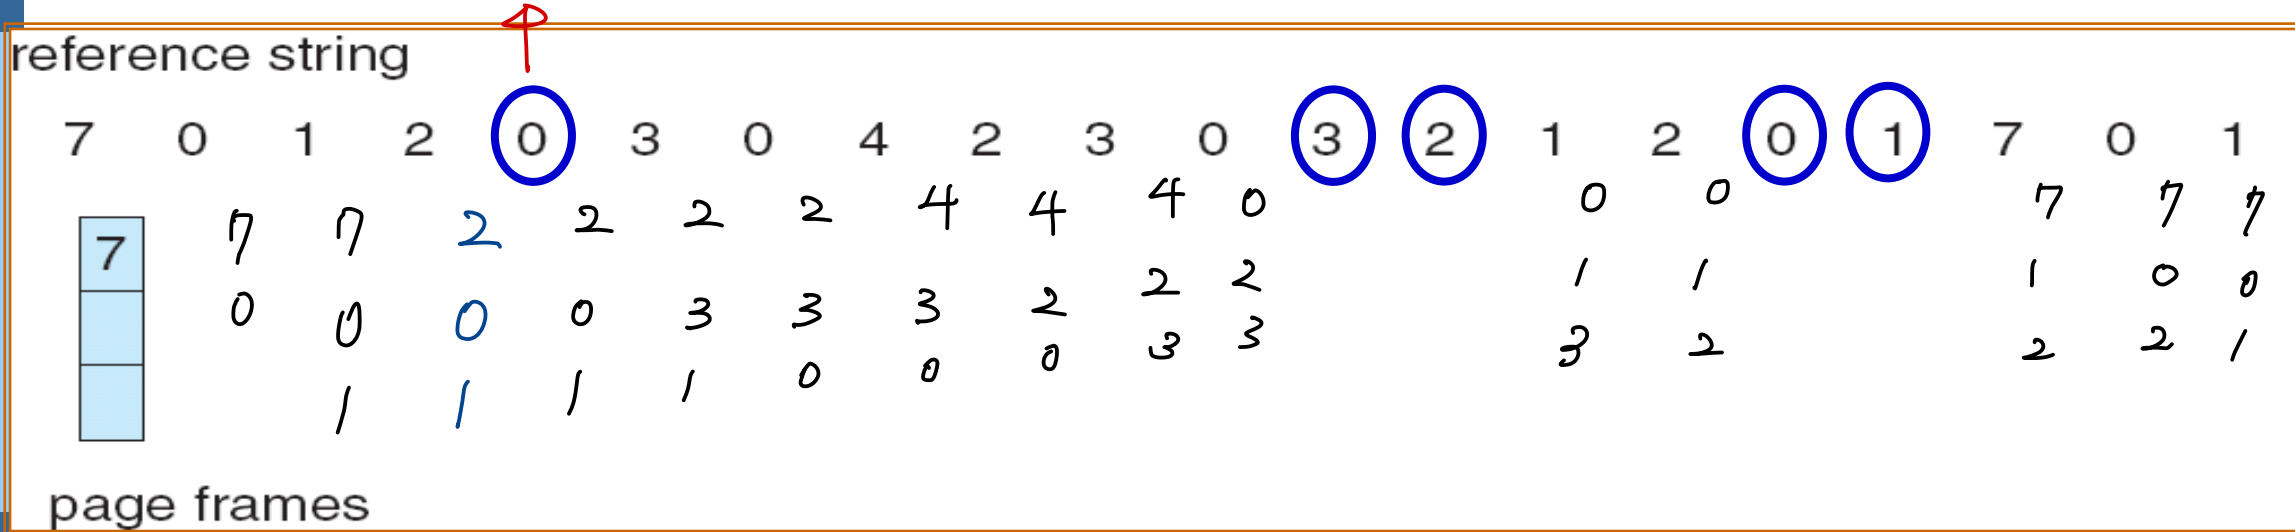
Page Replacement Algorithms

- **Page-replacement algorithm**
 - Which frame should be replaced?
 - Want lowest page-fault rate
 - Evaluate algorithm by computing the number of **page faults**
- FIFO page replacement
- Optimal page replacement
- LRU (Least-Recently-Used) algorithm
 - clock, reference-bit, second-chance algorithm

First-In-First-Out (FIFO) Algorithm

- When a page must be replaced, the oldest page is chosen

애초에 page-fault가 발생하지 않으면 page-replacement도 필요 없다.



12 Page Replacements!! (15 Page Faults)

- How to track ages of pages?
 - Just use a FIFO queue

Replacement Algorithm	Page replacements	Page faults
FIFO	12	15
OPT		
LRU		

Problem of FIFO Algorithm

■ Example 2

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

-3 frames (3 pages can be in memory at a time per process)

1	1	4	5
2	2	1	3
3	3	2	4

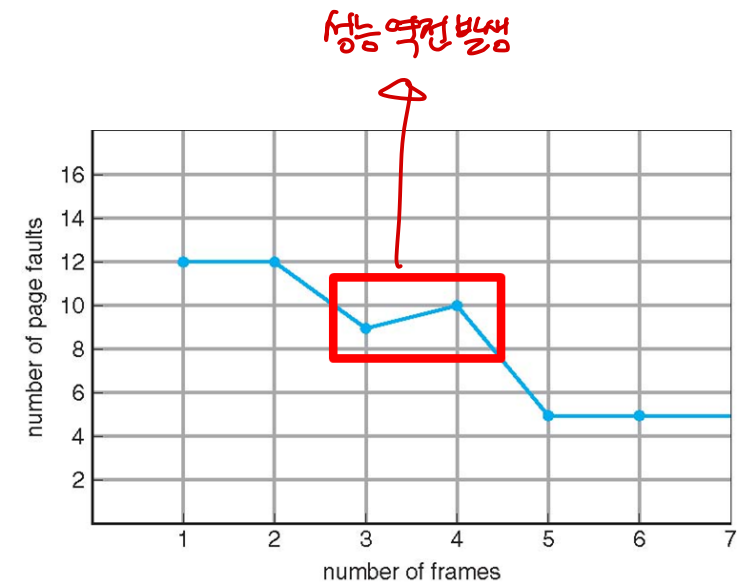
9 page faults

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

-4 frames:

1	1	5	4
2	2	1	5
3	3	2	
4	4	3	

10 page faults



Why did this happen???

FIFO Illustrating Belady's Anomaly

FIFO replaces the **oldest** page – is this always good?

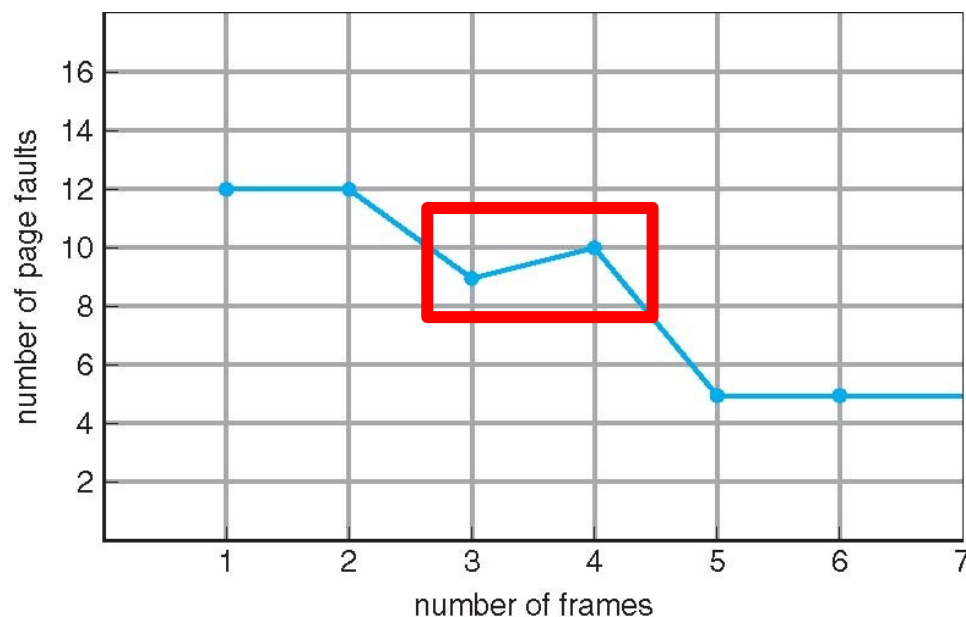
An old page may contain a **frequently used variable**



Belady's anomaly

예외적인 경우가 종종 발생할 수 있다.

Sometimes, page-fault may increase as the number of allocated frames increase!



Optimal Algorithm

- Algorithm that has the lowest page-fault rate
- Replace page that **will not** be used for longest period of time
 - This is a design to guarantee the lowest page-fault rate for a fixed number of frames

92개 안 "4" page. 다라기 안쓰 page...

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2	2
	0	0	0	0
		1	1	1

page frames

**6 Page Replacements!!
(9 Page Faults)**

Replacement Algorithm	Page replacements	Page faults
FIFO	12	15
OPT	6	9
LRU		

Optimal Algorithm

- Unfortunately, the optimal page-replacement algorithm is difficult to implement
 - Why? ? 어떻게 알아
 - ▶ Require **future knowledge** of the reference string – We do not know the future! (Notes: Have we seen this before?) → scheduling SJF
- Mainly used for comparison studies
 - Evaluating a new algorithm
 - e.g. “the new algorithm is within 12.3% of optimal at worst and within 4.7% on average”

Least Recently Used (LRU) Algorithm

- FIFO: Use time when a page **was brought** into the memory
- OPT: Use time when a page is to be **used**
- Let us predict the future by using the past!

가장 오래동안 쓰지 않은 것

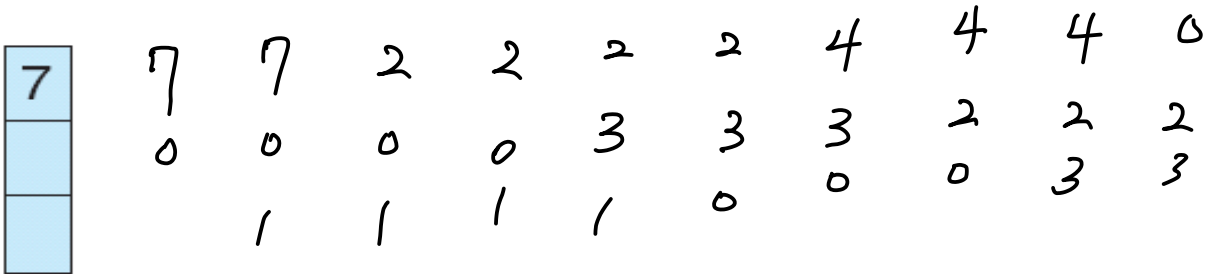
- **Least Recently Used (LRU) Algorithm**
 - Use past knowledge rather than future
 - **Replace** page that has **not been used** in the most amount of time
 - Generally good algorithm and frequently used

Least Recently Used (LRU) Algorithm

- LRU replacement associates with each page the time of that page's last use
- Replace the page that has not been used for the longest period of time

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

**9 Page Replacements!!
(12 Page Faults)**

Replacement Algorithm	Page replacements	Page faults
FIFO	12	15
OPT	6	9
LRU	9	12

Implementation of LRU Algorithm

- The major problem is **how to implement LRU replacement**
 - Use **counter**
 - LRU implementation requires H/W support (due to speed)
- Counter implementation example



Implementation of LRU Algorithm

- **Counter** implementation
 - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
 - When a page needs to be changed, look at the counters – replace page with **smallest** timer value
 - ▶ Con1: Need to **search** to replace page (complexity **$O(N)$** with N frames)
 - ▶ Con2: Counter must be updated for each memory reference

LRU Implementation

- Updating of **counter fields** must be done for **every** memory reference
- LRU needs special hardware and still slow
 - Counter approach use too much resources
- LRU Approximation: **Reference bit**
 - With each page associate a bit, initially = 0
 - When a page is referenced, then bit set to 1
 - Replace **any** with **reference bit = 0** (if one exists)
 - In this method, although we can not know the **exact order** of page reference, we can know that whether the page is **referenced or not**
 - Basis for many page-replacement algorithms that **approximate** LRU replacement

인번이로 참조되었다면 1로 변경

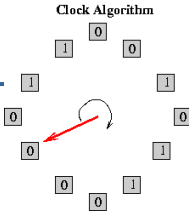
원번도 참조 안된 것들을 하나 골라

대체한다.

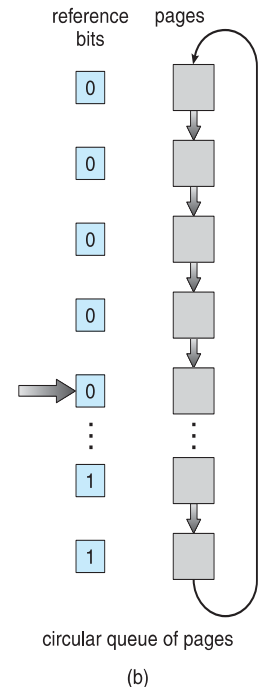
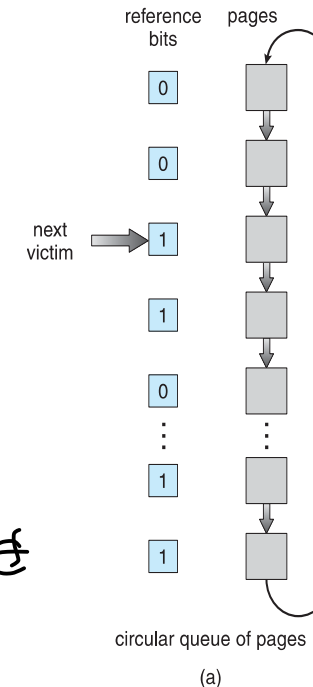
LRU Approximation: Additional-Reference-Bits Algorithm

- 8-bit byte for each page in a table in memory
- At regular intervals (e.g., 100ms), the reference bit for each page is shifted right by 1-bit
- Example 100ms 마다, 참조된 경우 1, 아니면 0. 그래서 가장 작은 값을 replace한다.
 - 00000000: page has not been used for 8 time periods
 - 11111111: page has been used at least once in each period
 - 11000100: page has been used more recently than (01110111)
- Page with lowest integer number is LRU – it can be replaced

Second-chance algorithm (Clock algorithm)



- Use one reference bit per page using **circular list**
- **Pointer** (hands on clock) indicates which page is to be replaced next
- If page to be replaced has
 - **reference bit = 1** then:
 - ▶ set **reference bit 0**, leave page in memory
 - ▶ move on to next page
 - **reference bit = 0** -> **replace it**
 - ▶ this was the second chance
 - ▶ if this page was reference for in the last cycle, then reference bit should have been 1



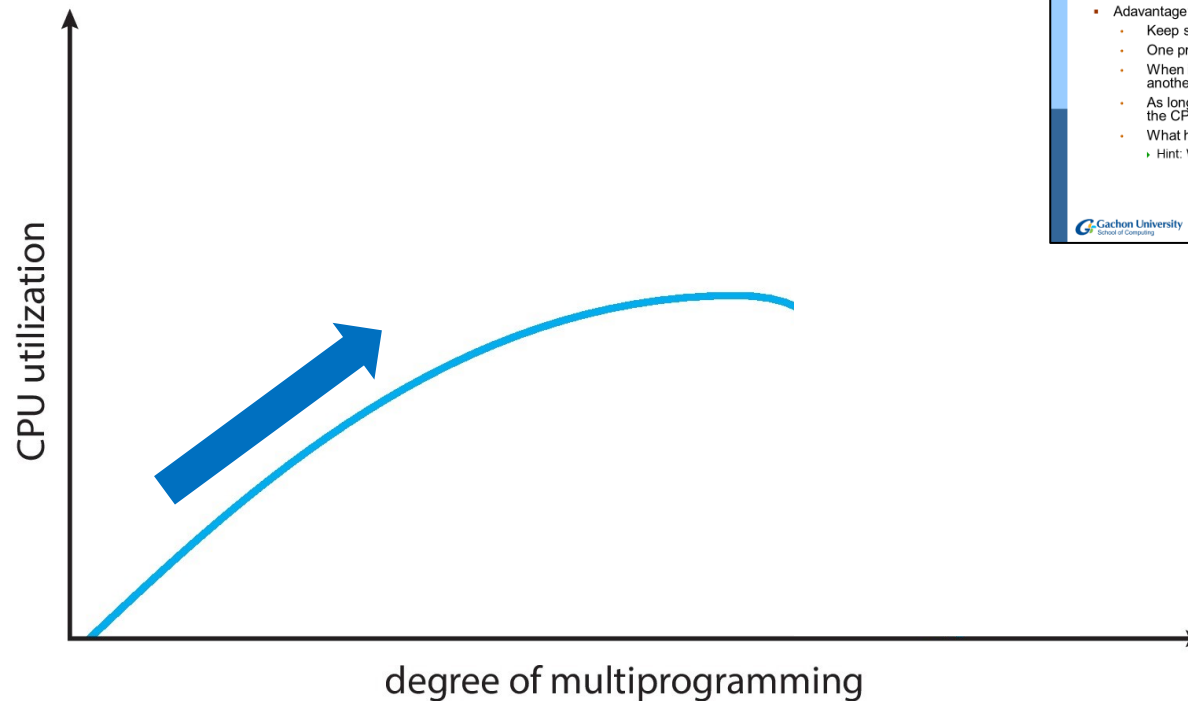
두번 연속 0이면 replace

Chapter 10: Virtual-Memory Management

- Background
- Demand Paging
- Page Replacement
- Page Replacement Algorithms
- **Thrashing**
- Swapping on Mobile Systems

Degree of Multiprogramming

- **Degree of multiprogramming:** Number of programs in memory.
- If CPU utilization is low, operating system thinks that it needs to **increase** the degree of multiprogramming
 - Another process added to the system. But, ...



Ch. 1

Multitasking

- **Multitasking** (=Multiprogramming, Multiprocessing) needed for efficiency
 - Users frequently have **multiple processes** loaded on **Main memory**
 - Single process *cannot* keep **CPU** and **I/O devices** busy at all times
- Advantage1: **CPU utilization**
 - Keep several processes (=job, task) in memory simultaneously
 - One process selected and run via **scheduling**
 - When it has to wait (e.g., for I/O operation), OS **switches** to another process
 - As long as at least one process needs to execute, the CPU never stays idle
 - What happens if it is not switched?
 - ▶ Hint: Waste CPU cycles (why?)

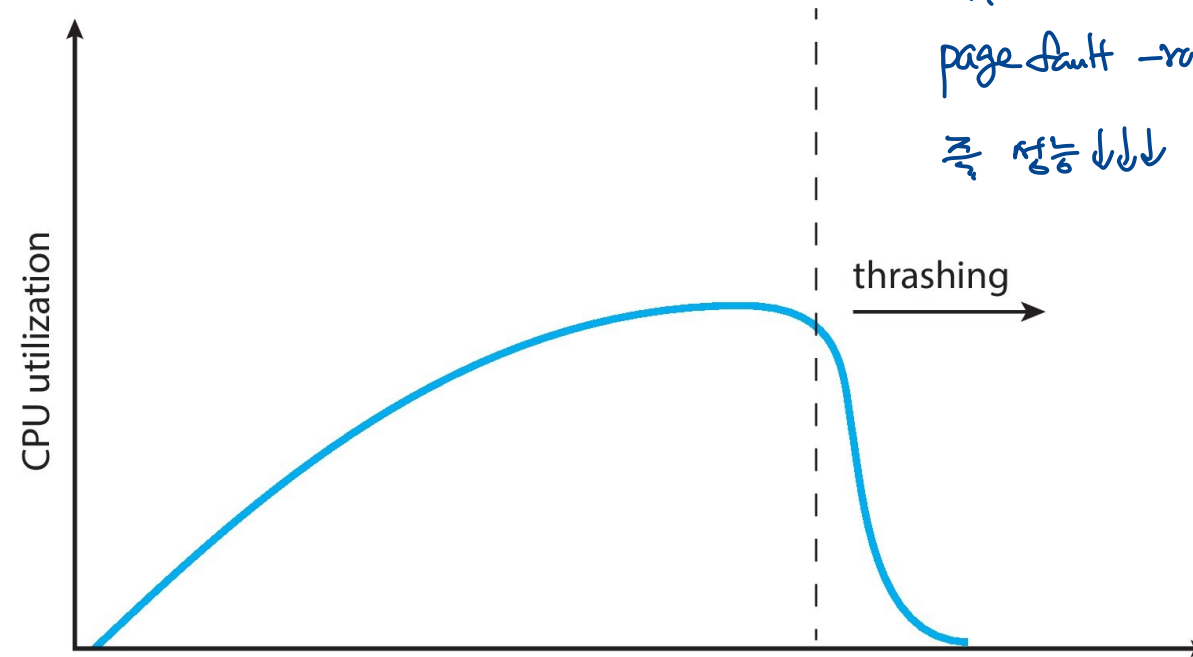
The diagram on the right shows a vertical stack of four boxes labeled 'job 1', 'job 2', 'job 3', and 'job 4'. To the left of the stack, a vertical axis is marked with '0' at the top and '512M' at the bottom. Below the diagram is a screenshot of the Windows Task Manager 'Performance' tab, showing the 'CPU' section with a graph and various statistics.

Gachon University School of Computing

1.8

Thrashing

- **Thrashing.** A process is busy dealing with page faults
 - If a process does not have “enough” pages, the page-fault rate is very high
 - Page fault to get page
 - Replace existing frame
 - But quickly need replaced frame back



너무 많이 실행시키면, frame이 부족해서
page fault -rate가 많이 올라간다.
즉 성능↓↓↓

Many pages per
process

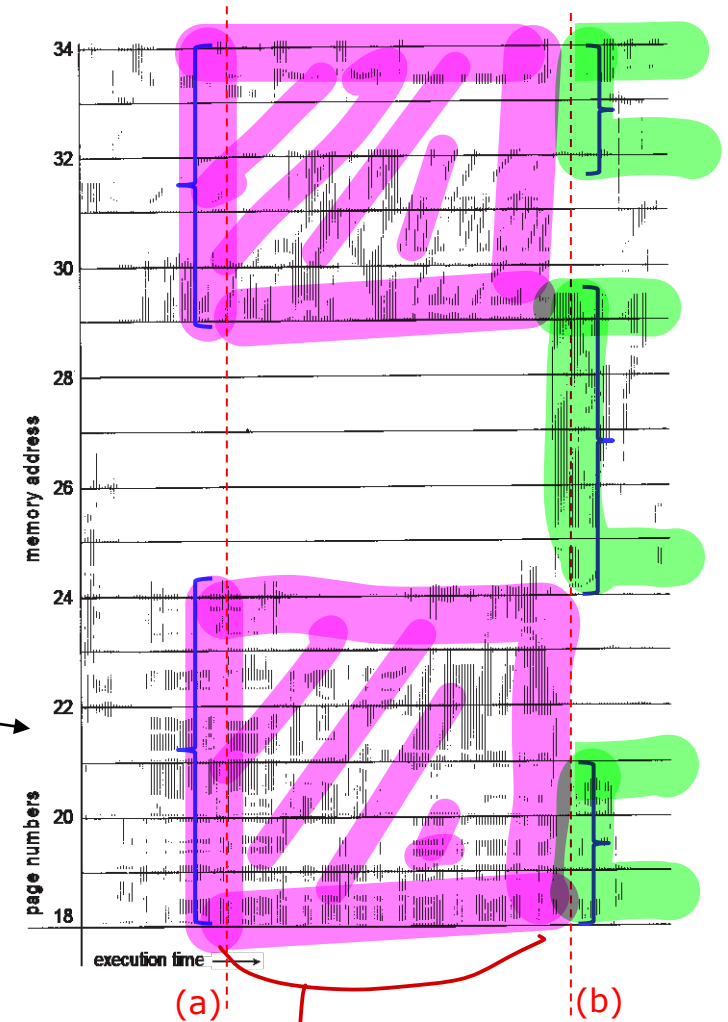
degree of multiprogramming

Fewer pages per process
– more page faults –
lower CPU utilization

Demand Paging and Thrashing

- To prevent thrashing
 - provide a process with as many frames as it needs
 - How many does it need?
- Locality model *메모리의 특정 지역을 특정 시간대에 점유하고 있다!*
 - Process migrates from one locality to another (e.g., from (a) to (b))
 - Localities may overlap
- Why does thrashing occur?
 $\Sigma \text{ size of locality} > \text{free frames}$

필요한 메모리 - locality 보다 free frame가 적으면 thrashing 발생



Locality In A Memory-Reference Pattern

비슷한 곳을 계속 사용!

Working-Set Model

앞으로 계속 사용할 것 같은 곳을 메모리에 놓는 것

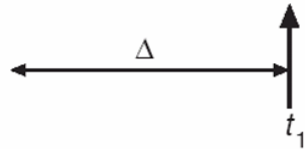
- **working-set window** (Δ): the most recent Δ pages

- an approximation of the program's locality
- Example: $\Delta=10$

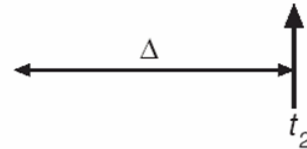
최근에 사용한 페이지는 무엇인가?

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$WS(t_1) = \{1, 2, 5, 6, 7\}$



$WS(t_2) = \{3, 4\}$

- if Δ too small will not encompass entire locality
- if Δ too large will encompass several localities
- if $\Delta = \infty \Rightarrow$ will encompass entire program

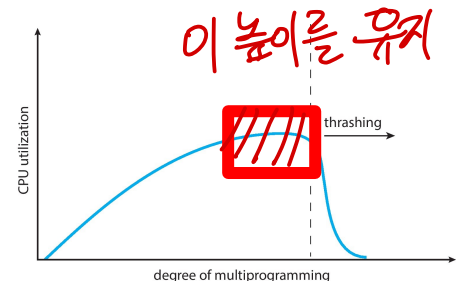
너무 작으면 locality를 다 저장하지 못함

예전 데이터 구간도 생각해야 함으로 낭비

의미없음.

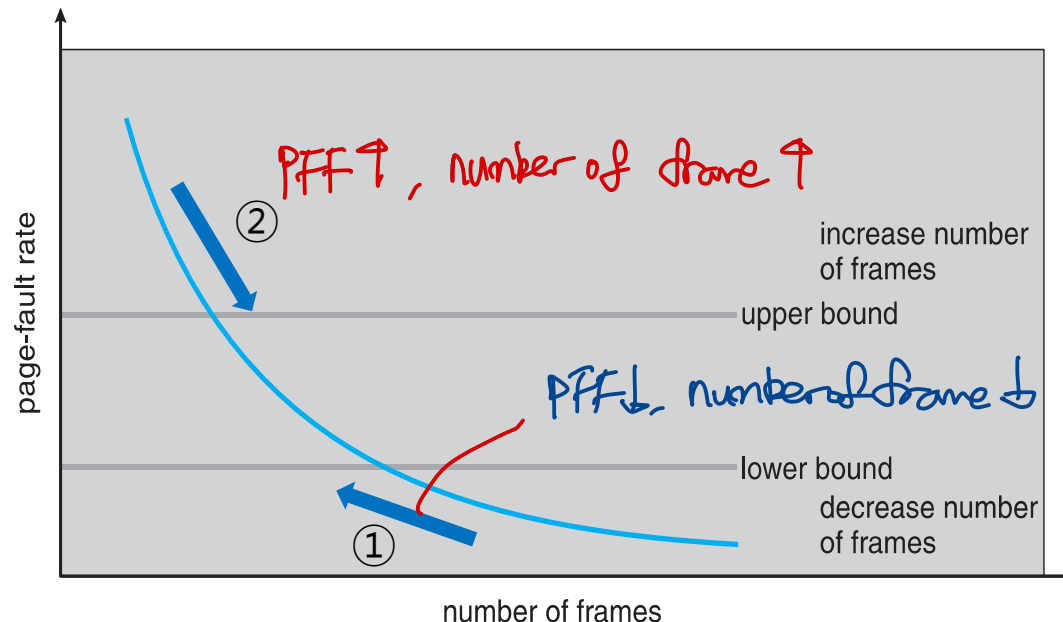
Working-Set Model (Cont.)

- $D = \sum WSS_i \equiv$ total demand frames 모두라는 frame 총 개수
 - WSS_i : working set size for each process
- Policy
 - if $D < m$ (m : available frames), then another process can be initiated
 - if $D > m$, then Thrashing! \Rightarrow suspend or swap out a process
 \Rightarrow degree of multiprogramming (thus thrashing) reduces
 - Suspended process can restart later
- Prevents thrashing while keeping degree of multiprogramming as high as possible
 - optimizes CPU utilization



Page-Fault Frequency (PFF)

- More direct approach than working set model
- Establish “acceptable” **page-fault frequency (PFF)** rate and use local replacement policy
 - If actual rate too low, process loses frame (①)
 - If actual rate too high, process gains frame (②)
- If page fault rate increases and not enough free frames, then process swapped out \Rightarrow degree of multiprogramming (thus thrashing) reduces



Chapter 10: Virtual-Memory Management

- Background
- Demand Paging
- Page Replacement
- Page Replacement Algorithms
- Thrashing
- **Swapping on Mobile Systems**

Swapping on Mobile Systems (Ch. 9.5.3)

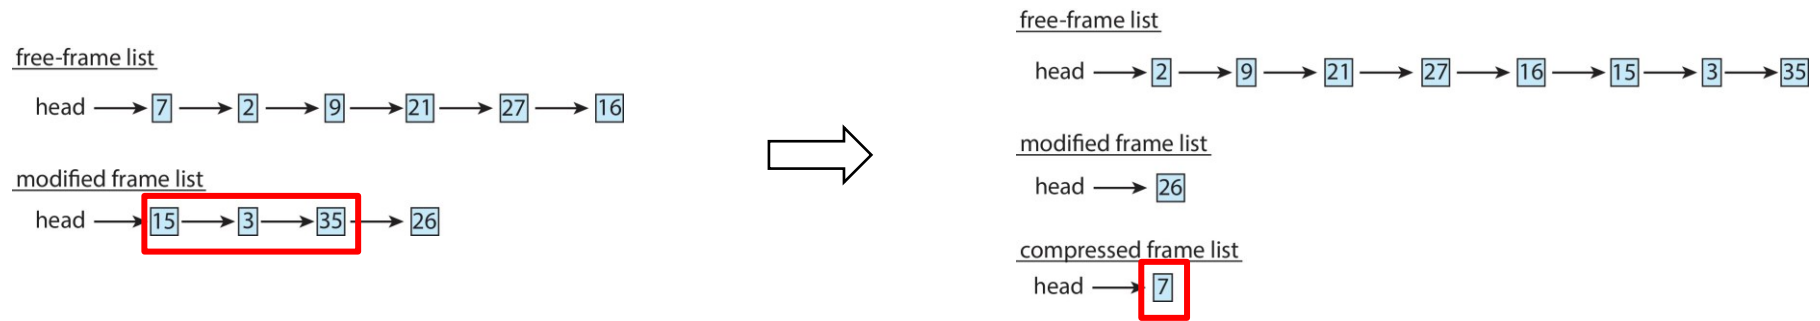


- Typically, mobile systems do not support swapping pages, since they use **Flash memory** instead of disks
 - flash memory space constraint
 - flash memory becomes unreliable and shows poor throughput after some number of writes
- Apple iOS: ask application to voluntarily relinquish allocated memory
 - Read-only data (e.g., code) are removed from main memory, later reload if necessary
 - Some applications may be terminated by OS
- Android: Similar to iOS, but Android writes application state to flash memory before App termination for quick restart.



Memory Compression (Ch. 10.7)

- **Memory compression** in mobile systems
 - Rather than paging out modified frames to swap space, compress several frames into a single frame



- **Compression ratio**: amount of reduction achieved by compression algorithm
 - High compression ratio: slower, more computation

Summary

- Virtual memory is how we stuff large programs into small physical memories.
- We perform this by using demand paging, to bring in pages only when they are needed.
- But to bring pages into memory, means kicking other pages out, so we need page replacement algorithms.



<https://www.istockphoto.com/kr/%EC%9D%BC%EB%9F%AC%EC%8A%A4%ED%8A%B8/qa>