


Chapter 3: Process Concept

School of Computing, Gachon Univ.
Joon Yoo



Objectives

- To introduce the notion of a process -- a program in execution, which forms the basis of all computation
- To describe the various features of processes, including scheduling, creation and termination, and communication
- To explore interprocess communication using shared memory and message passing

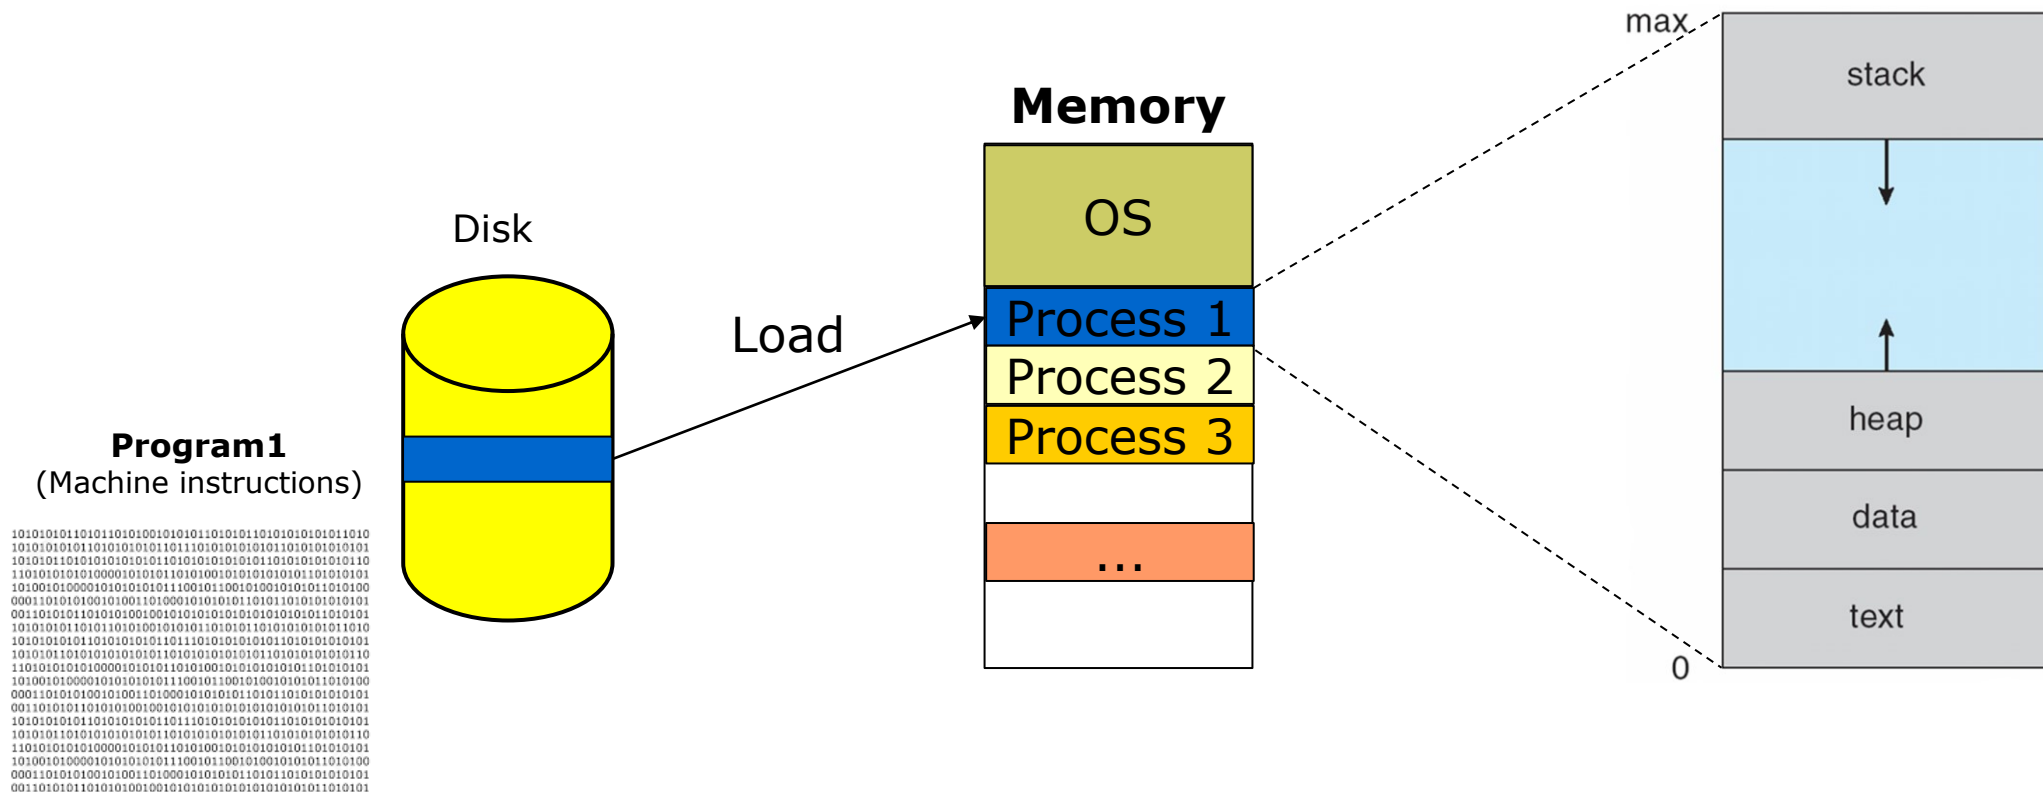
Chapter 3: Process Concept

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication

Process Concept

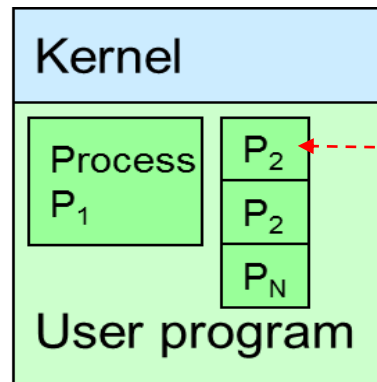



- **Process** (= job/task, ≠program)
 - a program *in execution*



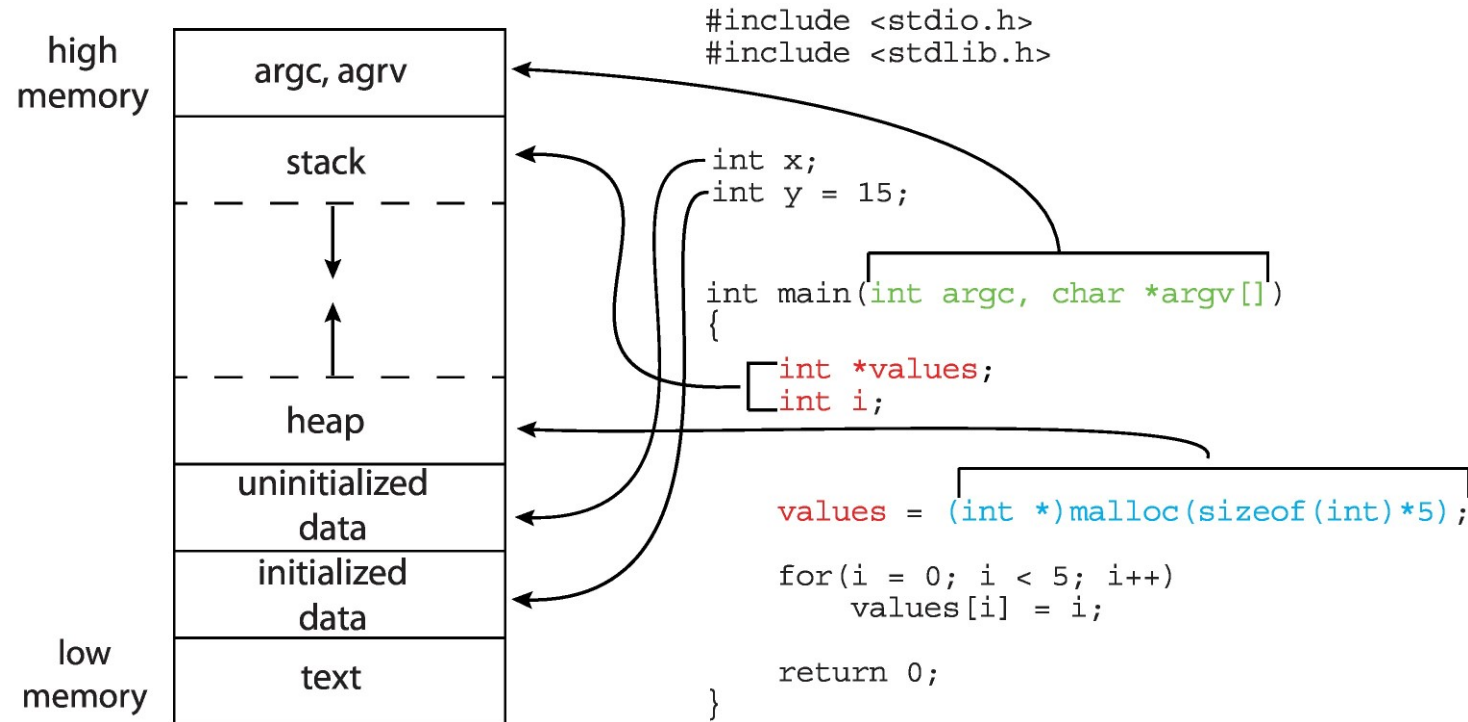
Process in Memory

- Process in memory
 - **Text (Code)**: The binary **program instructions**
 - **Data**: Static data. (e.g., global **variables**)
 - **Stack**: temporary local **data**
 - ▶ Function parameters, return addresses, local variables
 - **Heap**: memory **dynamically** allocated **during run time** (e.g., C malloc(), Java objects)



Memory

Memory Layout of a C Program



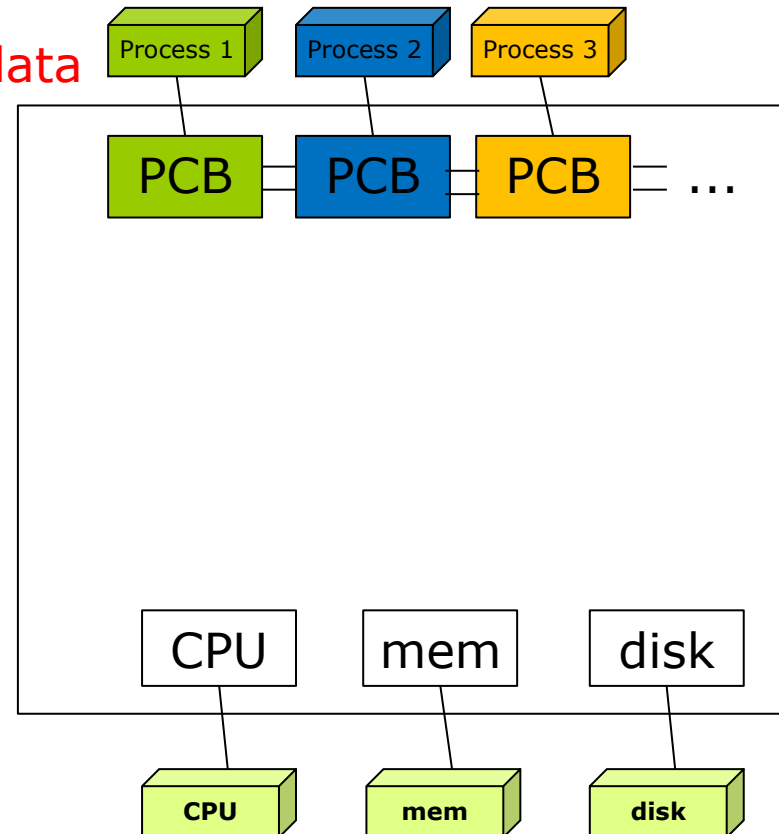
Kernel Memory Space


Kernel code (text)

Kernel code

- System call, Interrupt handling code
- Resource management code
- Others...

data




 : Table (Data Structure)

 : Object (HW or SW)

Process Control Block (PCB)



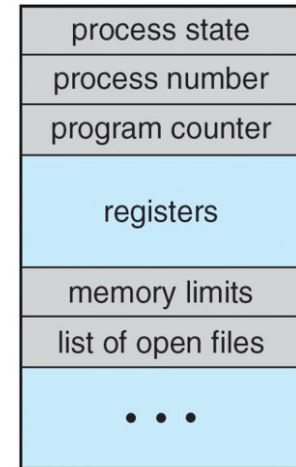
- Each process is registered to and managed by OS
 - manages and take care of all the processes.
 - Therefore, the OS needs to manage the current information (e.g., state) of each process – use a data structure called **PCB (Process Control Block)**



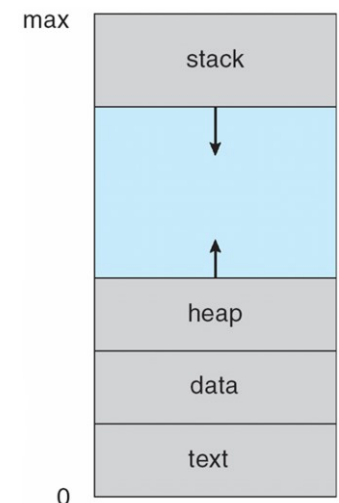
Process Control Block (PCB)

PCB: OS maintains the **information** for *each* process

- Process **state** (next slide)
- Process **number**: process id (pid)
- **Program counter (PC)** – next instruction address
- **CPU registers** – contents of registers (in CPU)
- CPU scheduling information (Ch. 5)
- Memory-management information (Ch. 9-10)
 - Where is the process located in **memory**?
- I/O status information (Ch. 12-15)
 - I/O devices allocated to process, list of open files

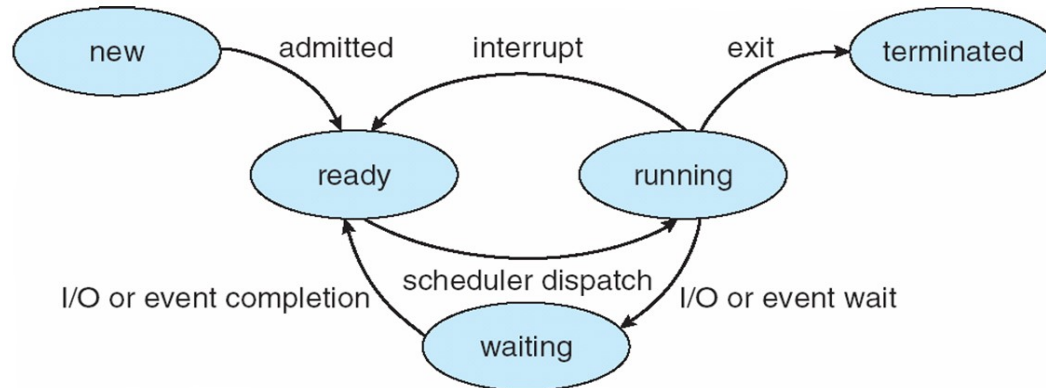


PCB



Process State

- As a process executes, it changes **state**
 - **new**: Before process is created (not a process yet)
 - **ready**: The process (in memory) is ready to be assigned to CPU
 - **running**: Process instructions are being executed by CPU
 - **waiting**: The process is waiting for some event (e.g., I/O operation) to occur
 - **terminated**: The process has *finished* execution (not a process anymore)



Important!

- Only **one** process is **running** on any CPU core at any instant
- All the other processes are waiting in **ready** or **waiting** states

Chapter 3: Process Concept


- Process Concept
- **Process Scheduling**
- Operations on Processes
- Interprocess Communication

Scheduling Queues

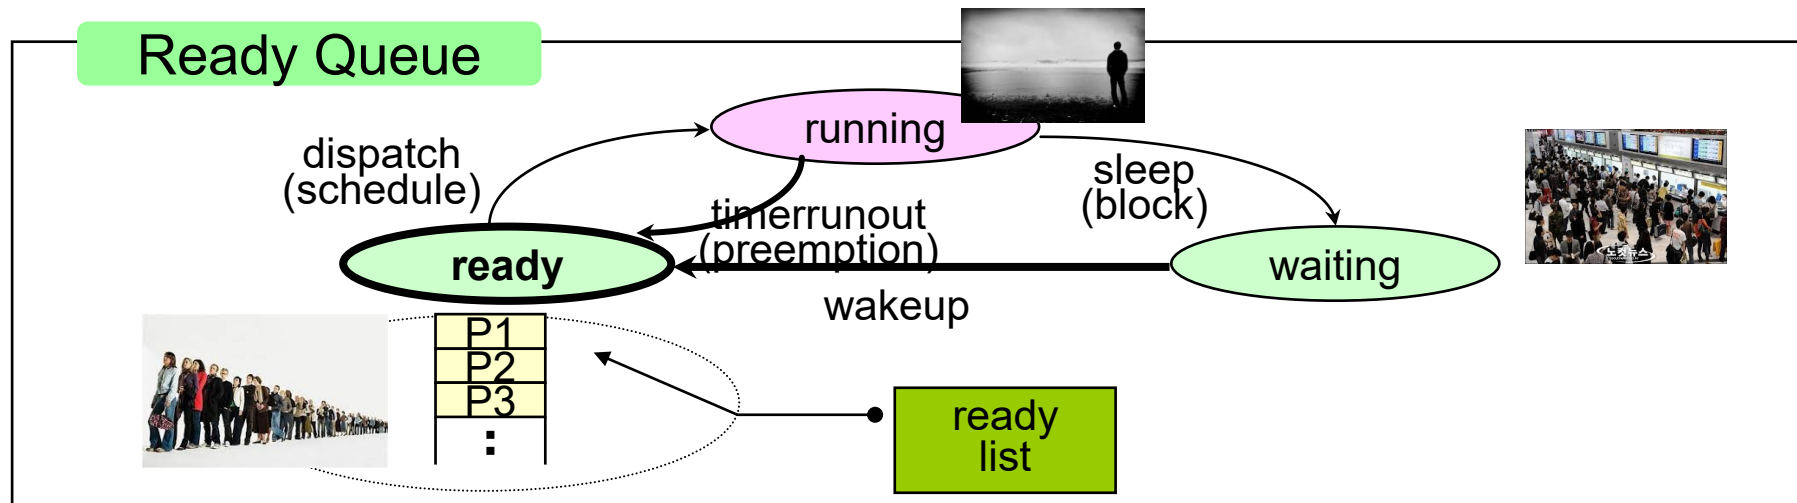


■ Ready queue

- The processes that are ready to execute on a CPU's core (in **Ready state**; waiting for **CPU**)




- Question: How many processes can be **running** at the same time?
- Question: How many processes can be **ready** at the same time?



CPU Scheduler

■ CPU Scheduler

- Selects from among the processes in ready queue, and allocate the CPU core to one of them.



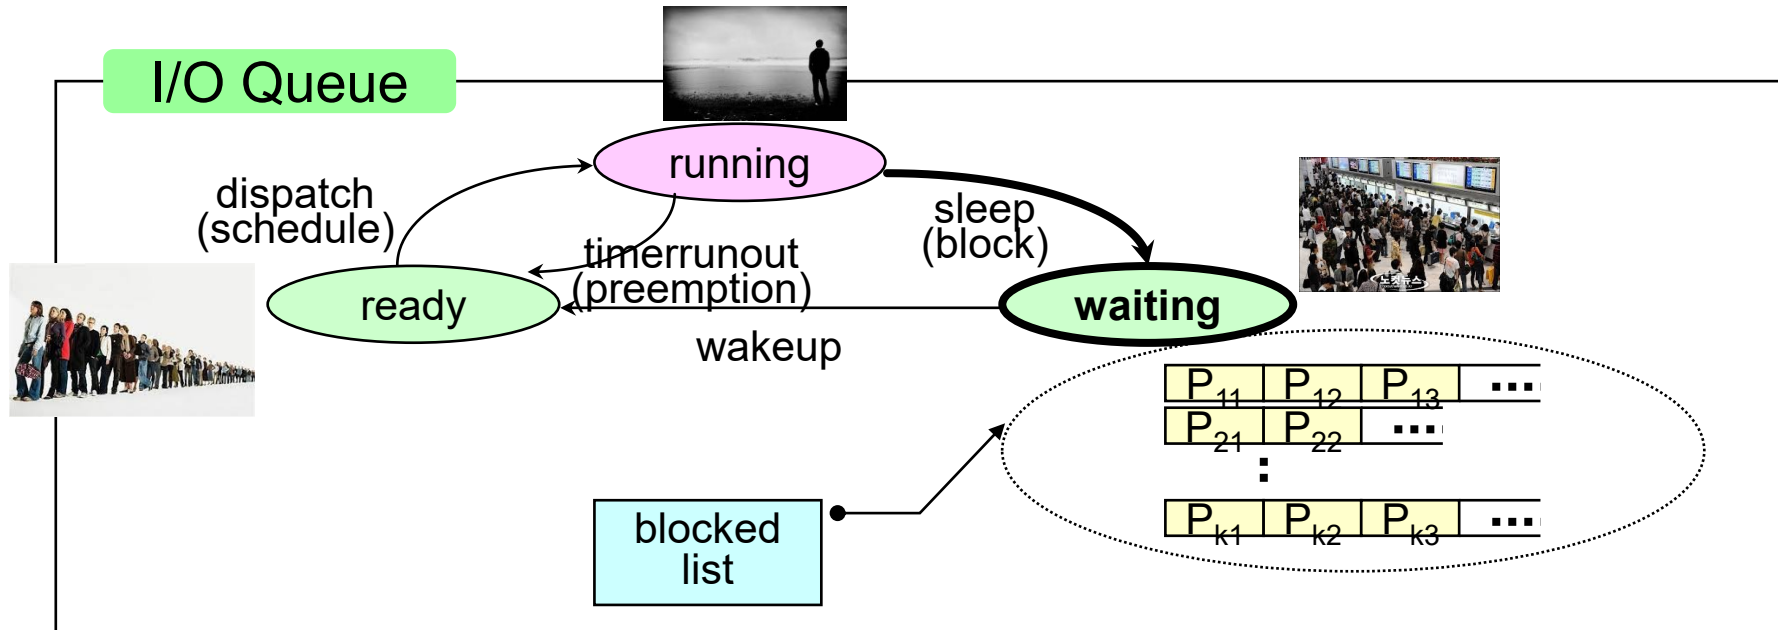
- Use CPU scheduling algorithms (Ch. 5)

Scheduling Queues




■ Waiting queue

- When a process is allocated the CPU, it **executes** for a while and eventually **quits** or **interrupted**
- or **waits for an event** (e.g., I/O completion)
 - ▶ The process has to be **waiting** (or **blocked**) in the **wait queue**
 - ▶ Also called Blocked list (block queue, I/O queue)




Ready Queue and Wait Queues Linux Data Structure



I/O
(device)
queue

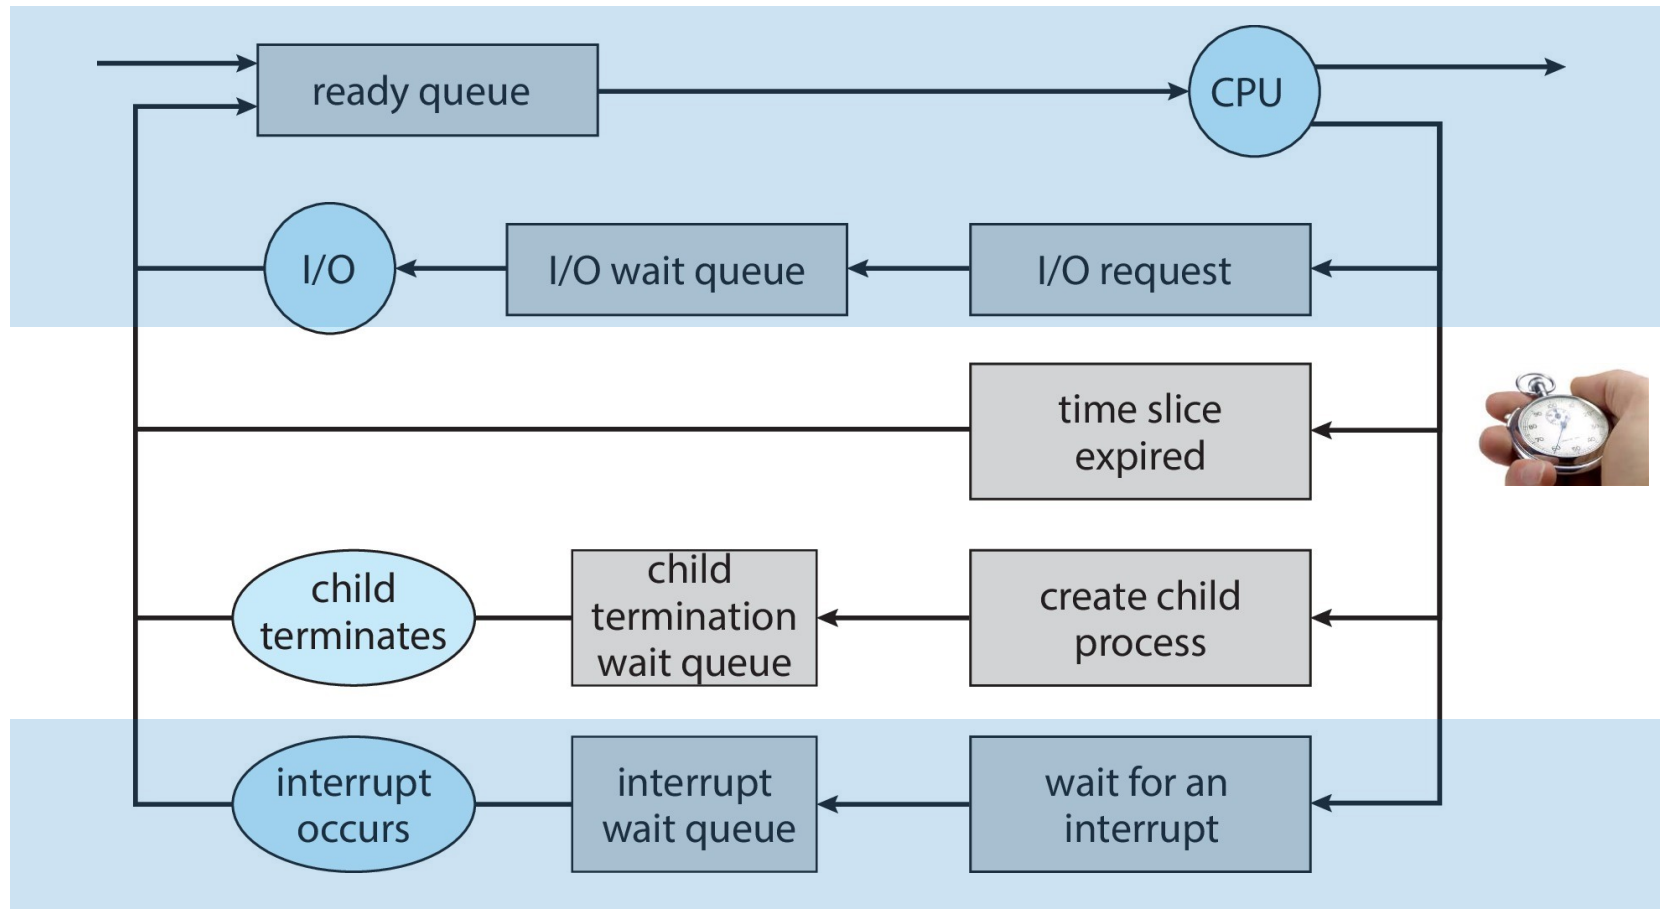
- `<include/linux/sched.h>`
- Queues are usually implemented with linked lists
 - Queue header → pointing the first and last **PCB** structures
 - Each **PCB** structure has the address of its next **PCB** structure



PCB

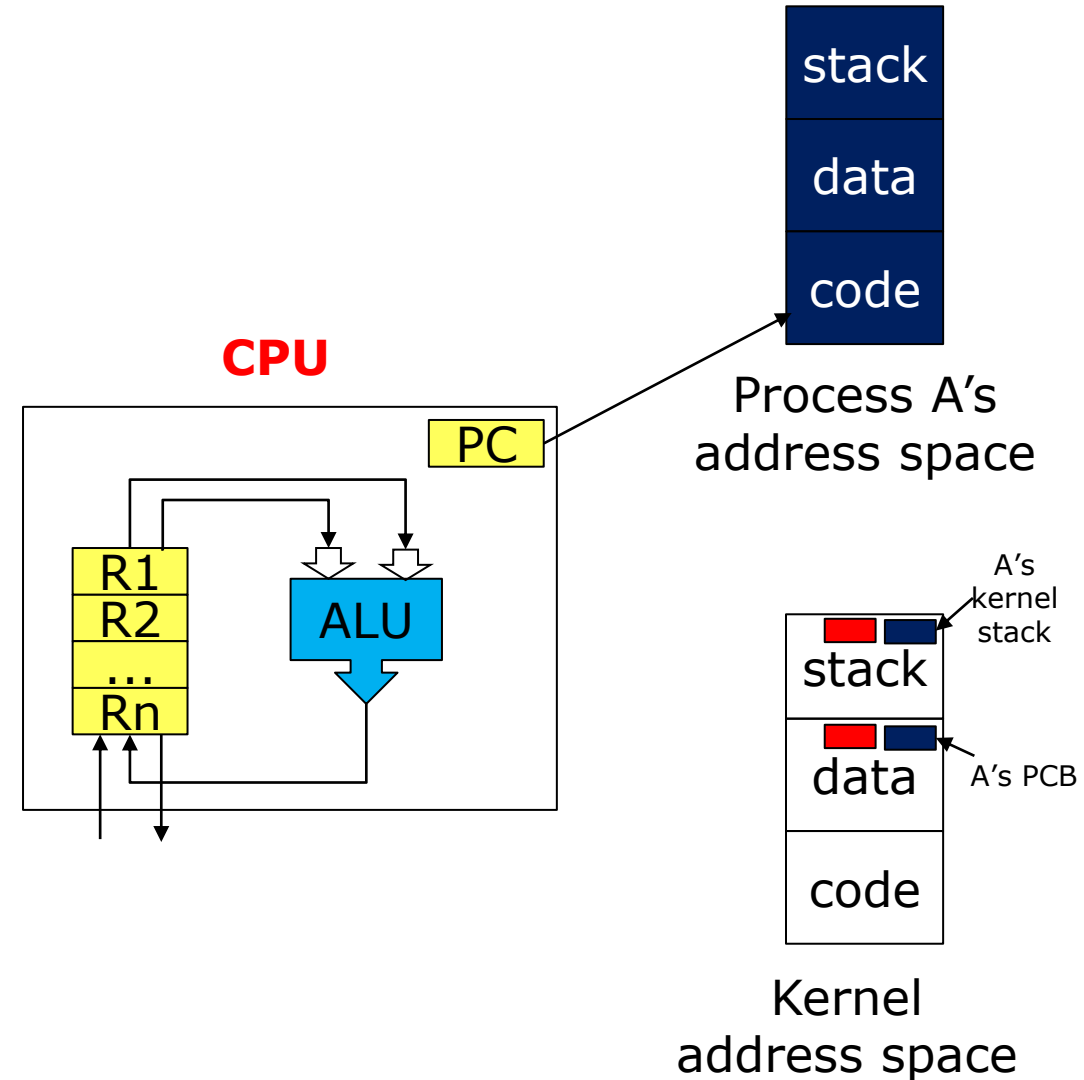
Representation of CPU Scheduling

Queuing diagram



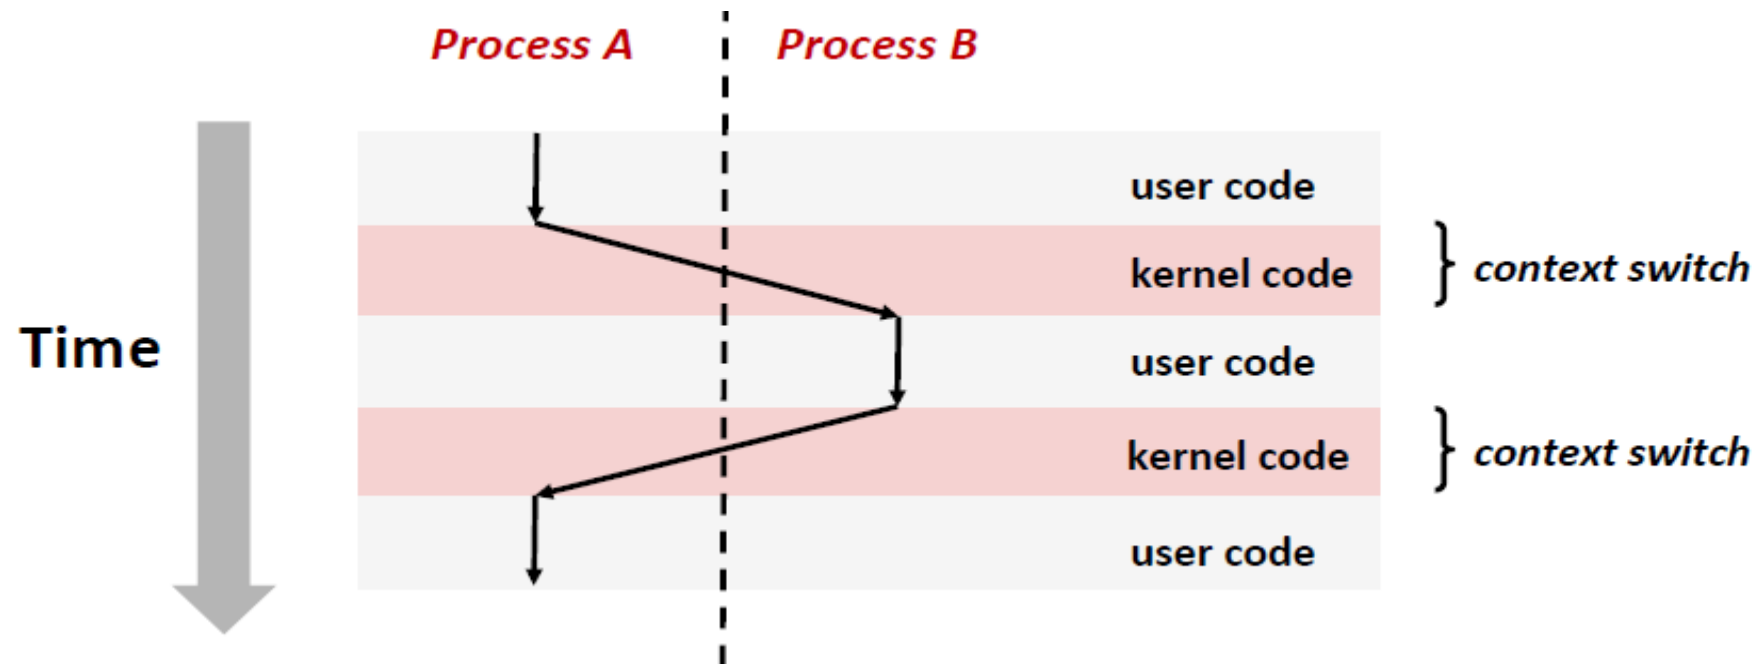
Process Context

- CPU execution context
 - Program Counter (PC)
 - Registers
- Process memory space
 - code, data, stack
- Process management in OS
 - Process Control Block (PCB)
 - Kernel stack



Context Switch

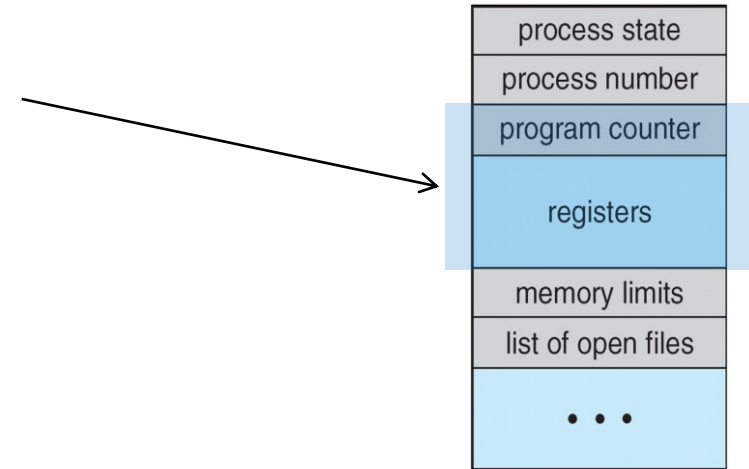
- Process switching from one process to another by OS !
 - called **Context switch**



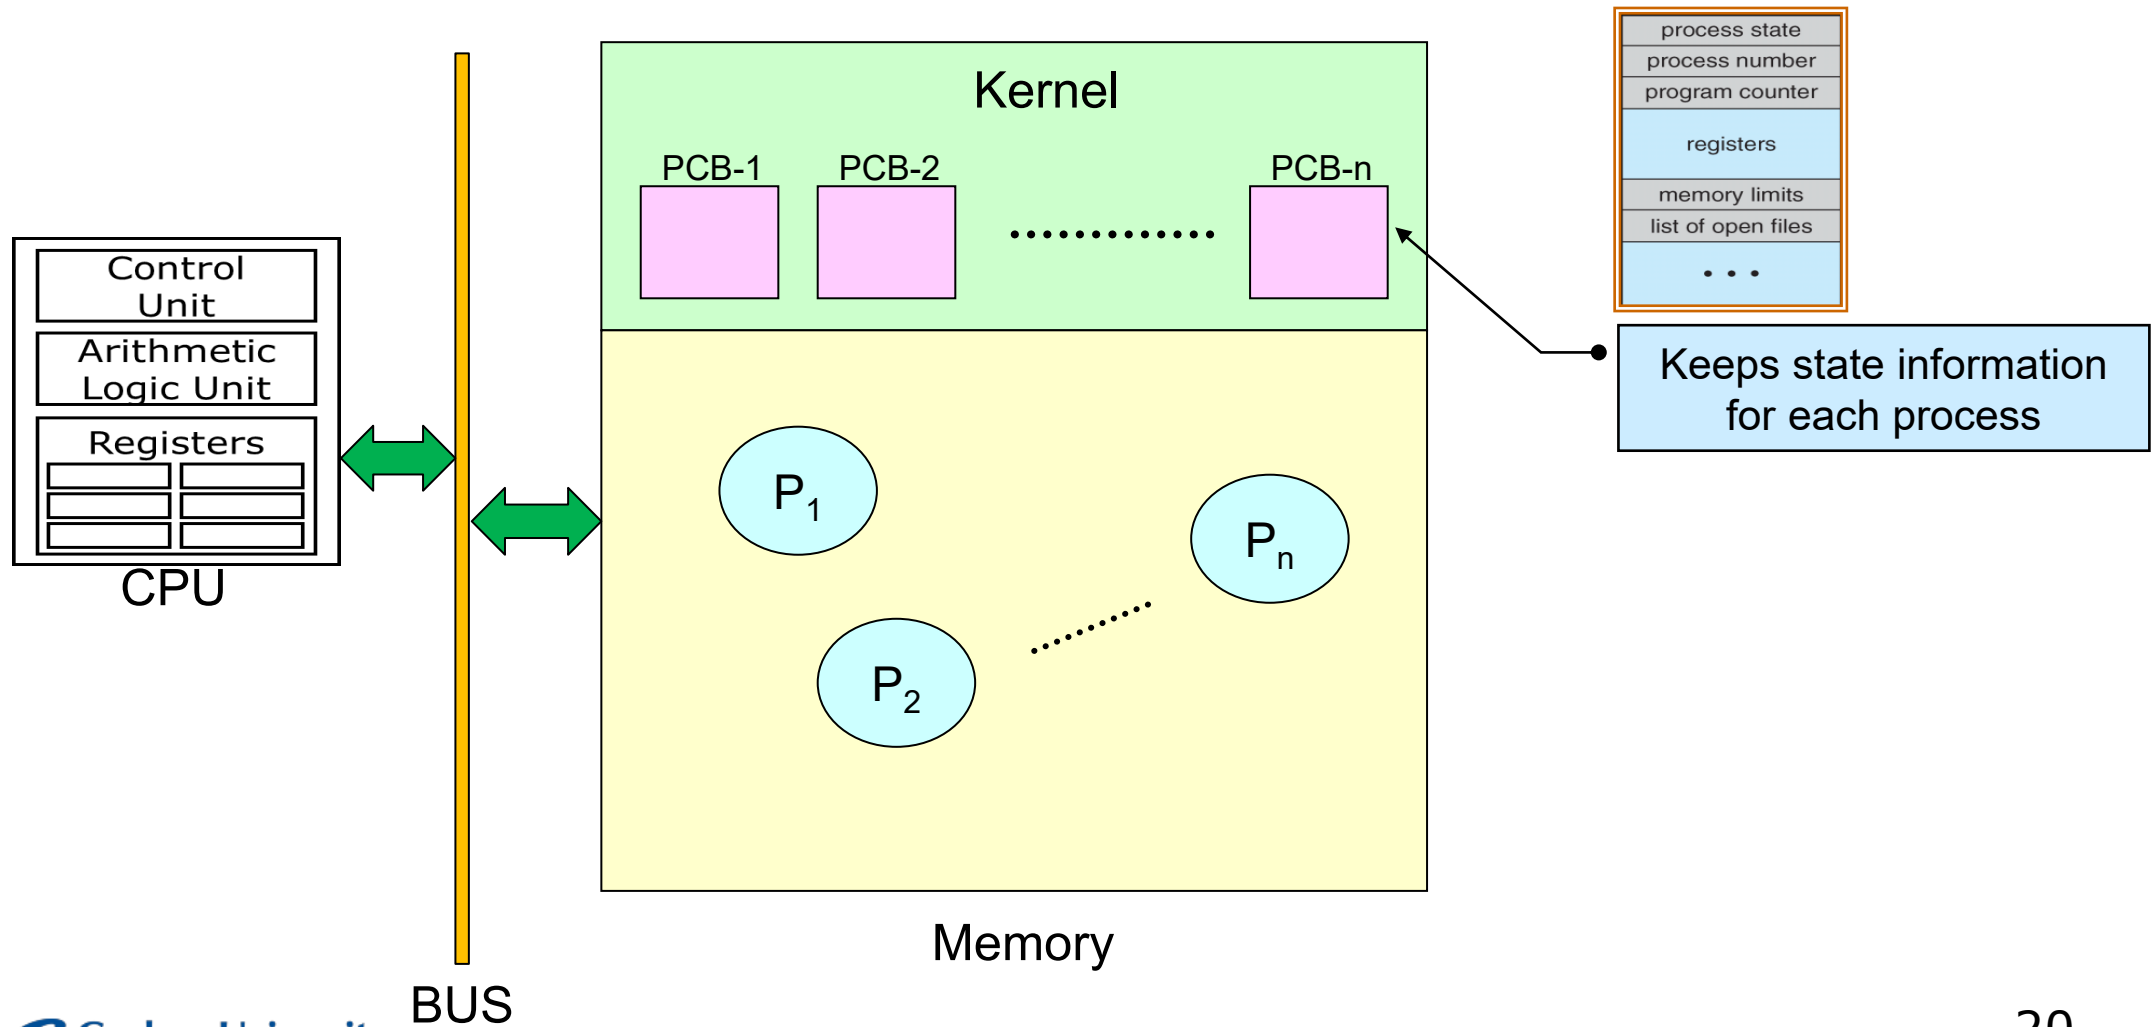
Context Switch

- When CPU switches to **run** another process
 - the system must **save the state (context)** of the old process (to resume where we stopped) and
 - load the **saved state** for the new process via a **context switch**

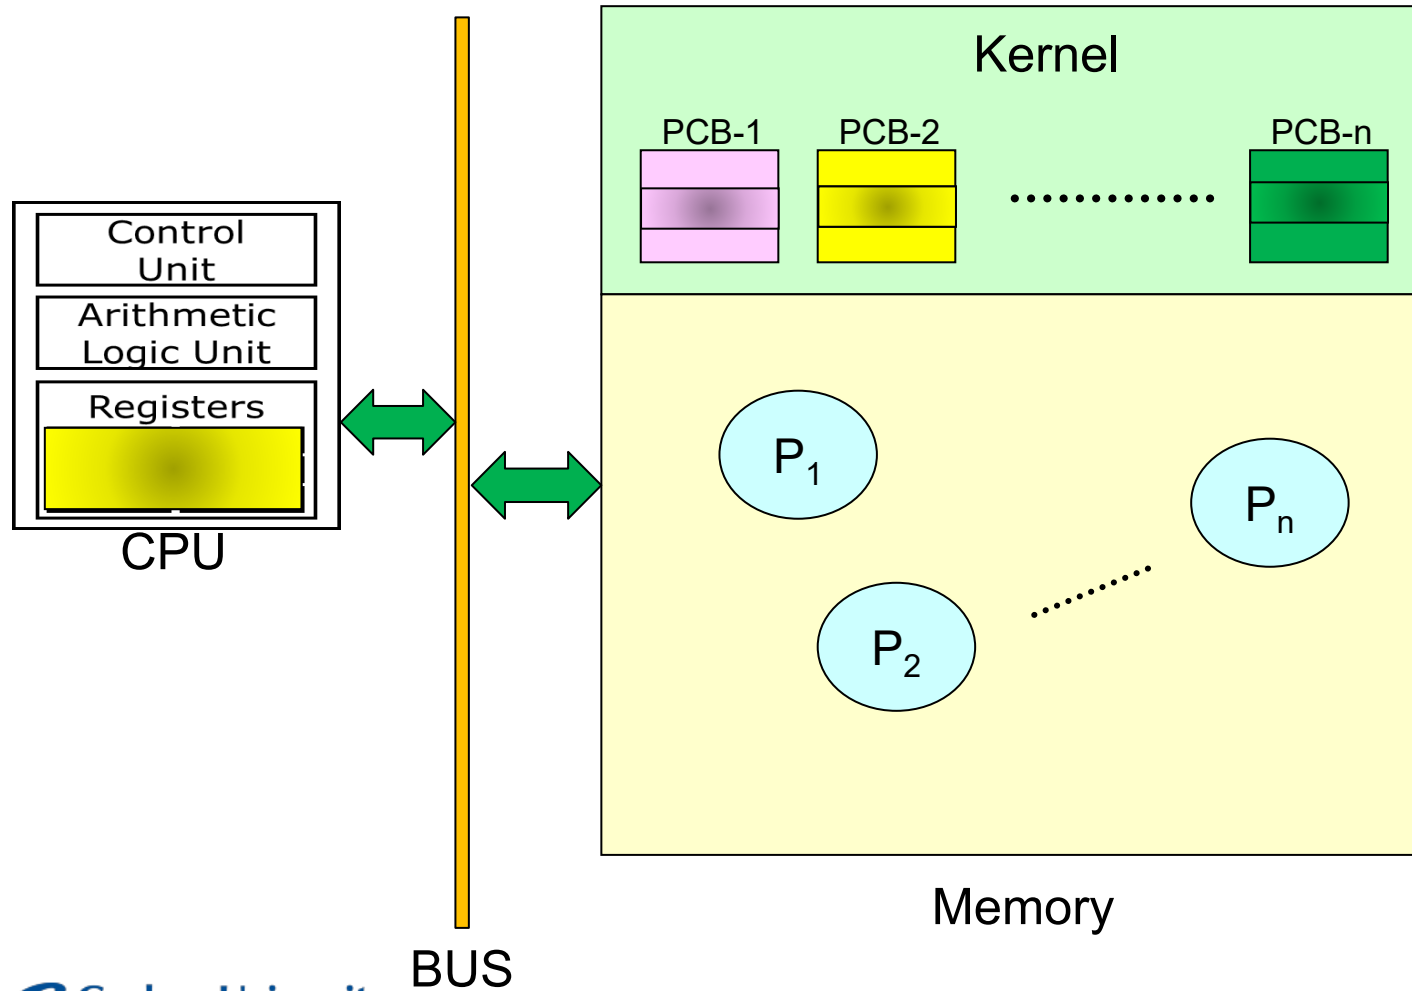
- **Context** of a process represented in the PCB



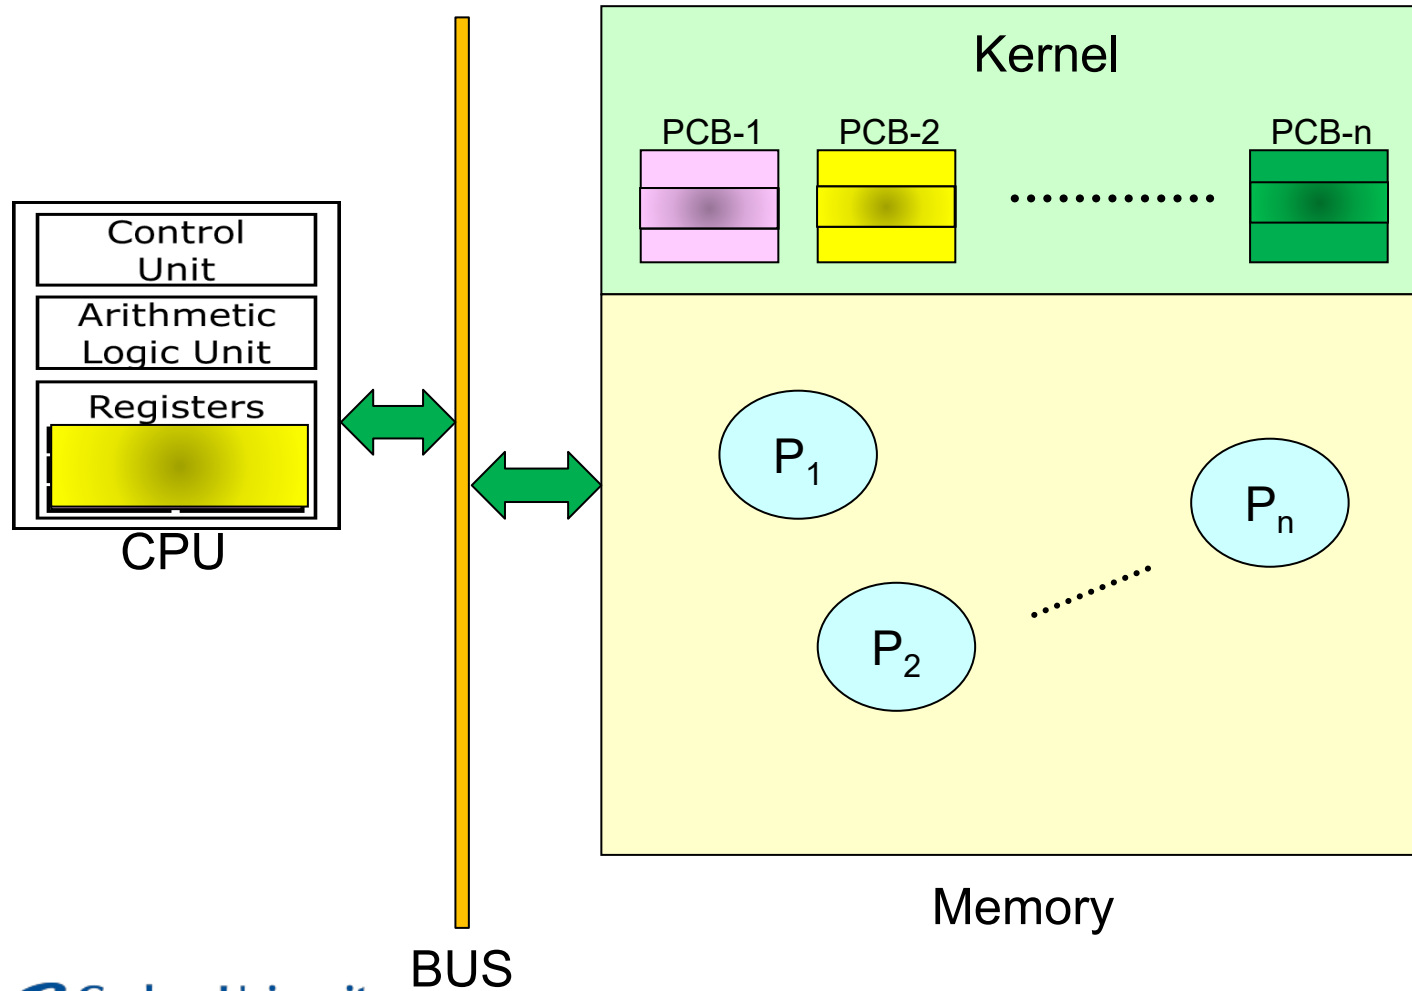
(Cont.)



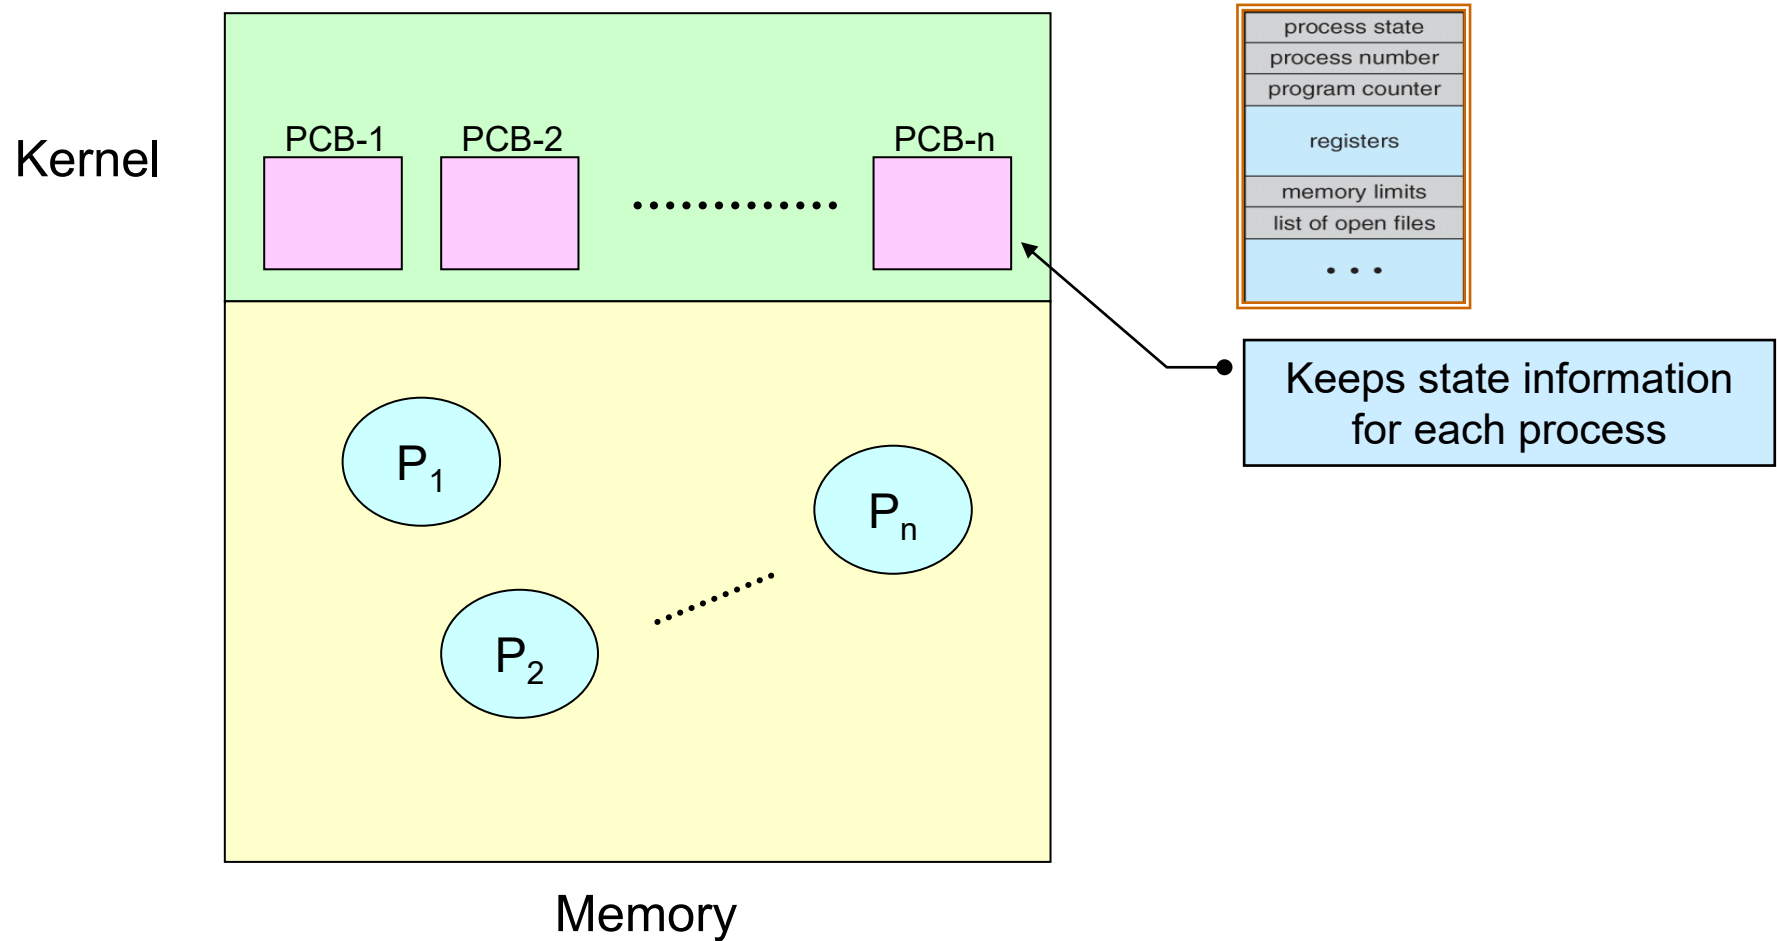
(Cont.)



(Cont.)



(Cont.)



Context Switch Overhead


- Context-switch time is **overhead**; the system (or CPU) cannot do any other work while switching
 - context switch time depends on memory speed, number of registers that need to be copied, ...
 - typically takes several microseconds
- Context-switch overhead is an important factor in CPU scheduling

Chapter 3: Process Concept

- Process Concept
- Process scheduling
- **Operations on Processes**
- Interprocess Communication

Operations on Processes

- Processes can execute concurrently, thus they may be **created** and **deleted** dynamically.
 - Operating System must provide mechanisms for process creation, termination, and so on.
 - Generally, process identified and managed via a unique **process identifier** (**pid**), typically an integer number
- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes

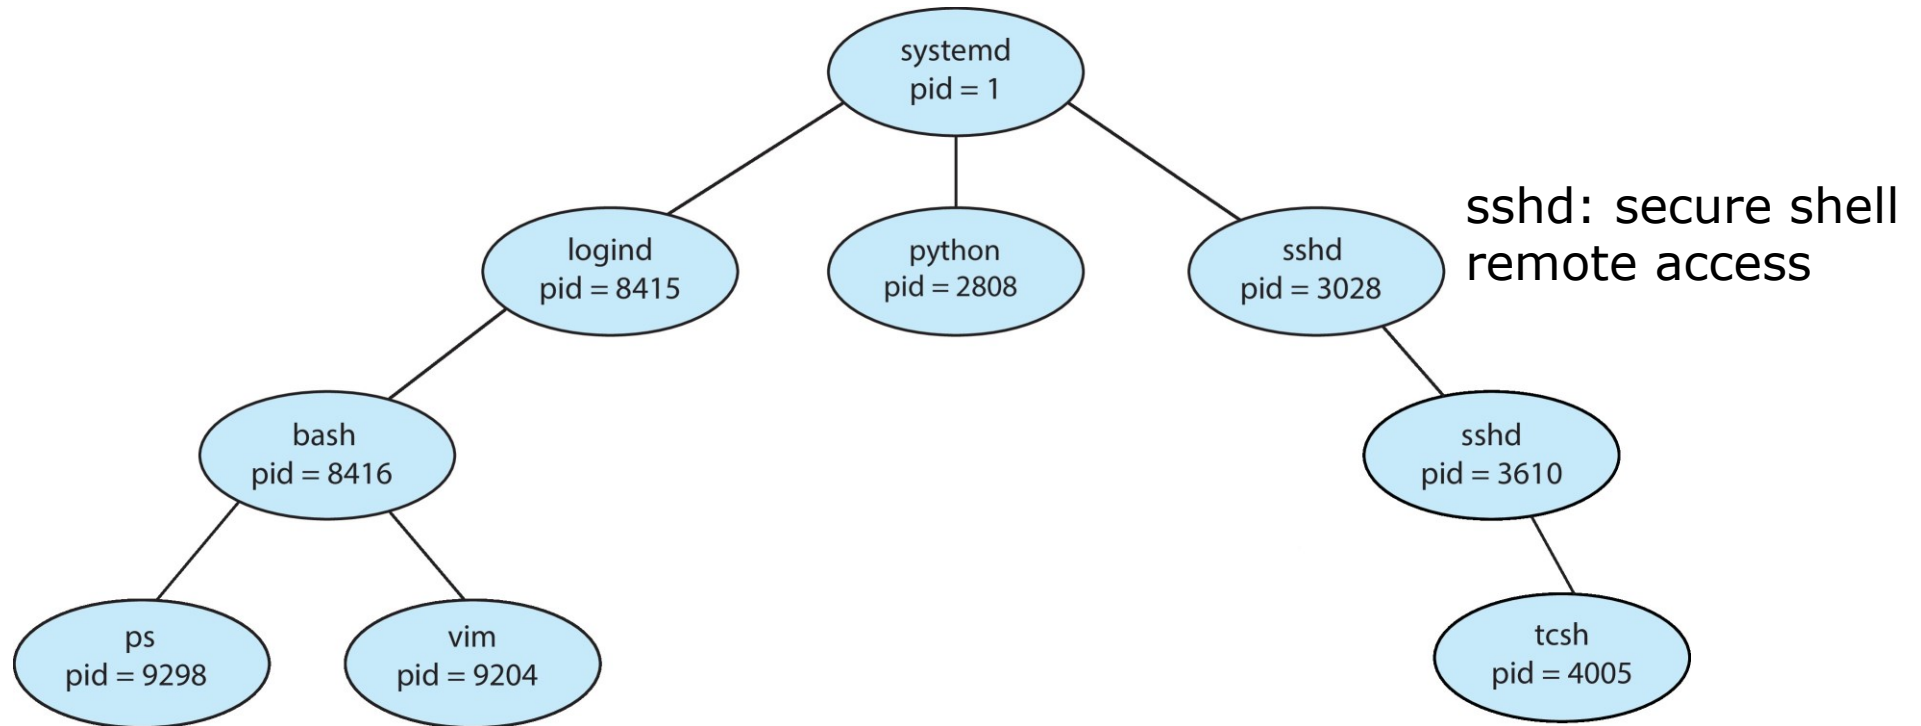


- So who created this universe (i.e., the first process)?



A Tree of Processes in Linux

Root process for all user processes

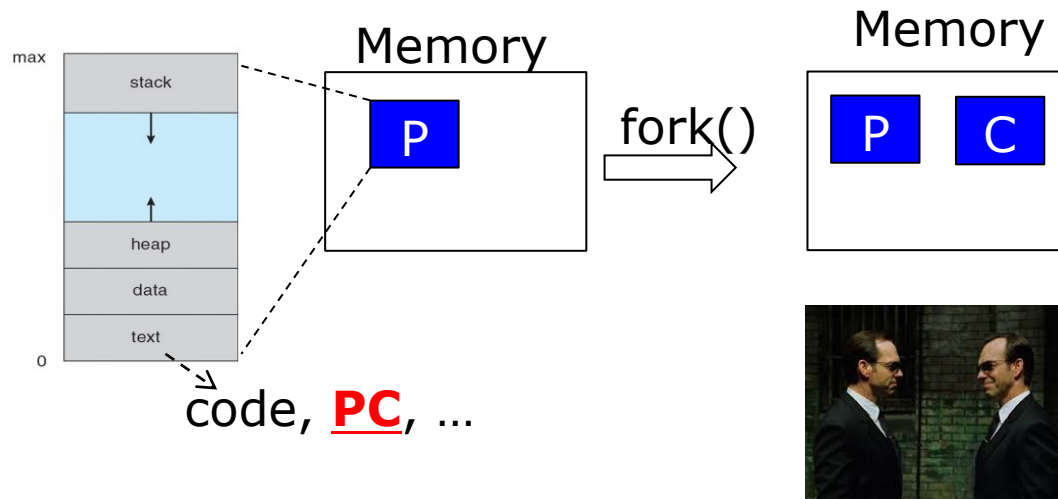


type ps -el in Linux shell

UNIX examples: Process Creation


■ fork()

- **fork()** system call creates a new process
- New process consists of a copy of parent's memory space
- Both processes continue execution



■ exec()


- **exec()** system call used after a **fork()**, to *replace* the process' memory space with a new program
- loads a new binary file into memory (deletes original memory – copy of parent)



UNIX examples: Process Creation

■ UNIX examples

- `fork()` system call creates a new process
- Both processes (parent and child)
 - continue execution after `fork()`,
 - with one difference: `fork()` return value:
 - Child process: 0
 - Parent process: `pid` of child process (>0)
process identifier (`pid`)



```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
```

```
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
```

```
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
```


```
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
```

```
    return 0;
}
```

Parent

Child

Fork Timeline Example



Parent (100)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Parent

Child (150)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();


    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Child

Process Creation

- Memory address space
 - Child duplicate of parent: `fork()`
 - Child has a program loaded into it: `exec()` system call
 - ▶ `exec()` system call used after a `fork()` to *replace* the process' memory space with a new program



```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
```

```
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
```

```
    } else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL)
```

```
    } else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
```

```
    return 0;
}
```

Parent

Child


Process Creation

■ Execution options

- Parent and children: Which executes first?
- Parent can
 - ▶ execute concurrently with child or
 - ▶ wait until children terminate: `wait()` system call returning the pid:

```
pid_t pid; int status;  
pid = wait(&status);
```

```
#include <sys/types.h>  
#include <stdio.h>  
#include <unistd.h>  
  
int main()  
{  
    pid_t pid;  
  
    /* fork a child process */  
    pid = fork();  
  
    if (pid < 0) { /* error occurred */  
        fprintf(stderr, "Fork Failed");  
        return 1;  
    }  
    else if (pid == 0) { /* child process */  
        execlp("/bin/ls", "ls", NULL);  
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        wait(NULL);  
        printf("Child Complete");  
    }  
  
    return 0;  
}
```



Process Termination

- Process asks the operating system to delete itself:
`exit()` system call
 - Return status value (integer) to parent process (via `wait()`)
- Parent may **terminate** execution of children processes:
`abort()` system call
 - E.g., child has exceeded allocated resources, task assigned to child is no longer required
 - If parent is exiting, some operating systems do not allow child to continue if its parent terminates
 - ▶ All children terminated - **cascading termination**



Chapter 3: Process Concept

- Process Concept
- Process scheduling
- Operations on Processes
- Interprocess Communication


Interprocess Communication



- In many cases, processes can **cooperate** with each others
 - Reasons for cooperating processes:
 - ▶ Information sharing – e.g., shared file
 - ▶ Computation speedup – e.g., parallel computing in multi-core environment
- Why can't processes just communicate with each other?
 - A process cannot directly access other process's **memory** – why not?
 - ▶ A: For system _____
- Cooperating processes need **interprocess communication (IPC)** provided by OS via system call

Example: Chrome Browser

- Many web browsers ran as single process (some still do)
 - If one web site causes trouble (e.g., JavaScript, Flash, HTML5), entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 categories
 - **Browser** process manages user interface, disk and network I/O (1 process created)
 - **Renderer** process renders web pages (deals with HTML, Javascript), new process for each website opened in a new tab
 - **Plug-in** process for each type of plug-in (e.g. Flash, QuickTime)



Each tab represents a separate process.

Communications Models

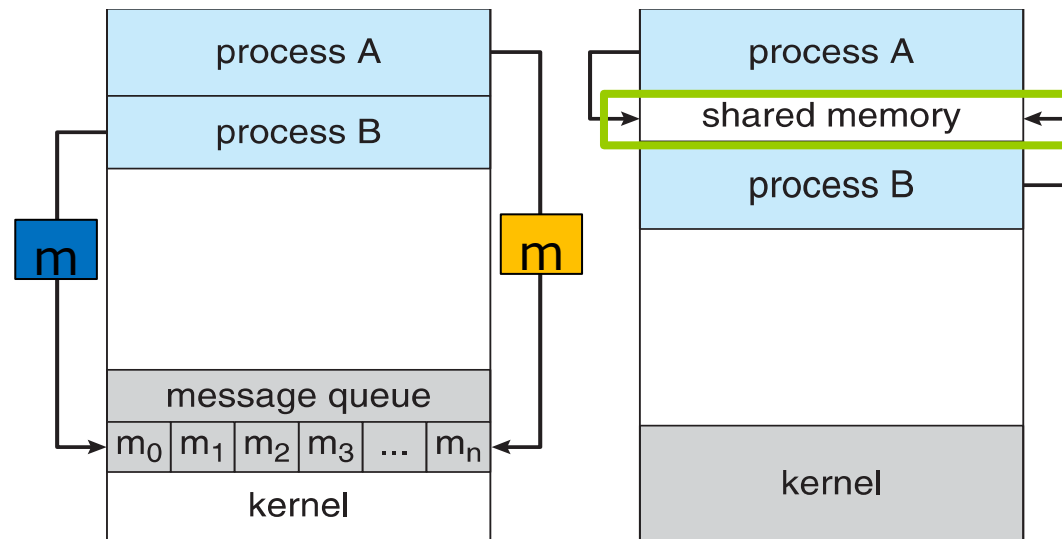
Mechanism

for processes to **communicate** & to **synchronize** their actions

Two fundamental models of IPC

Shared memory

Message passing




(a) Message passing

(b): Shared memory

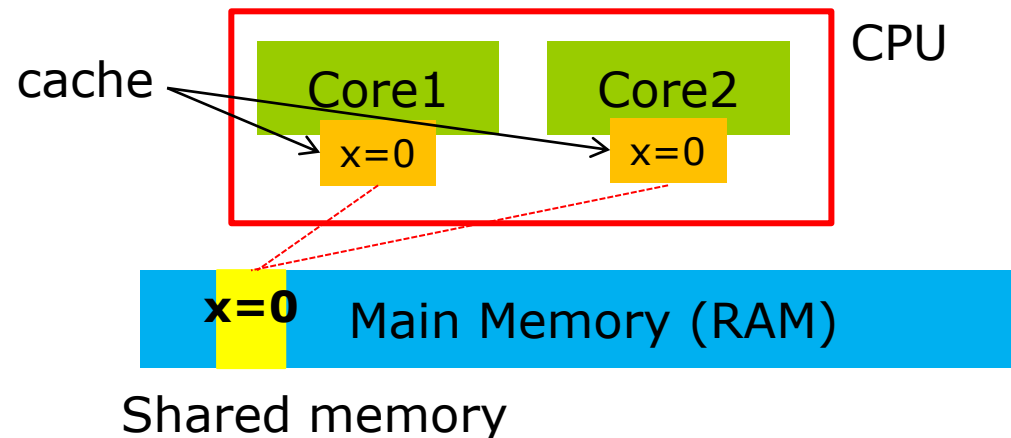
Shared-Memory Systems

- OS prevents one process from accessing another process's memory – Protection
 - **Shared memory:** Two or more processes agree to remove this restriction
- A region of memory that is shared is created by system call
 - Processes exchange information by reading and writing data on shared area
 - All access to shared memory are treated as routine memory access – no need for system call



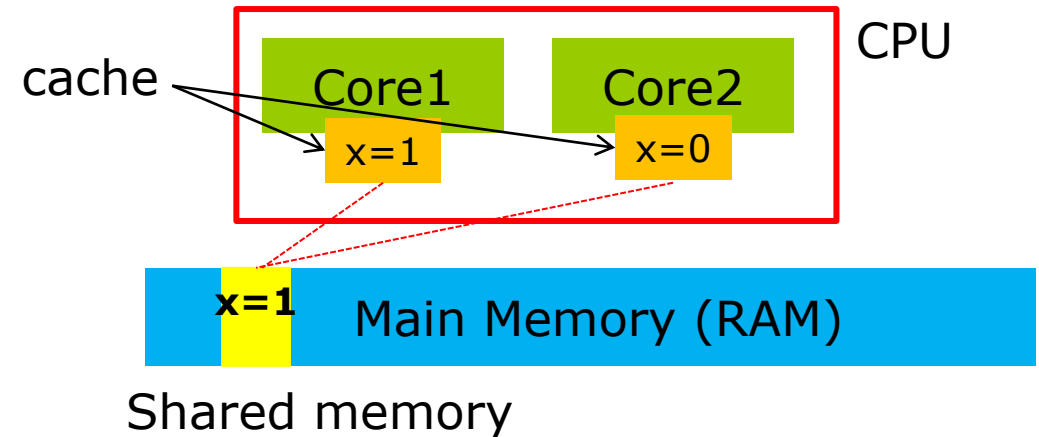
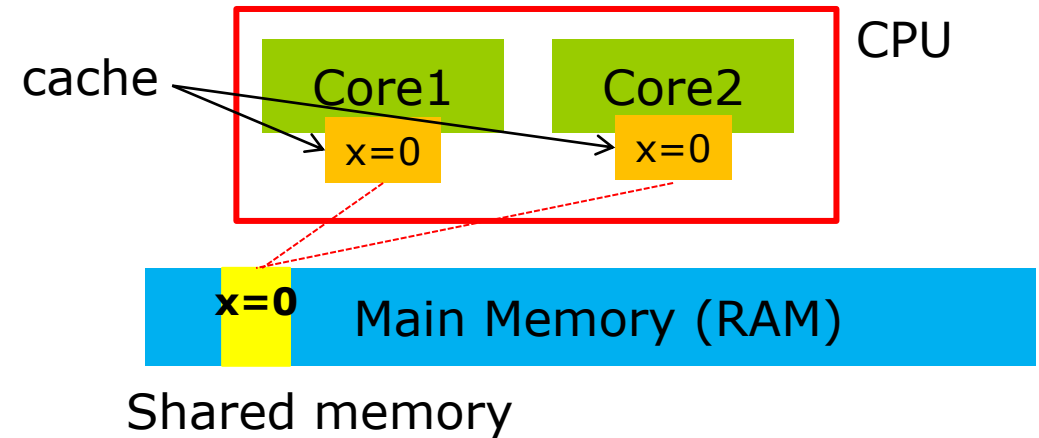
Shared Memory Systems

- **Bigger problem:** What happens if two processes attempt to access the shared memory concurrently?
 - Process synchronization (Ch. 6)
- **Another problem:** Multicore processors
 - Each core have separate cache – cache coherence problem




Multicore processor: Cache coherence problem

1. process1 in core1 reads “x=0” from shared memory – stored in core1’s cache
2. process2 in core2 reads same “x=0” from shared memory – stored in core2’s cache
3. process1 in core1 changes data to “x=1”
 - updated core1’s cache and shared memory to “x=1”
 - But core2’s cache still has “x=0”
 - **Cache coherence problem!**
 - ▶ Usually solved by CPU cache hardware



Interprocess Communication – Message Passing

- Message system – processes communicate with each other by **messages** without resorting to shared variables
- IPC facility provides two operations via system call:
 - **send(message)** **receive(message)**
- If P and Q wish to communicate, they need to:
 - establish a **communication link** between them
 - exchange messages via send/receive system call
 - ▶ need **system call** for every message
 - ▶ More time-consuming compared to shared memory
 - e.g., **Microkernel structure**
 - e.g., **Sockets** (networking), Remote Procedure Call (RPC: cloud computing)





<https://www.proactiveinvestors.co.uk/>