


Chapter 4: Multithreaded Programming

School of Computing, Gachon Univ.
Joon Yoo



Chapter 4: Multithreaded Programming

- Overview
- Benefits of Multithreading
- Multithreading Models
- Thread Libraries
- Process vs. Thread

Objectives

- To introduce the notion of a **thread**—a fundamental unit of **CPU utilization** that forms the basis of multithreaded computer systems
- To discuss the APIs for the Pthreads and Java thread libraries

Multi-Process (Multi-tasking)

- User opens the same program two times
 - e.g., opens two web browsers, opens two Word files, Web server executes similar tasks for each user
- The two programs will be
 - executing **same code**,
 - may want to **share data**


Multi-Process (Multi-tasking)

- Processes are not very efficient
 - Each process has its own PCB and OS resources
- Processes don't (directly) share memory
 - Each process has its own address space
 - Need IPC (shared memory, message passing)
- Can make it more efficient by **sharing**?


Process A




Process A's PCB



Process B




Process B's PCB




What can we do? Let us share...

- What can we **share** across all of these processes?
 - Same code: generally running the same or similar programs
 - Same data
- What is private to each process? (i.e., what can we **not share**?)
 - Execution context: CPU registers, stack, and program counter (PC)

Process A




Process A's PCB



Processes and Threads

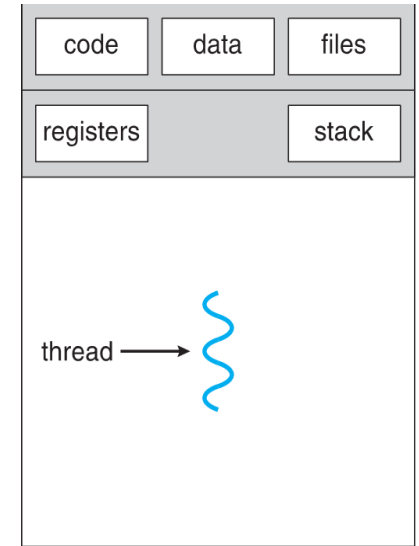
- **Thread?**
 - A thread (or lightweight process) is a basic unit of **CPU scheduling**
 - A process is just a “container” for its threads
 - Each thread is bound to its containing process
- Each thread has its own (**not share**)
 - **stack, CPU registers, PC**
- All threads within a process **share** memory space
 - **text, data, and OS resources**
 - **Threads in same process** can communicate directly via shared memory (no need for IPC)




Process

Single-threaded vs. Multi-threaded Process

- Simple programs can have one thread per process
 - *single-threaded* process
- Complex programs can have multiple threads
 - *multi-threaded* process
 - Multiple threads running in same process's address space



single-threaded process



multithreaded process

Chapter 4: Multithreaded Programming

- Overview
- **Benefits of Multithreading**
- Multithreading Models
- Thread Libraries
- Process vs. Thread

Benefits of Multithread


■ Resource Sharing

- All threads in one process share memory resources (code, data) of process
- easier to share than IPC


■ Lighter weight

- lighter weight than process
 - ▶ creation/deletion, context-switching faster
- Easier to create/delete 10 threads than 10 processes

Process A




Process A's PCB



Thread vs. Process creation

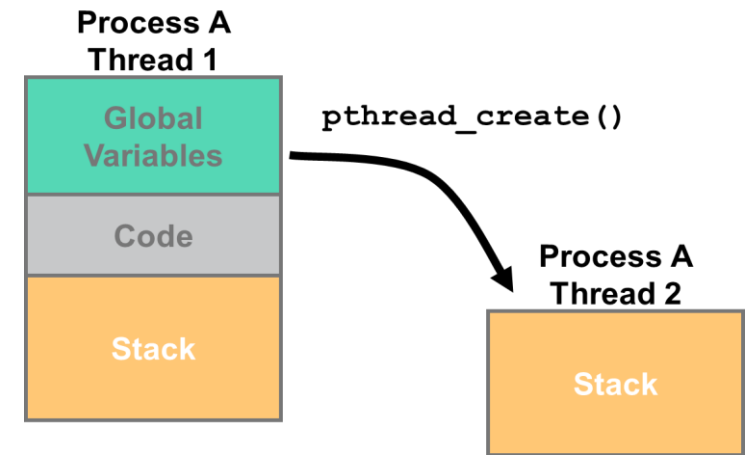
Creation of a new process using *fork()* is *expensive* (time & memory).

Need new PCB



A thread creation using `pthread_create()` does **not** require a lot of memory or startup time.

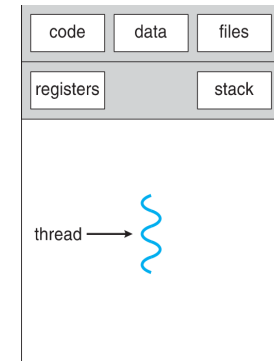
No need for new PCB




Benefits of Multithread

■ Non-blocking System Call (Responsiveness)

- may allow continued execution if part of **process** is blocked
 - ▶ why blocked? – time consuming operation (e.g., I/O such as printing, network)
 - ▶ single-threaded process: wait until I/O operation is complete
 - ▶ multithreaded process: one thread must wait, but another thread in same process can continue




single-threaded process





multithreaded process

Benefits of Multithreading: Multicore Programming

- **Concurrent** Execution on a Single-core System




- **Parallel** Execution on a Multicore System



Multicore Programming

- Multicore Programming provides parallelism!
 - **Concurrency** supports more than one task making progress – *Single-core* processor can provide concurrency
 - **Parallelism** implies a system can perform more than one task simultaneously – Need *Multi-core* processor



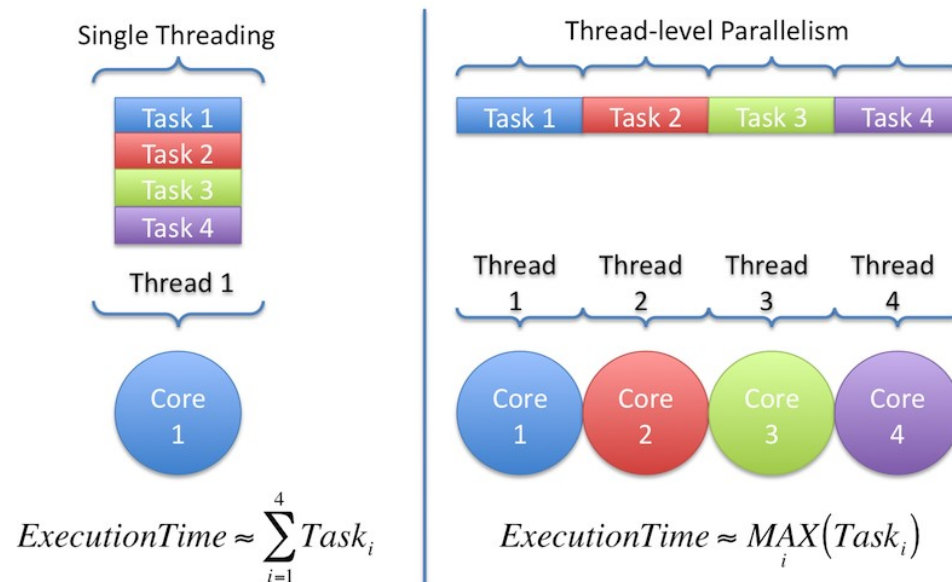
- By using multithreading, one process can use multiple cores!
 - Q: Multithreading gives (concurrency/parallelism) to a single process

Multicore Programming



- Allows one process to use **multiple cores**

- A multithreaded process can take advantage of Multicore CPU architectures
- A process can run many threads in parallel on different processor cores



- Task: Add 1 to 4,000,000
 - Divide into 4 subtasks
 - ▶ Task1: add 1 to 1,000,000
 - ▶ Task2: add 1,00,001 to 2,000,000
 - ▶ Task3: add 2,00,001 to 3,000,000
 - ▶ Task4: add 3,00,001 to 4,000,000
- Say a core takes 10 seconds to add 1 million numbers
- Single threading: Give all 4 tasks to 1 core
 - Takes _____ seconds
- Multithreading: Give one task to each core (core 1~4)
 - Takes _____ seconds


Chapter 4: Multithreaded Programming

- Overview
- Benefits of Multithreading
- **Multithreading Models**
- Thread Libraries
- Process vs. Thread

Kernel Threads

■ Kernel threads

- Threads supported by the **Kernel**
 - ▶ created by `thread_create` system call
 - ▶ each thread needs **thread control block (TCB)**
- **Pros:** kernel knows the thread, so...
 - ▶ **Parallelism:** Can run multiple threads on **multi-core**
 - ▶ **Concurrency:** another thread can run when one thread makes blocking system call (e.g., I/O request)
- **Cons:** kernel knows the thread, so...
 - ▶ every thread operation must go through kernel; heavy weight
 - ▶ Syscall 10x-30x slower than user threads



Control Blocks

Thread Control Block (TCB)

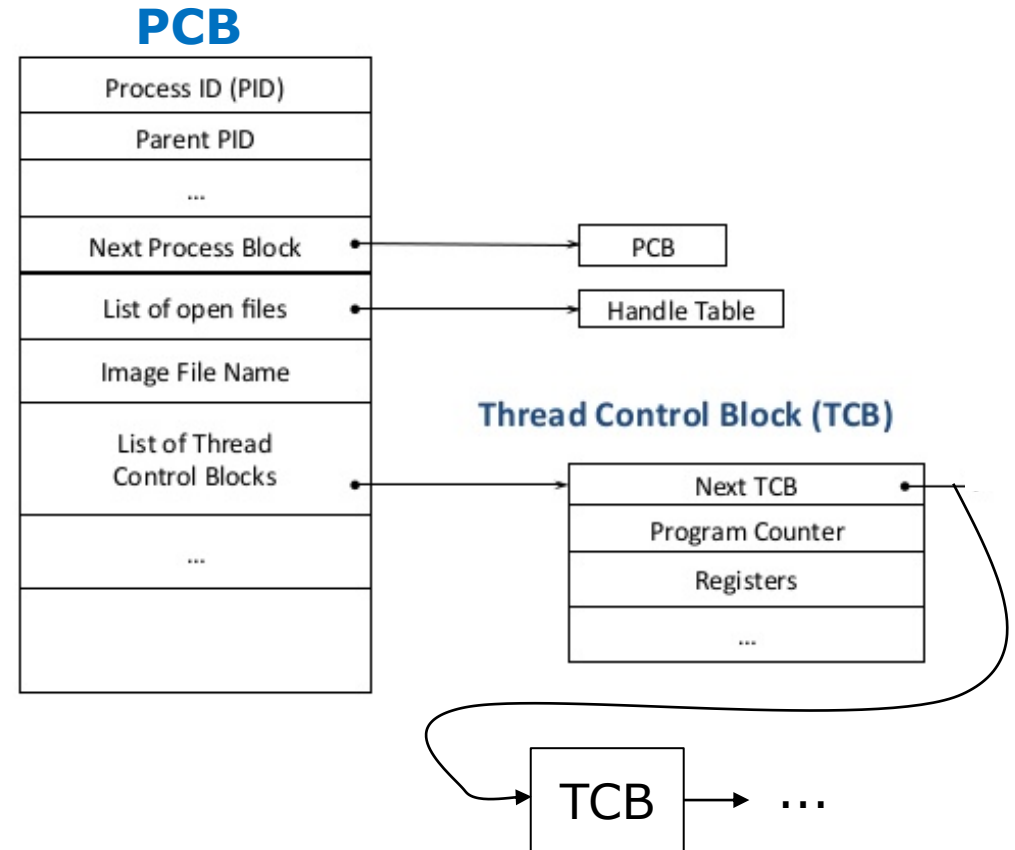
Created for each kernel thread

Contains Program Counter (PC), and registers

Other resources shared

Ready queue is now a list of **TCBs** waiting for CPU resource

CPU Context switching is done for **Threads**



User thread

■ User Thread (=green threads)

- Implement thread in user library
- created/managed *without* kernel support: no need for system call
- One kernel thread per process, many user threads mapped to single kernel thread


Pros: kernel doesn't know the user thread, so...

Thread management is done by the thread library in user space

Fast and efficient (10x-30x faster than kernel thread) – no syscall

■ **Cons:** kernel doesn't know the user thread, so...

- A thread makes a system call - one thread blocking causes all threads in process to **block**
- Multiple user-level threads may **not** run in parallel on multicore system



} why?

Chapter 4: Multithreaded Programming

- Overview
- Benefits of Multithreading
- Multithreading Models
- Thread Libraries (Pthread)
- Process vs. Thread

Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- Common in UNIX operating systems (Linux & Mac OS)

Multithreaded C program using the Pthreads API

```
#include <pthread.h>
#include <stdio.h>
```

```
int sum;          /* this data is shared by the thread(s) */
void *runner( void * param ); /* the thread */
```

②

```
int main( int argc, char * argv[] )
{
```

```
    pthread_t tid;          /* the thread identifier */
    pthread_attr_t attr;     /* set of thread attributes */

    /* set the default attributes */
    pthread_attr_init( &attr );
    /* create the thread */
    pthread_create( &tid, &attr, runner, argv[1] );
    /* wait for the thread to exit */
    pthread_join( tid, NULL );
```

①

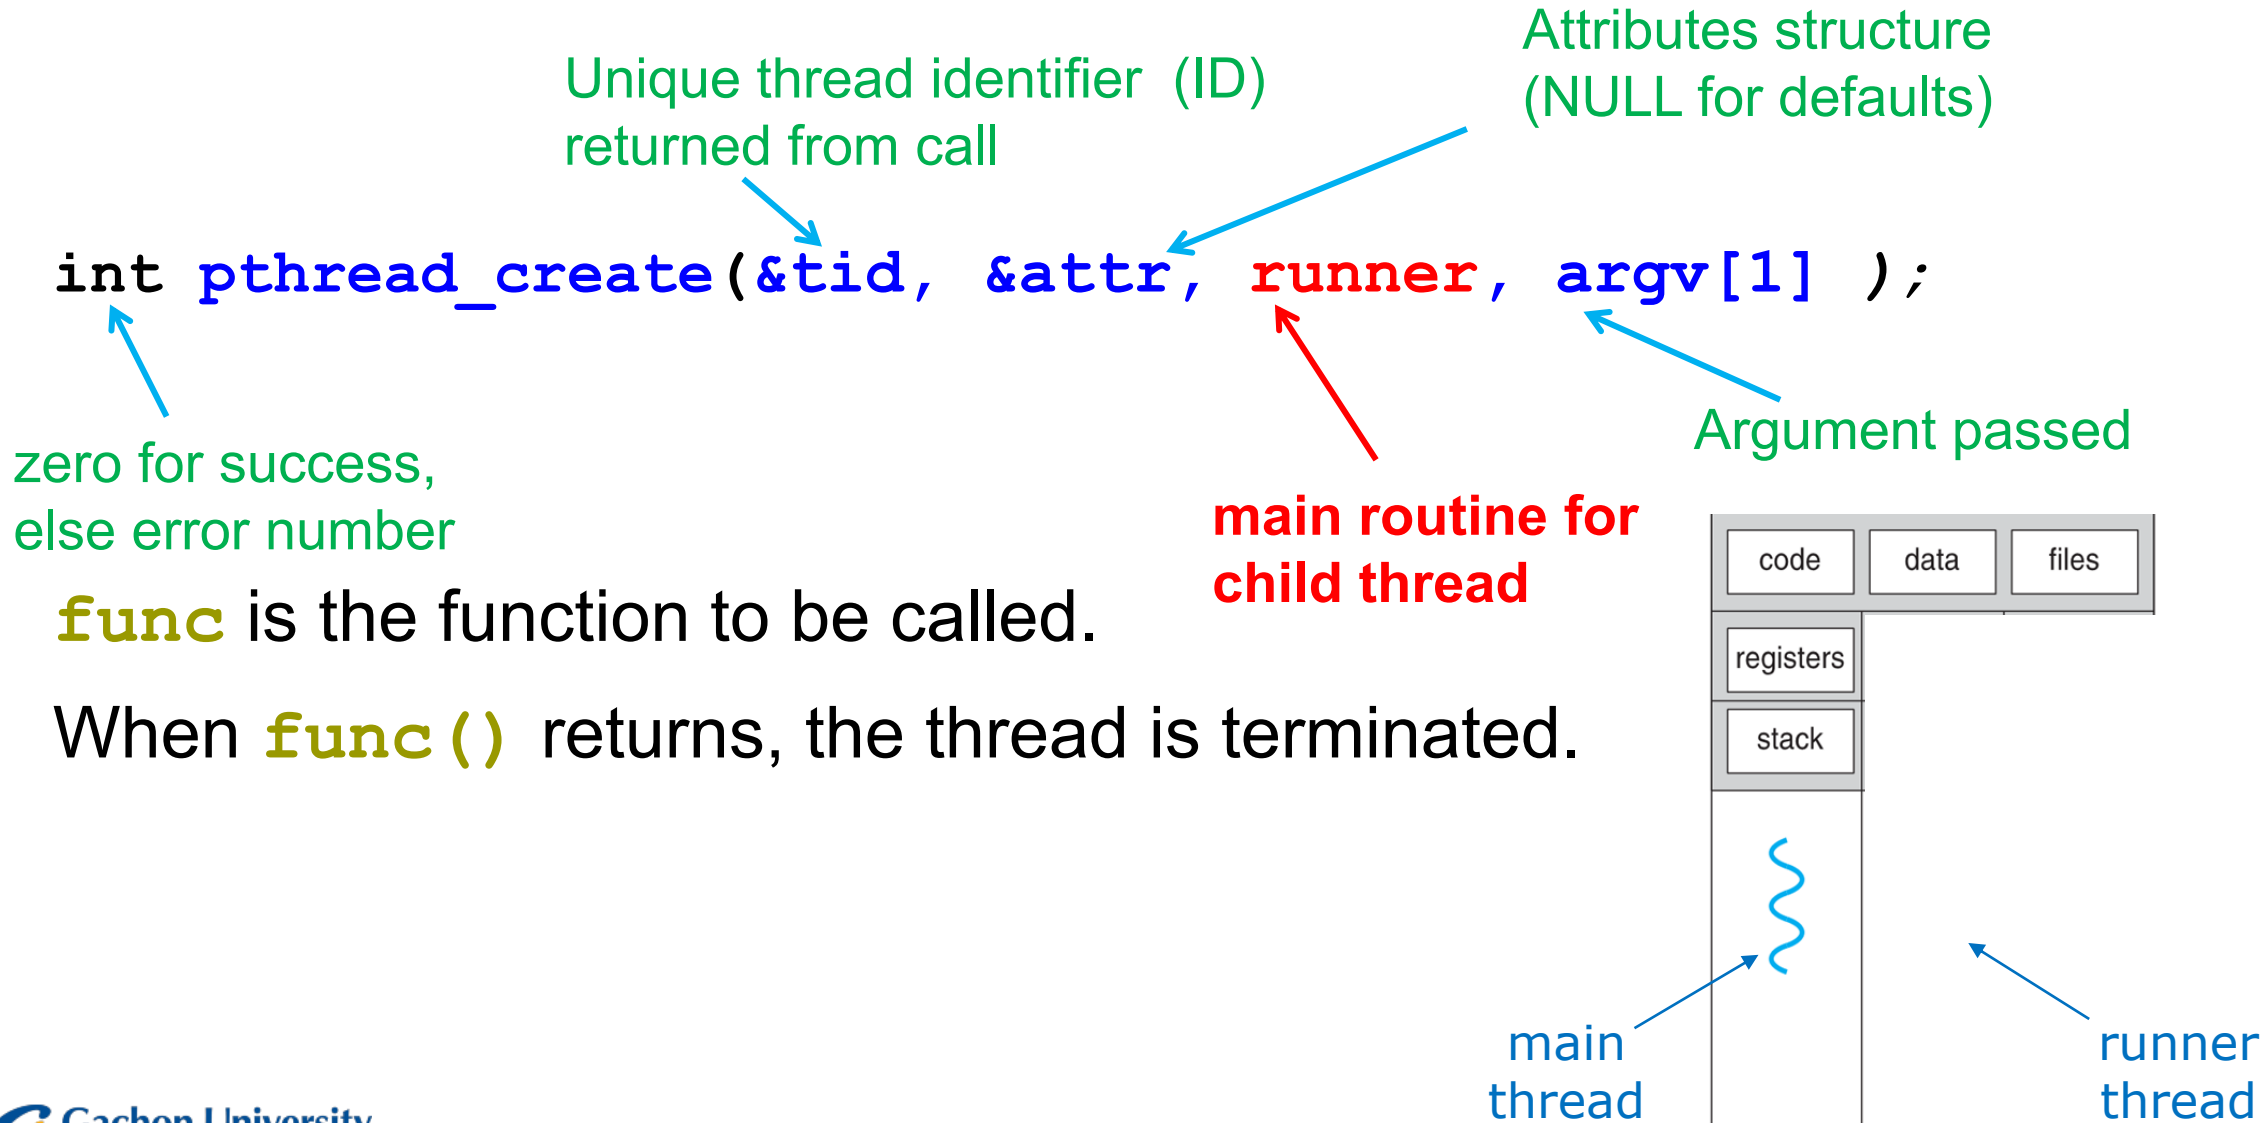
```
    printf( "sum = %d\n", sum );
}
```

③

$$sum = \sum_{i=0}^N i$$

Separate Thread
does this..

① Thread Creation




② Thread Function

`/* The thread will begin control in this function */`

```
void *runner( void * param )
{
    int i, upper = atoi( param );
    sum = 0;

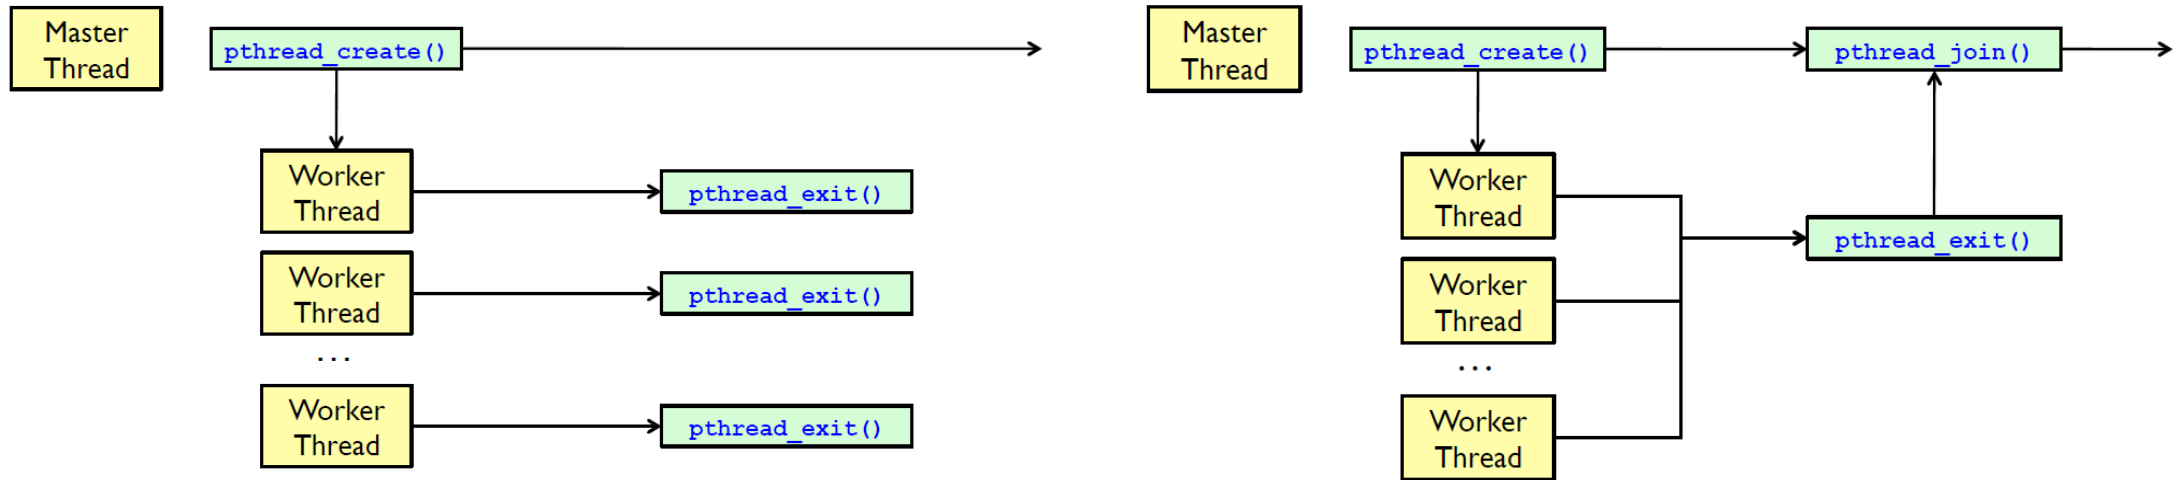
    for( i = 1; i <= upper; i++ )
        sum += i;

    pthread_exit( 0 );
}
```



③ pthread_join()

```
pthread_create( &tid, &attr, runner, argv[1] );  
/* wait for the thread to exit */  
pthread_join( tid, NULL );
```



Suspends parent thread until child thread terminates
similar to _____ system call in process

Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

...


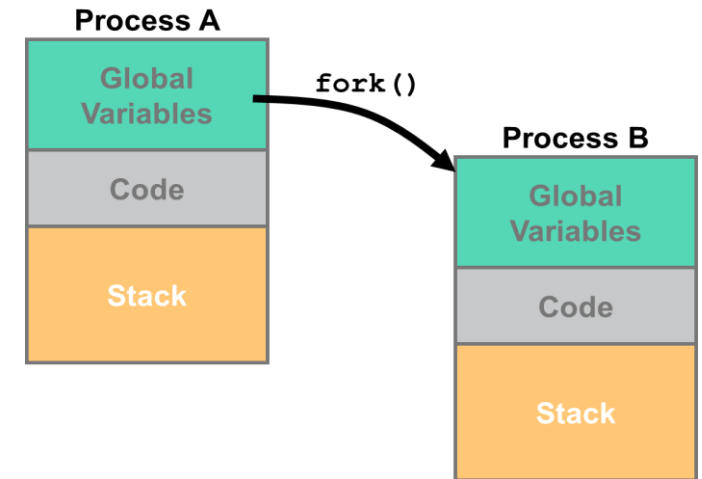
for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

Chapter 4: Multithreaded Programming

- Overview
- Benefits of Multithreading
- Multithreading Models
- Thread Libraries
- Process vs. Thread

Thread vs. Process Creation

- **fork ()**
 - Two separate processes
 - Child process starts from same position as parent (clone)
 - Independent memory space for each process
- **pthread_create ()**
 - Two separate threads
 - Child thread starts from a function
 - Share memory



Process vs. Thread Example

- Shared code:

```
int x = 1; //global variable
void* func(void* p){
    x = x + 1;
    printf("x is %d\n", x);
    return NULL;
}
```

- fork version:

```
main(...) {
    fork();
    func(NULL);
}
```

- threads version:

```
main(...) {
    pthread_t tid;

    pthread_create(&tid, NULL, func, NULL);
    func(NULL);
}
```

Possible output: fork() case 1

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

time



Parent process

Child process

Possible output: fork() case 2

```
int x = 1; //global variable
```

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

time


Parent process

Child process

Possible output: threads case 1

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```



```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Parent thread

Child thread

Possible output: threads case 2

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;
```

```
    printf("x is %d\n", x);  
    return NULL;  
}
```

Parent thread

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Child thread

time



Possible output: threads case 3

```
int x = 1; //global variable
```

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
  
    // interrupted during printf()  
  
    printf("x is %d\n", x);  
  
    return NULL;  
}
```

Parent thread

```
void* func(void* p){  
    x = x + 1;  
    printf("x is %d\n", x);  
    return NULL;  
}
```

Child thread

time



Possible output: threads case 4

Output:

x is 2

x is 2

- Is it a possible output for this example ??
 - Hint: translate $x = x + 1$ into assembly instructions
 - ▶ lw \$t0, 0(\$gp)
 - ▶ addi \$t0, \$t0, 1
 -
 - ▶ sw \$t0, 0(\$gp)
- Bottom line: We cannot predict the results!
 - We need process (thread) synchronization (Ch. 6)

<p>\$t0: data register \$gp: memory address of x lw: load word (from memory) sw: store word (to memory)</p>

Up Next

- Which thread gets to go next when a thread exits running state?
 - Scheduling Algorithm (Ch. 5)
- What happens when multiple threads want to use the shared resource?
 - Synchronization (Ch. 6)



<https://stock.adobe.com/kr/search?k=q%26a>